



CSE 325-Operating Systems


Lecture: 1-4

Dr. Shamim Akhter

Associate Professor

Computer Science and Engineering





Text book





Lecture 1: Introduction





Objective

Overview of the major OS components.





Components of a Computer System



Banking system	Airline reservation	Web browser
Compilers	Editors	Command interpreter
Operating system		
Machine language		
Microarchitecture		
Physical devices		



Application programs

services to the user
(e.g. Web Browser, MS Office)

System programs

services to the hardware
(Disk Driver, assembler, Linker, loader)

Hardware





Microarchitecture- Computer Organization


- The way a given instruction set architecture (ISA) is implemented on a processor

The ISA includes :

- the execution model,
- processor registers,
- address and data formats.

The microarchitecture includes:

- the parts of the processor and
- how these interconnect and
- interoperate to implement ISA.

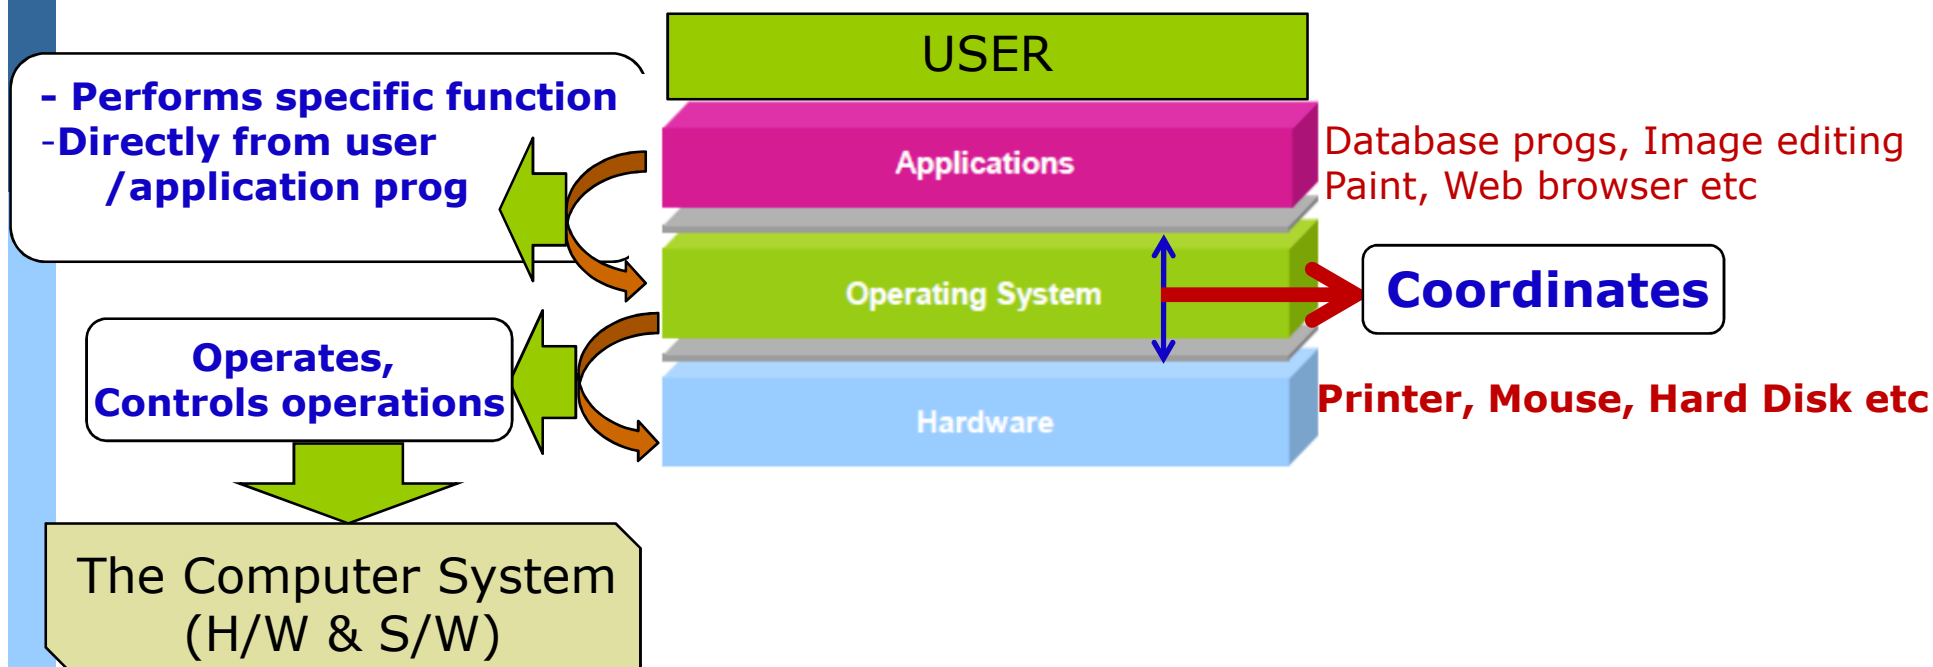


Single Bus Organization



What is an Operating System?

- A **program** that **acts as an intermediary** betⁿ a user & H/W



GOALS:

- » Convenient to use
- » Execute user programs & make solving user problems easier
- » Use H/W in an efficient manner





Operating System Definition

■ OS is a **resource allocator**

- Manages all resources
- Decides between conflicting requests for efficient and fair resource use.




One goal of the operating system is to increase the utilization of resources.

Utilization = useful time / total time

For example, the OS should avoid wasting CPU time because the disk is rotating or wasting switching among tasks. Waiting for I/O.

■ OS is a **control program**

- Controls execution of programs to prevent errors and improper use of the computer. e.g. Memory Protection.






An example comparing life with/without OS

Life with an OS

```
file = open ("test.txt",  
            O_WRONLY);  
write (file, "test", 4);  
close (file);
```

Life *without* an OS

- Blocks, platter, track, and sector
- Where is this file on disk?
Which platter, track, and sectors?
- Code needs to change on a different system





How does a program execute?

■ C Program

```
int A[2]; int C[2]; int B[2];  
P=A[0];  
Q=B[0];  
C[0]=P+Q;
```

First.c

Compile

■ Assembly

```
Load $t0, [A]  
Load $t1, [B]  
ADD $t3, $t0, $t1  
STORE $t3, [C]
```

First.asm

First.o

Linker

■ Executable

```
0001.....00  
0000.....01  
00011.....11  
100.....11
```

Loader

PC=0

IR= 0001.....00

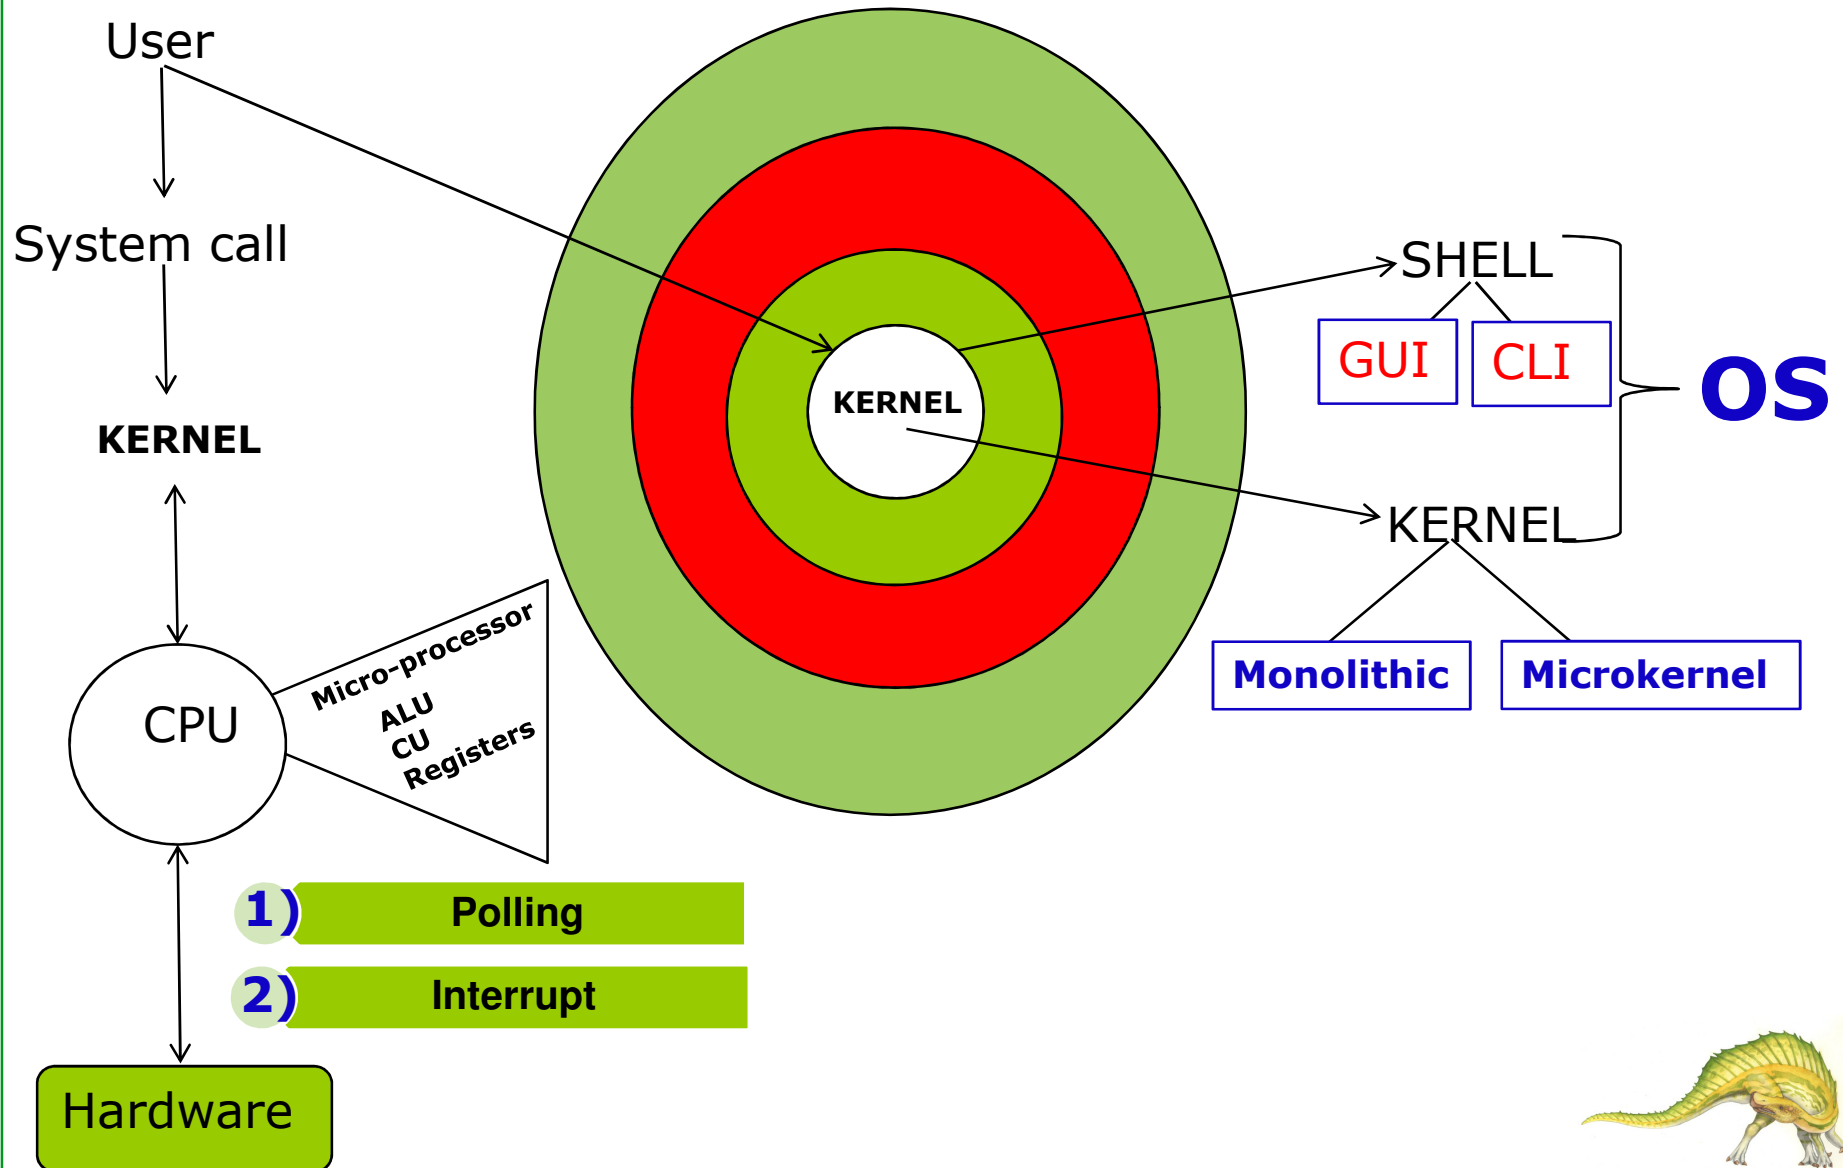
■ Memory

0	0001.....00
4	0000.....01
8	00011.....11
12	100.....11





How does user communicate with OS






The Kernel and The Shell

An Operating system has two parts: **kernel** and **shell**

KERNEL (resides in main memory):

- is the **heart and soul** of an OS. Useful applications and utilities are added over the kernel, then the complete package becomes an OS.
- directly controls the computer hardware by performing **OS services using specific system calls or requests or hardware interrupts.**
- hides the hardware details from the user and application programs.

Example,



- **Linux is a kernel** as it does not include applications-file-system utilities, windowing systems and graphical desktops, system administrator commands, text editors, compilers etc.
- So, various companies added these kinds of applications over Linux kernel and provide their operating system like **ubuntu, centOS, redHat** etc.






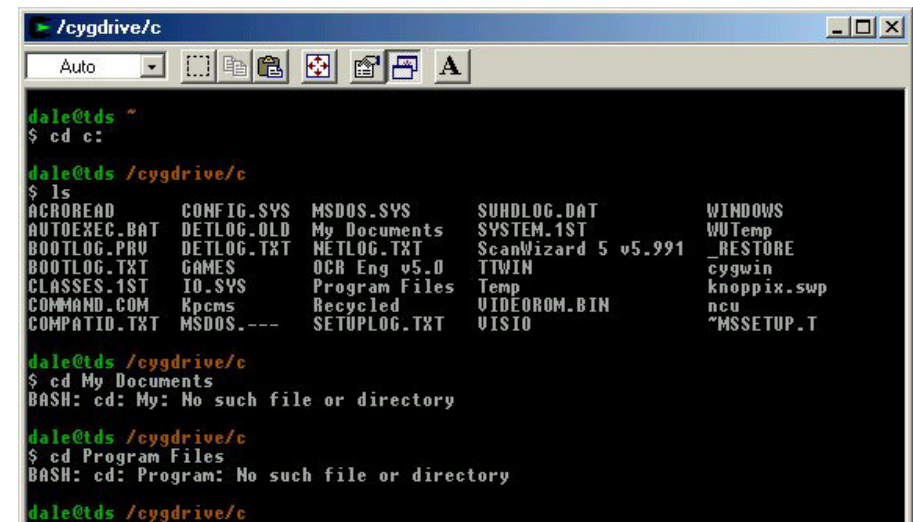
An Operating system has two parts: **kernel** and **shell**
shell (command interpreter): Part of OS

- Serves as the interface between user and kernel.
- User enters commands via shell in order to use the kernel service.

Type of shells



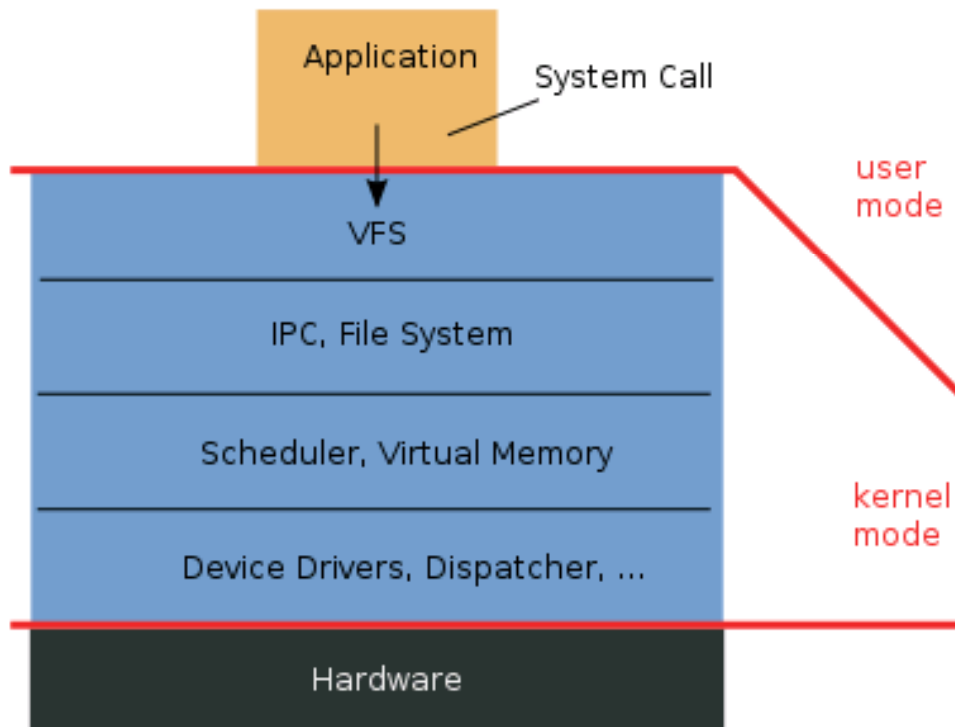
Graphical User Interface (GUI)
Windows/Mac OS



Command line interface (CLI)
Dos/Linux

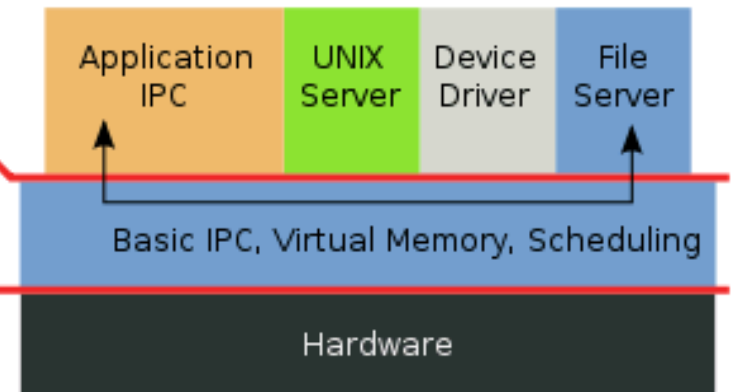


Monolithic Kernel based Operating System



Microkernel based Operating System

Close





Kernel Types

Monolithic Kernel	Micro Kernel
<p>All the parts of a kernel Space</p> <ul style="list-style-type: none"> - Scheduler, file system, memory, device drivers etc. are in one unit within the kernel. 	<p>Only the important parts are in kernel space</p> <ul style="list-style-type: none"> - IPC, basic scheduler, basic memory management, basic I/O primitives etc. - Others (including device drivers) are in user space. Except clock driver.
Signals & sockets to ensure IPC	-Message Passing system to ensure IPC
<p>Advantages</p> <ul style="list-style-type: none"> - Faster processing 	<p>Advantages</p> <ul style="list-style-type: none"> - Crash resistant - Smaller size - Portable - No need recompiling(for a new feature)
<p>Disadvantages</p> <ul style="list-style-type: none"> -Crash insecure -Porting Inflexibility -Add a new feature needs recompiling whole -Kernel size explosion 	<p>Disadvantages</p> <ul style="list-style-type: none"> - slower processing due to additional message passing - 2%-4% slower
DOS, Unix, Linux	Symbian, Mac OS X, MINIX, Windows NT


Tanenbaum VS Torvalds Debate (Micro Kernel VS Monolithic Kernel)

<http://www.oreilly.com/openbook/opensources/book/appa.html>,

<http://www.cs.vu.nl/~scst/reliable.cs/>



OS Organization

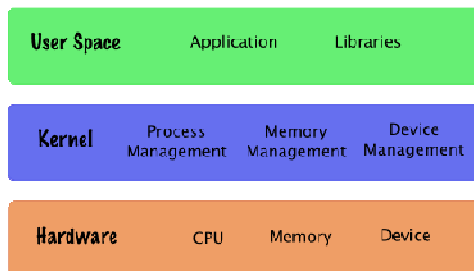




System call and dual mode operation

Dual-mode operation allows OS to protect itself and other system components.

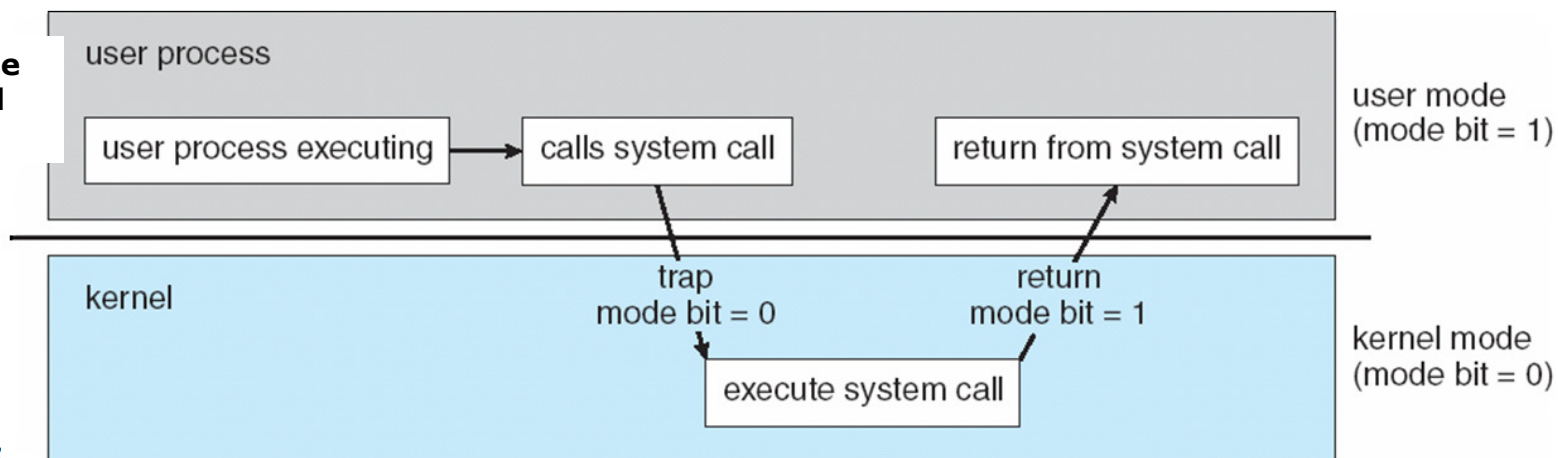
- User mode
- Kernel/Supervisor/Monitor/System Mode.



Mode bit @ Process Status Word (PSW)-Resister

- Provides ability to distinguish when system is running user code or kernel code.
- Some instructions designated as **privileged**, only executable in kernel mode. Example: load/save from protected memory, initiate I/O, timer management, and interrupt management etc.
- System call changes mode to kernel, return from call resets it to user.

System Boot time the H/W starts in kernel mode





Working with User and Kernel mode

PC

@Dr. Md. Shamim Akhter, CSE, EWU

`scanf("%d",&a); // $t0 holds a value`

0	move \$v0, 5
1	syscall
2	store \$v0, 0(\$t0)

IVT- Interrupt Vector Table @ RAM-0000:0000H	
Service #	Code Address
....
005	0058
019	0090

	RAM
0150	syscall (SWI)
0160	
0170	store \$v0, 0(\$t0)
0058	Device driver (SWI)
0070	Allocates kernel buffer address To receive data
0074	GOTO RA=090
0078	GOTO RA=0160
0090	I/O command(control, test, read, write), Communication with H/W Controller
	GOTO RA=0078

1
kernel I/O
subsystem
Service call
handler
(code)

2

mode bit is set to kernel mode(0).

3

4


Modified figure-13.13 @ galvin





Different Method of CPU Interruptions

An operating system is interrupt driven



Interrupts (Hardware events)	<ul style="list-style-type: none"> external hardware event (external to the CPU). Higher priority
Exception (Software events)	<ul style="list-style-type: none"> automatically generated unexpected events occur in response to some exceptional condition. illegal program actions that generate an interrupt.
Traps (Software Interrupts)	<ul style="list-style-type: none"> a programmer initiated and expected transfer of control to a special handler routine. user program can ask for an operating system service unconditional –control always transfer to pre-defined procedure.
System Call (Software Interrupts)	<ul style="list-style-type: none"> simply sets up some registers with the needed system call/ OS service number, and execute the Trap/software interrupt.



How does process pass data to controller?

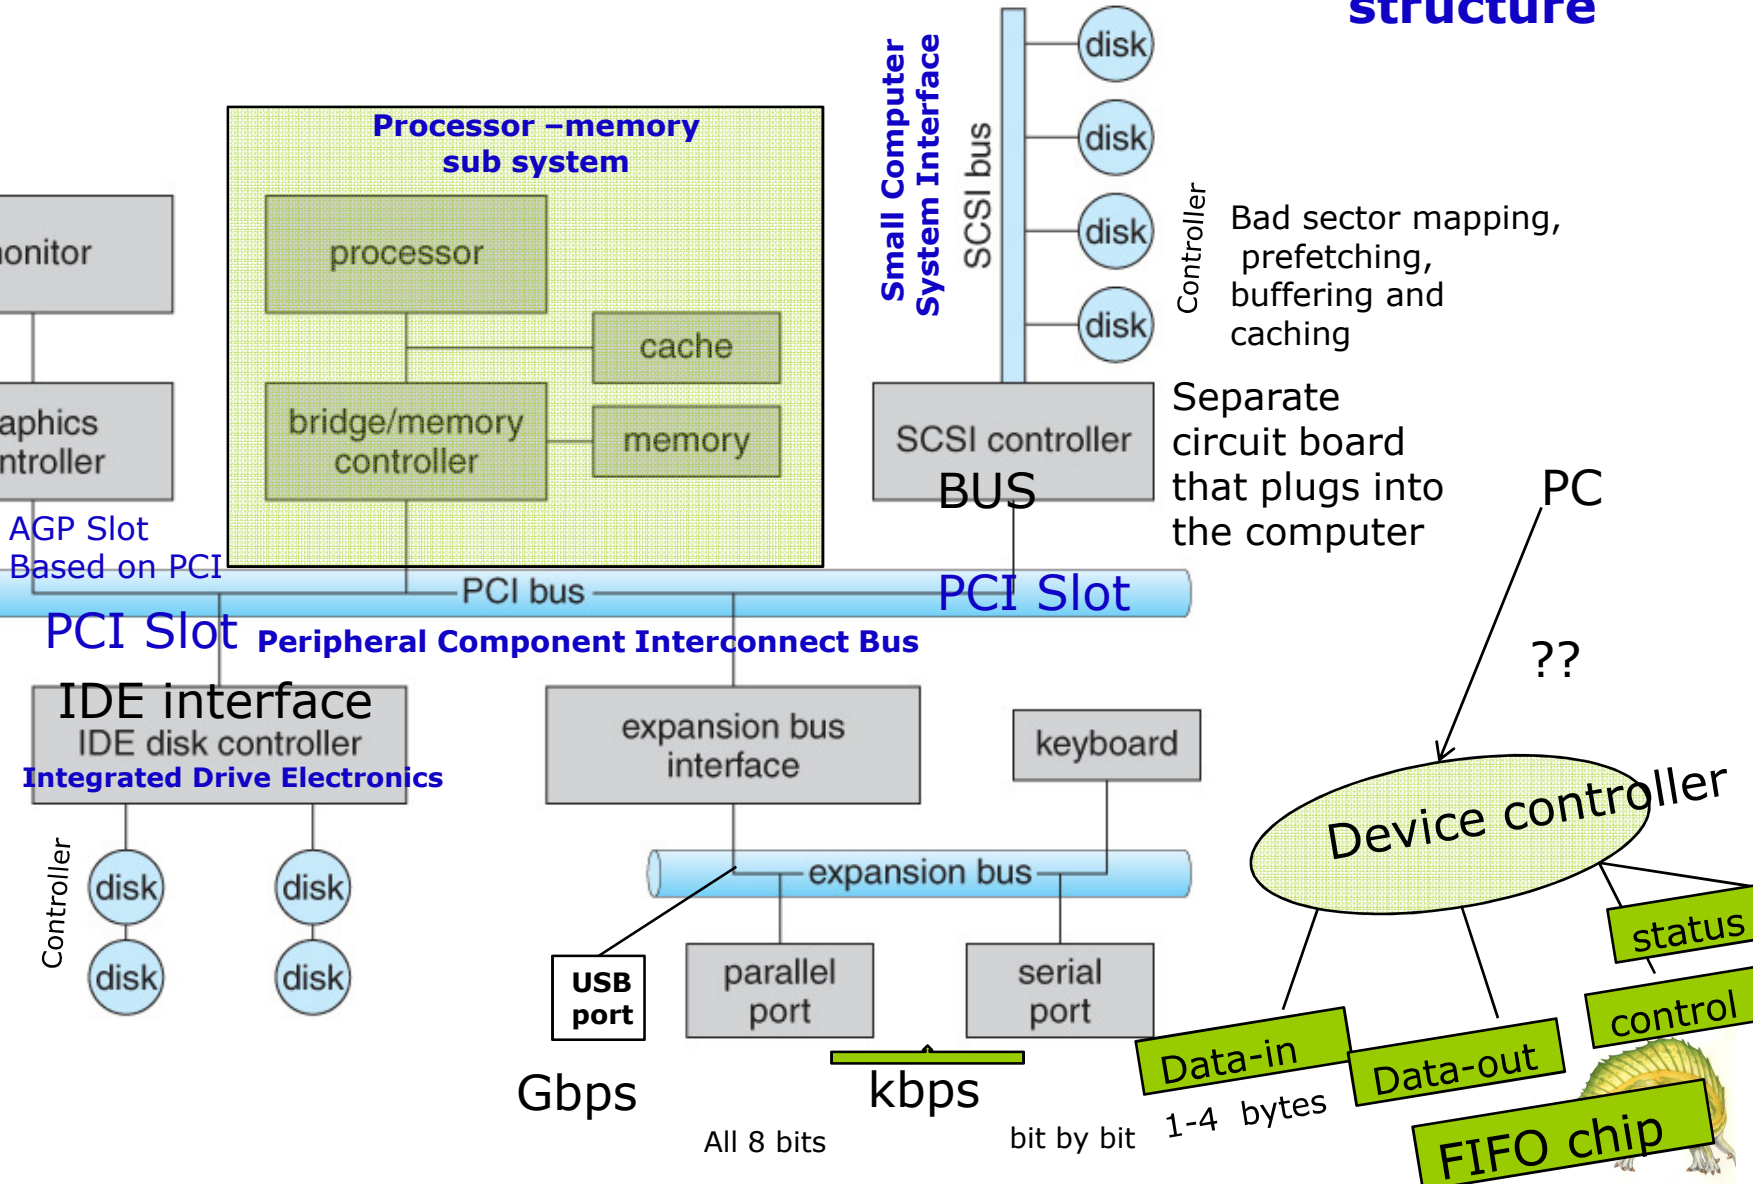




I/O Hardware

Typical PC BUS structure

Accelerated Graphics Port AGP








Device I/O port location on PCs

I/O address range (hexadecimal)	device
000–00F	DMA controller
020–021	interrupt controller
040–043	timer
200–20F	game controller
2F8–2FF	serial port (secondary)
320–32F	hard-disk controller
378–37F	parallel port
3D0–3DF	graphics controller
3F0–3F7	diskette-drive controller
3F8–3FF	serial port (primary)





I/O Mapped: Special I/O instruction in/out instruction

Example from the Intel architecture: `out 0x21, AL`

000–00F DMA Controller


020–021 Interrupt Controller

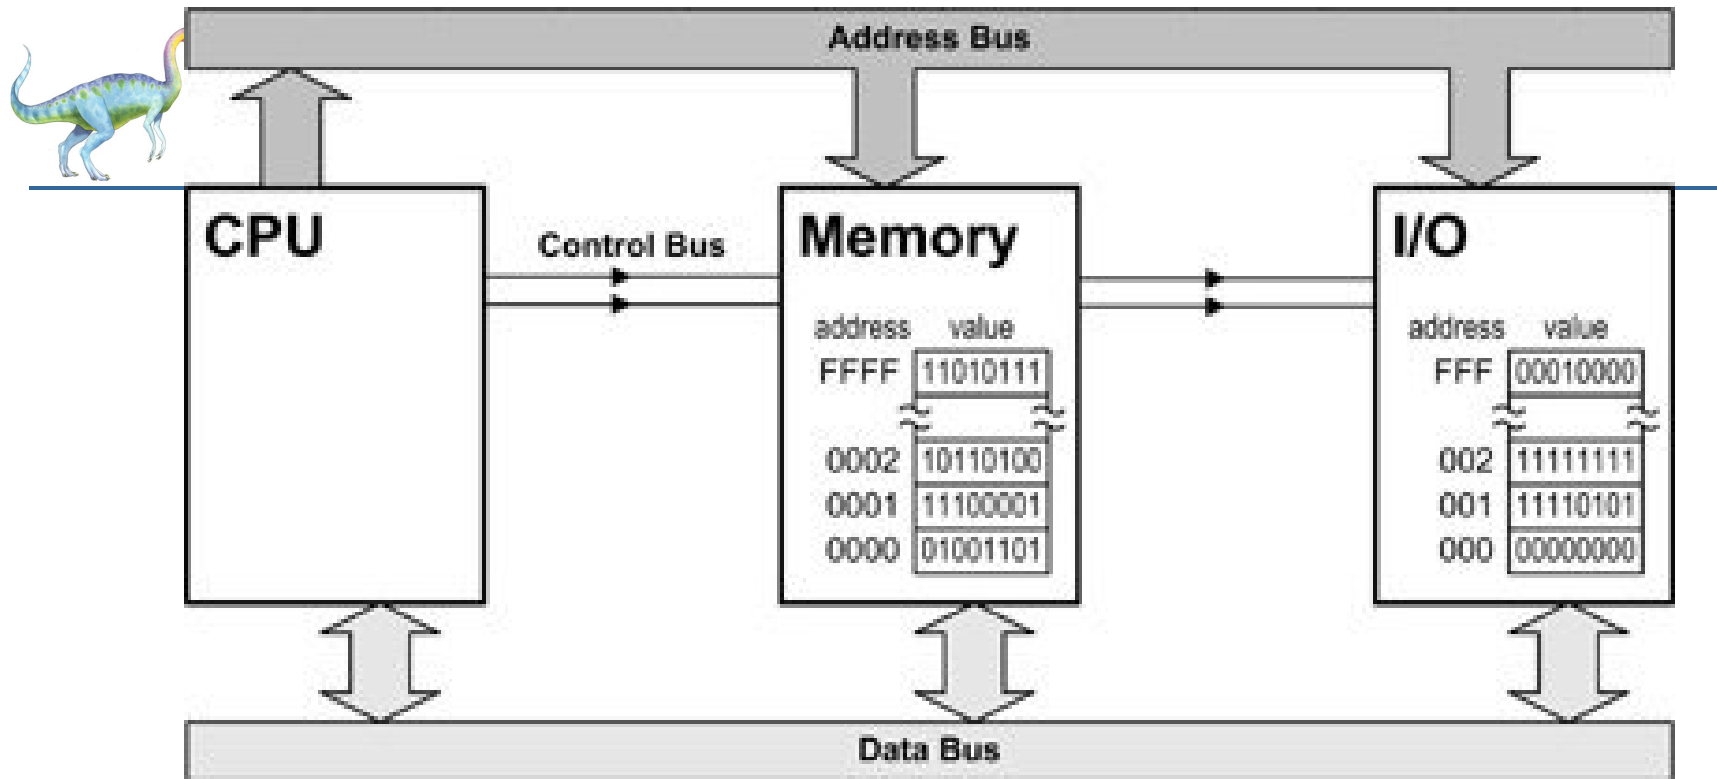
040–043 Timer

3D0–3DF Graphics Controller

Output to port

IN instruction reads from an
I/O device, OUT writes.





- CPU communicates either memory or I/O device at a time, never both.
- CPU uses one of the control lines called "IO/MEM"
 - asserts (line=high) the CPU wants to work to memory
 - (line = low) for an I/O device.
- The IN instruction reads from an I/O device, OUT writes.
 - IO/MEM is not asserted (held low), so memory doesn't respond and the I/O chip does.
- For the memory-oriented instructions (Load/Store),
 - IO/MEM is asserted so CPU talks to the RAM, and IO devices stay out of the communication.





Memory mapped I/O:

1-Busy
0-not

32 bits

IE

B

32 bits Data

@ Physical memory


No Special Instruction load/store instructions

Controller registers maps to address space of the processor.
I/O accomplished with load and store instructions

Example: **load 0xFFFF0000, AL (Control)**
load 0xFFFF0004, BX (Data)

Address Decoding Circuitry

Hybrid : Graphics controller
Control operation via I/O port
Data-out memory mapped region





How does H/W communicate with Kernel/vice versa?

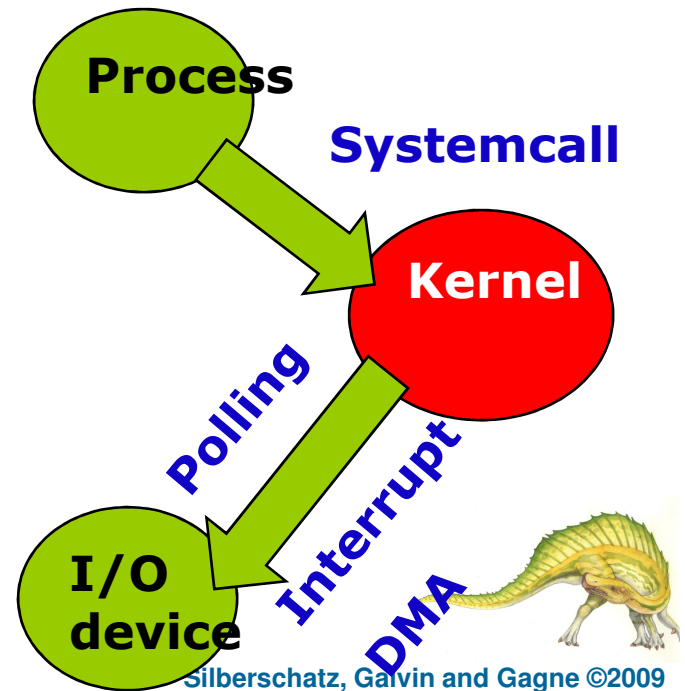
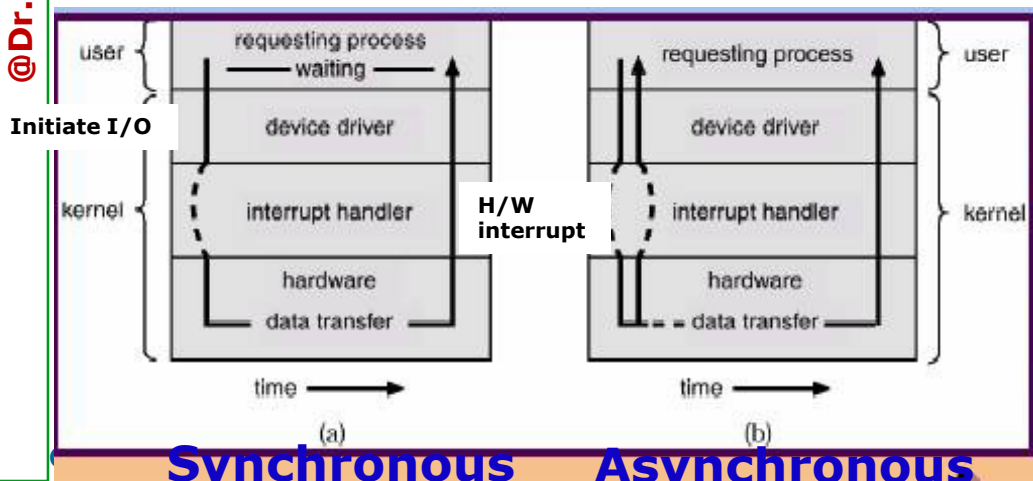
@Dr. Md. Shamim Akhter, CSE, EWU

Process requests I/O

Synchronous

Asynchronous

H/W Interrupt
After I/O completion





Communication Between CPU & H/W Controller


■ Three (3) techniques

Polling, Interrupt and Direct Memory Access

Polling (example : Mobile Phone)

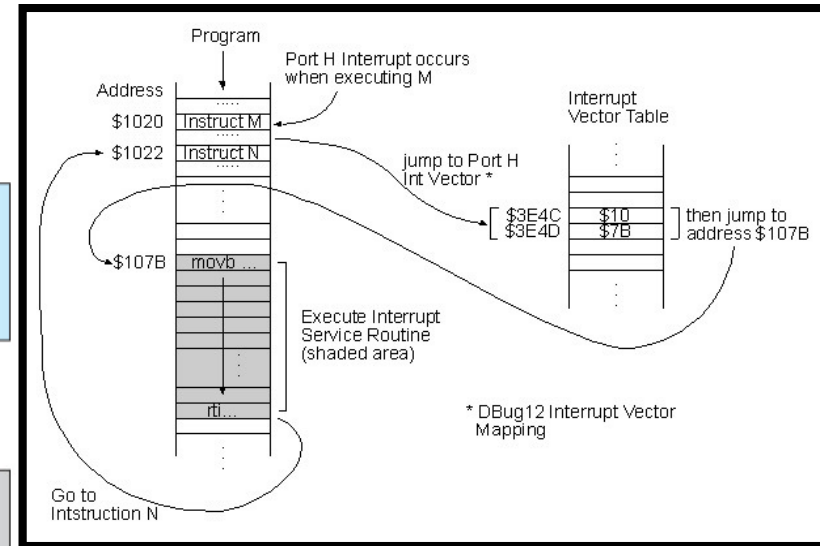
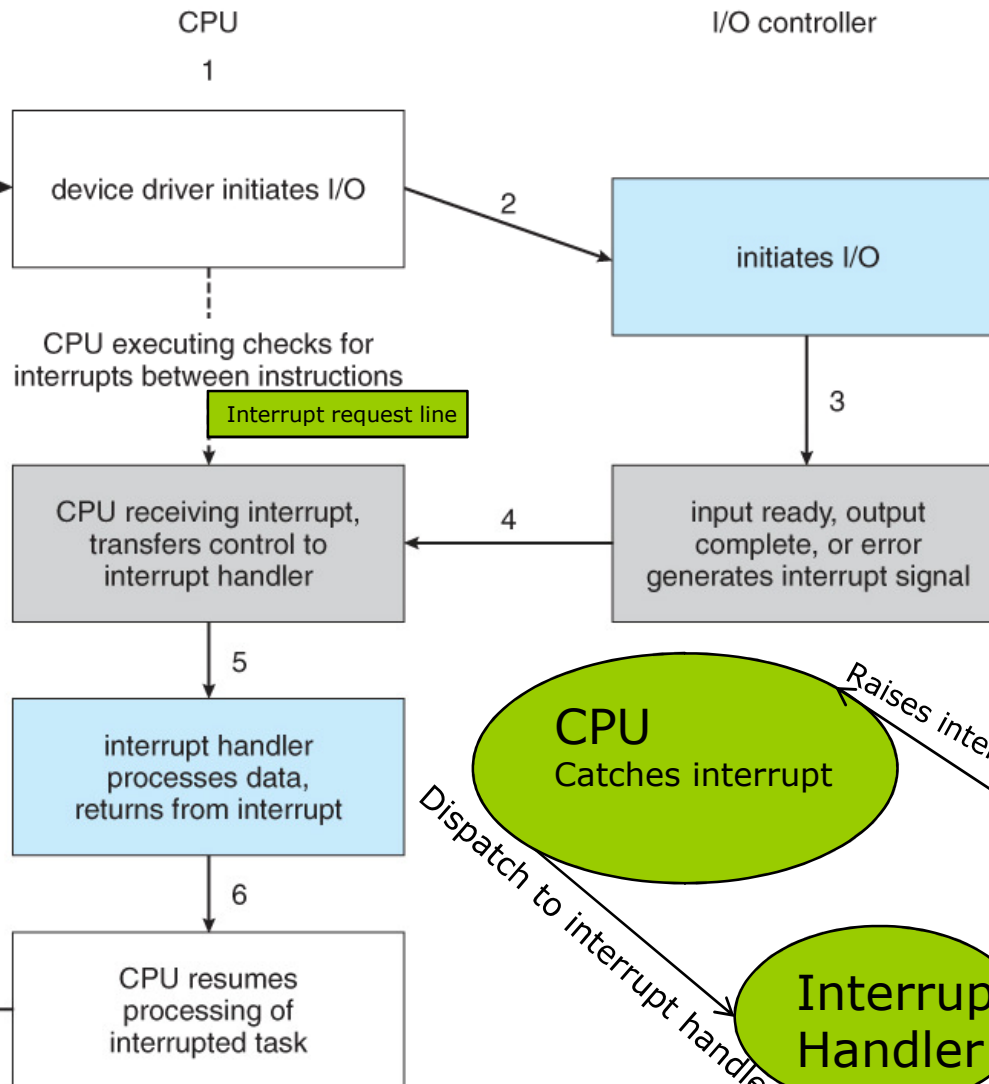
Processor	Controller
Sets the write bit in the command register and writes a byte into the data-out register	
Sets command-ready bit in the command register	Controller notices the command-ready bit is set
BUSY-WAITING	Sets busy bit
	Read command register and see write bit set/write command given
	Gets data from data-out register and does I/O to the device
	Clears the command-ready bit, clears the error bit in the status register to indicate I/O succeed, and clears the busy bit to indicate that it is finished.
Poll a device takes three (3) CPU cycles <ul style="list-style-type: none">-read device register-do logical-and to extract a busy bit @ status register-bit =0 then submit job to device	







Communication with CPU & H/W Controller

Interrupt (example : Mobile Phone)



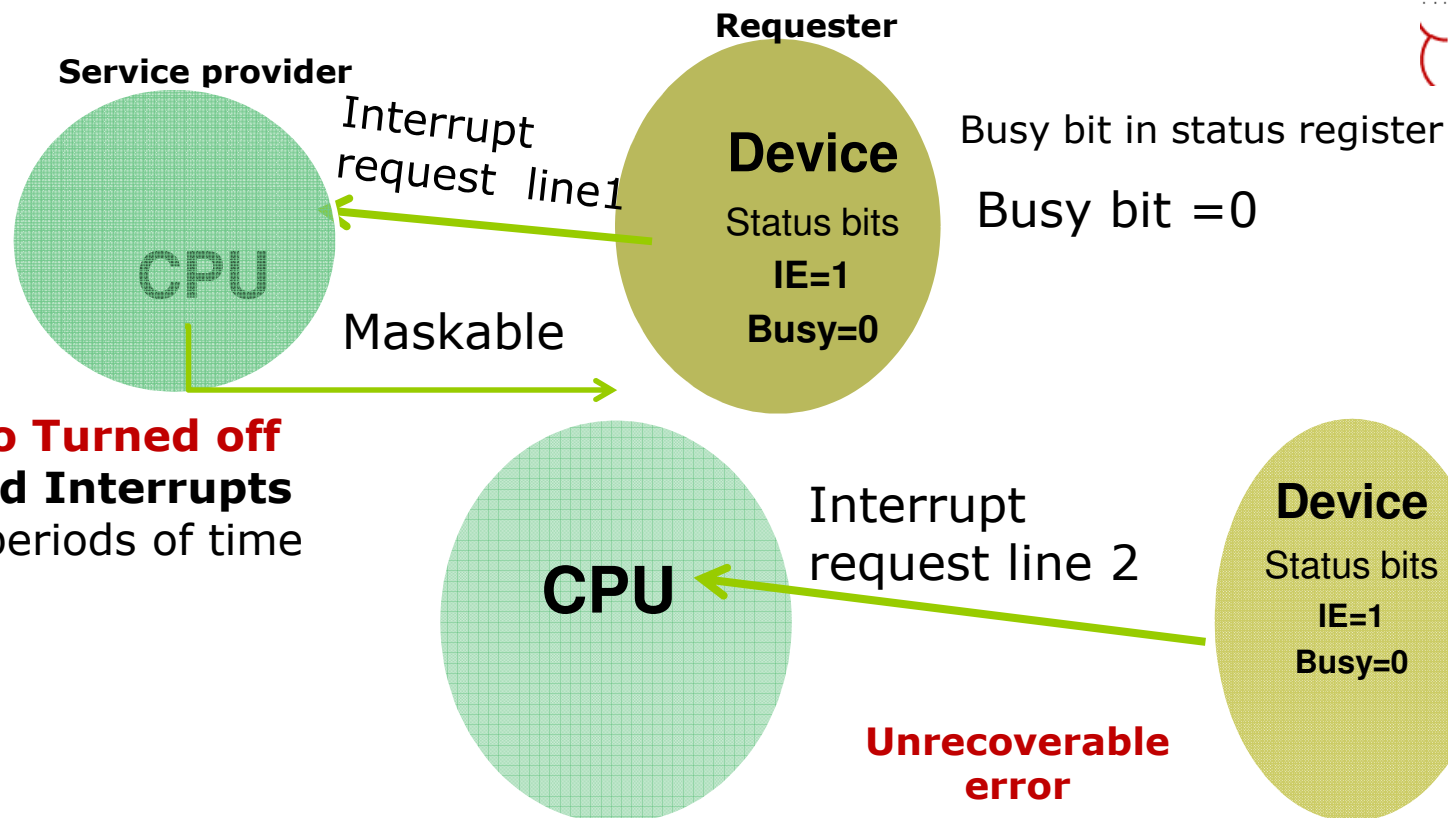
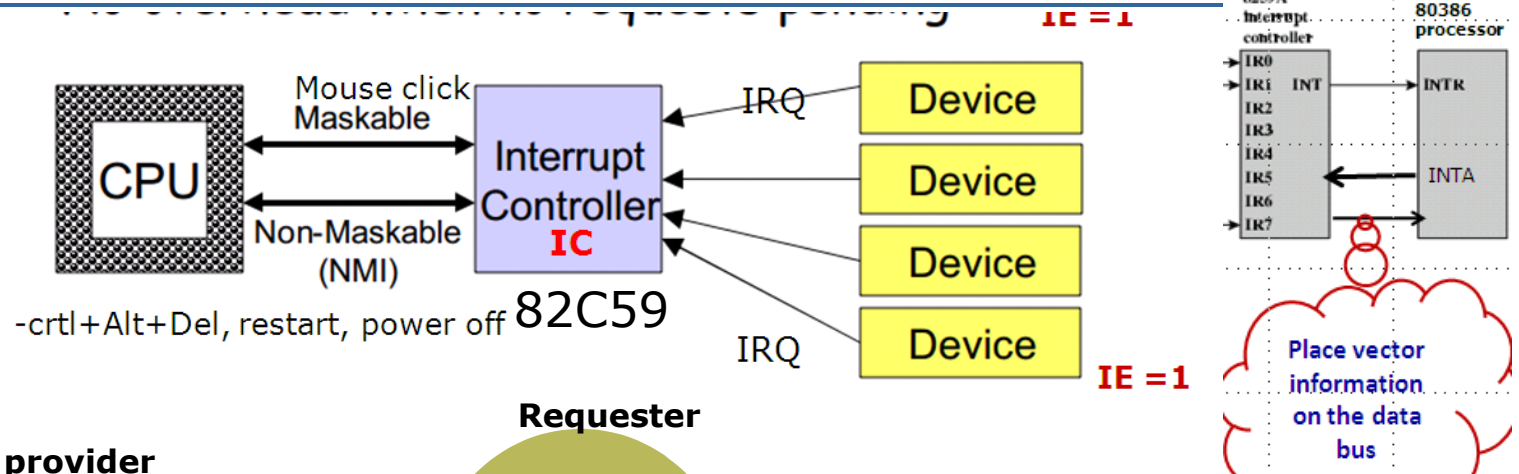
Program Execution with simple Interrupt





Communication with CPU & H/W Controller

Interrupts Work





Interrupt Vector Table

Non-
maskable

vector number	description
0	divide error
1	debug exception
2	null interrupt
3	breakpoint
4	INTO-detected overflow
5	bound range exception
6	invalid opcode
7	device not available
8	double fault
9	coprocessor segment overrun (reserved)
10	invalid task state segment
11	segment not present
12	stack fault
13	general protection
14	page fault
15	(Intel reserved, do not use)
16	floating-point error
17	alignment check
18	machine check
19–31	(Intel reserved, do not use)
32–255	maskable interrupts





DMA – Direct Memory Access

- Device does large transfer – disk drive
 - CPU takes long time (checking status bit, feed data register one byte) Programmed I/O-PPIO
- Without burdening the CPU (only beginning and end)
 - some of this task transfers to special purpose processor Direct Memory Access (DMA) controller

DMA Operation

- CPU wishes to read or write a block of data @memory


- CPU provides command to DMA controller
- Read/Write (IOR, IOW, MEMR, MEMW)
- Device address
- Starting address of memory block of data
- Amount of data to be transferred

DMA Controller

- Status
- Source
- Destination
- Length


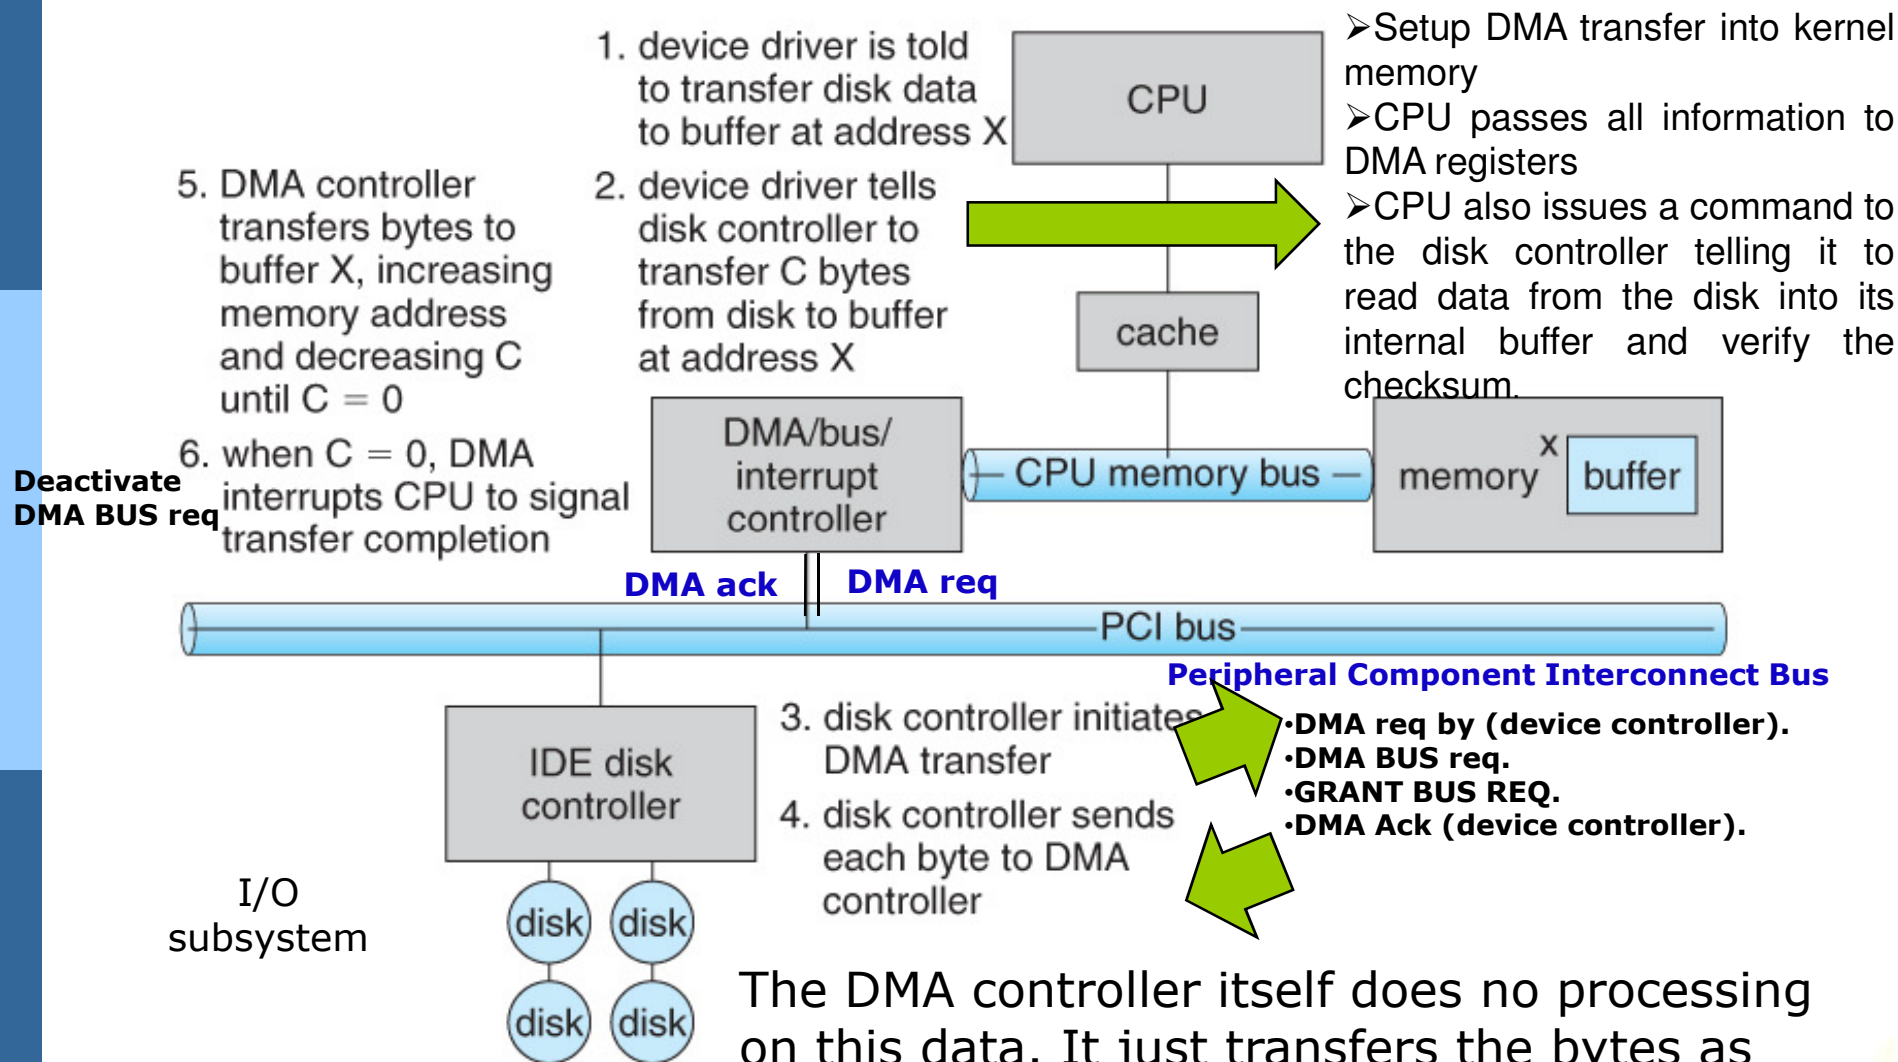
- DMA Takes control over bus

- DMA places **bus request** to CPU
- CPU accepts request by asserting the **bus grant signal** [Both asserts DMA can transfer]
- When data transfer (words) completes, DMA **deasserts request signal** and interrupts to CPU, CPU then deasserts the bus grant signal and takes control over bus.





Handshaking: DMA and Disk Controller





Cyclic Stealing :

- DMA forces CPU to stop its activities and handover the memory bus to DMA. Thus DMA steals a bus cycle from CPU.
- DMA controller transfers one word at a time after which it must return the control of the buses to the CPU.
- The CPU merely delays its operation for one memory cycle to allow the direct memory I/O transfer to “steal” one memory cycle.

DMA is capable to transfer data (Using 4 channels) between:

- Memory and I/O device or vice versa
- Memory to memory using (DMA control/command register loaded by CPU) Channel 0 is from memory to register and channel 1 is from register to memory.





OS Evaluation

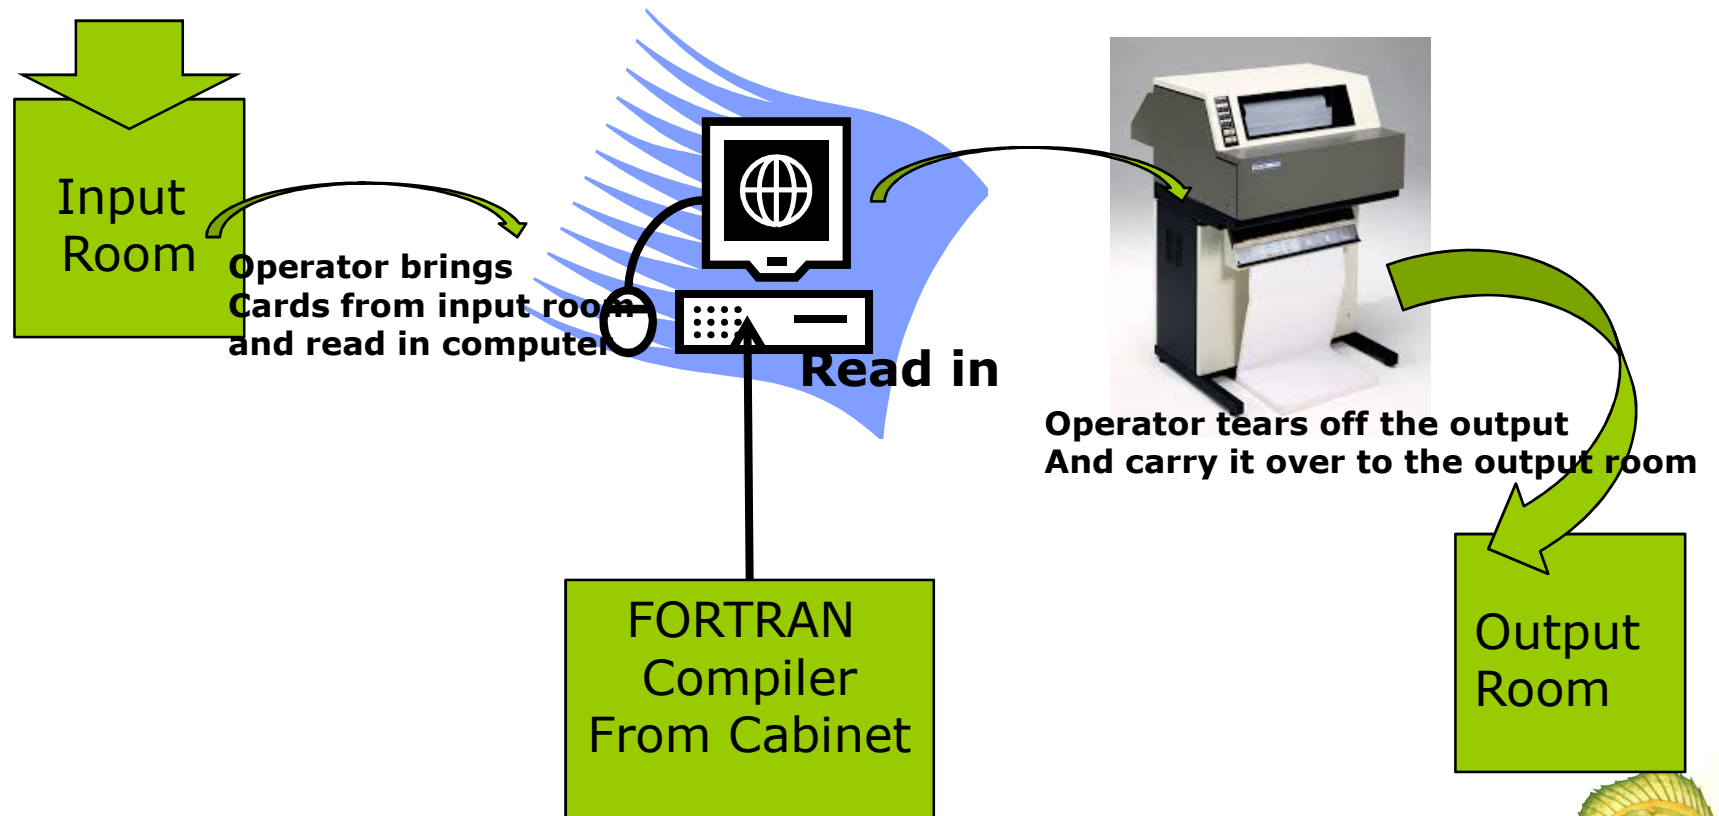




Computers were in Universities / Large companies

Programmers wrote programs (FORTRAN/ Assembly) on paper.

Punched the program into CARDS



Wasting time while operators were walking around the machine room







Supervisor/Operator Control

CR ----> MT MT ----> CPU CPU ----> LP

IBM 7094



An early batch system. (a) Programmers bring cards to 1401. (b) 1401 reads batch of jobs onto tape. (c) Operator carries input tape to 7094. (d) 7094 does computing. (e) Operator carries output tape to 1401. (f) 1401 prints output.

- Problems

- ▶ Long turnaround time - up to 2 DAYS!!!
 - Debugging...
- ▶ Low CPU utilization
 - I/O and CPU could not overlap;
 - slow mechanical devices.



From John Ousterhout slides


31

Offline Processing



BATCH Processing (One by one) Similar jobs

Process 1			Process 2			Process 3			
I1			I2			I3			
	C1			C2			C3		
		O1			O2			O3	
2	4	2	2	3	2	2	7	2	26



I= Input
C=Computation
O=Output

OS resides in memory and the task was simple
- transfer control from one job to another

Problem: Low CPU utilization;
- During I/O CPU is idle. WHY?
- @Class notes






Batch Systems – Upgrading Issues

- **Solutions to speed up I/O:**

- Offline Processing (Previous)

- ▶ load jobs into memory from tapes,
- ▶ card reading and line printing are done offline.



- Online Spooling (*simultaneous peripheral operations on-line*)


- ▶ A type of buffering, to disk (present)
 - Random access device, unlike tape
- ▶ Use disk as large storage for reading as many input files as possible and storing output files until output devices are ready to accept them.
- ▶ Allows overlap - I/O of one job with computation of another.
- ▶ Introduces : job pool (Queues-job queue, ready queue)
 - allows OS to choose next job to run so as to increase CPU utilization





Multiprogramming

- Use interrupts to run multiple programs simultaneously
 - When a program performs I/O,
 - ▶ instead of polling, execute another program till interrupt is received.
 - Efficiency



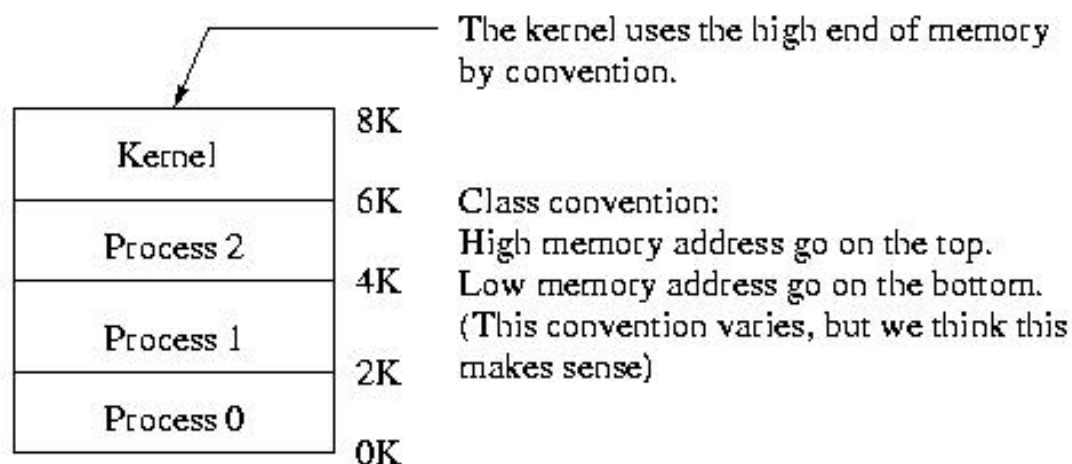
A multiprogramming system with three jobs in memory.

- Requires:
 - **secure memory and** I/O for each program.
 - **intervention, if program loops indefinitely.**
 - CPU scheduling to choose the next job to run.





Simple Model for Memory Protection



To obtain true protection,
need more hardware support.

- The goal is to make sure that, for example, process 2 can't go outside of the 4K-6K region.
 - This requires hardware support -- we can't do it in software alone.
 - The simplest model is to add two registers: **BASE** & **LIMIT** (assigned by OS)
 - This is the simplest model, most computers are more complex.

In order to keep process 2 within it's 2K box, we can do one of two things:

Option #1:

Set BASE = 4K

Set LIMIT=2K

A memory access is legal,

iff $\text{BASE} \leq X < \text{BASE} + \text{LIMIT}$,

where X is the actual memory address

Option #2:

Set BASE = 4K

LIMIT = 2K

more frequent scheme

A memory access is legal,

iff $0 \leq X < \text{LIMIT}$,

where X is the actual memory
address = BASE + X



I/O Protection

- Illegal I/O commands
 - I/O instructions are privileged

CPU Protection

- Timer device, can be set to interrupt the CPU after a specific periods
- Fixed and variable (counter decrement)





Multiprogrammed Processing

I= Input
C=Computation
O=Output

Process 1	I1 2	C1 4	O1 2			
Process 2		I2 2		C2 3	O2 2	
Process 3			I3 2	C3 7	O3 2
	2	4	3	7	2	= 18

8 units
Time efficiency

Problem: Convoy effect!

OS supports

- SPOOL
- Overlapping I/O and CPU [I/O of one job CPU of another job]
- Interleaving – mixed of CPU and I/O operations of several programs and alternating between them





Timesharing

- A variant of multiprogramming.
- Supports multiple users at the same time
- Programs queued for execution in FIFO order.
- A technique to protect CPU.
- **Like multiprogramming, but timer device interrupts** after a quantum (timeslice).
 - ▶ Interrupted program is returned to end of FIFO
 - ▶ Next program is taken from head of FIFO





Time Shared Processing

Process 1	I1 2	C1 2	...	C1 2	O1 2					
Process 2		I2 2	C2 2	C2 1	O2 2			
Process 3			I3 2	...	C3 2	...	C3 2	C3 2	C3 1	O2 2
	2	2	2	2	2	3	2	1	2	=18

Concept of Virtual Processors
Time Slice = 2





Virtualization

■ Time sharing system provides

- Multiprogramming
- **Extend machine** with more convenient interface

■ Virtualization

- Not real but virtual existence
- Two type virtualizations- H/W level abstraction & S/W level abstraction
-

-Virtual H/W- processor,
memory, chipset

-Interface
- Virtual environment
-Service virtualization – cloud
-Application- virtualization

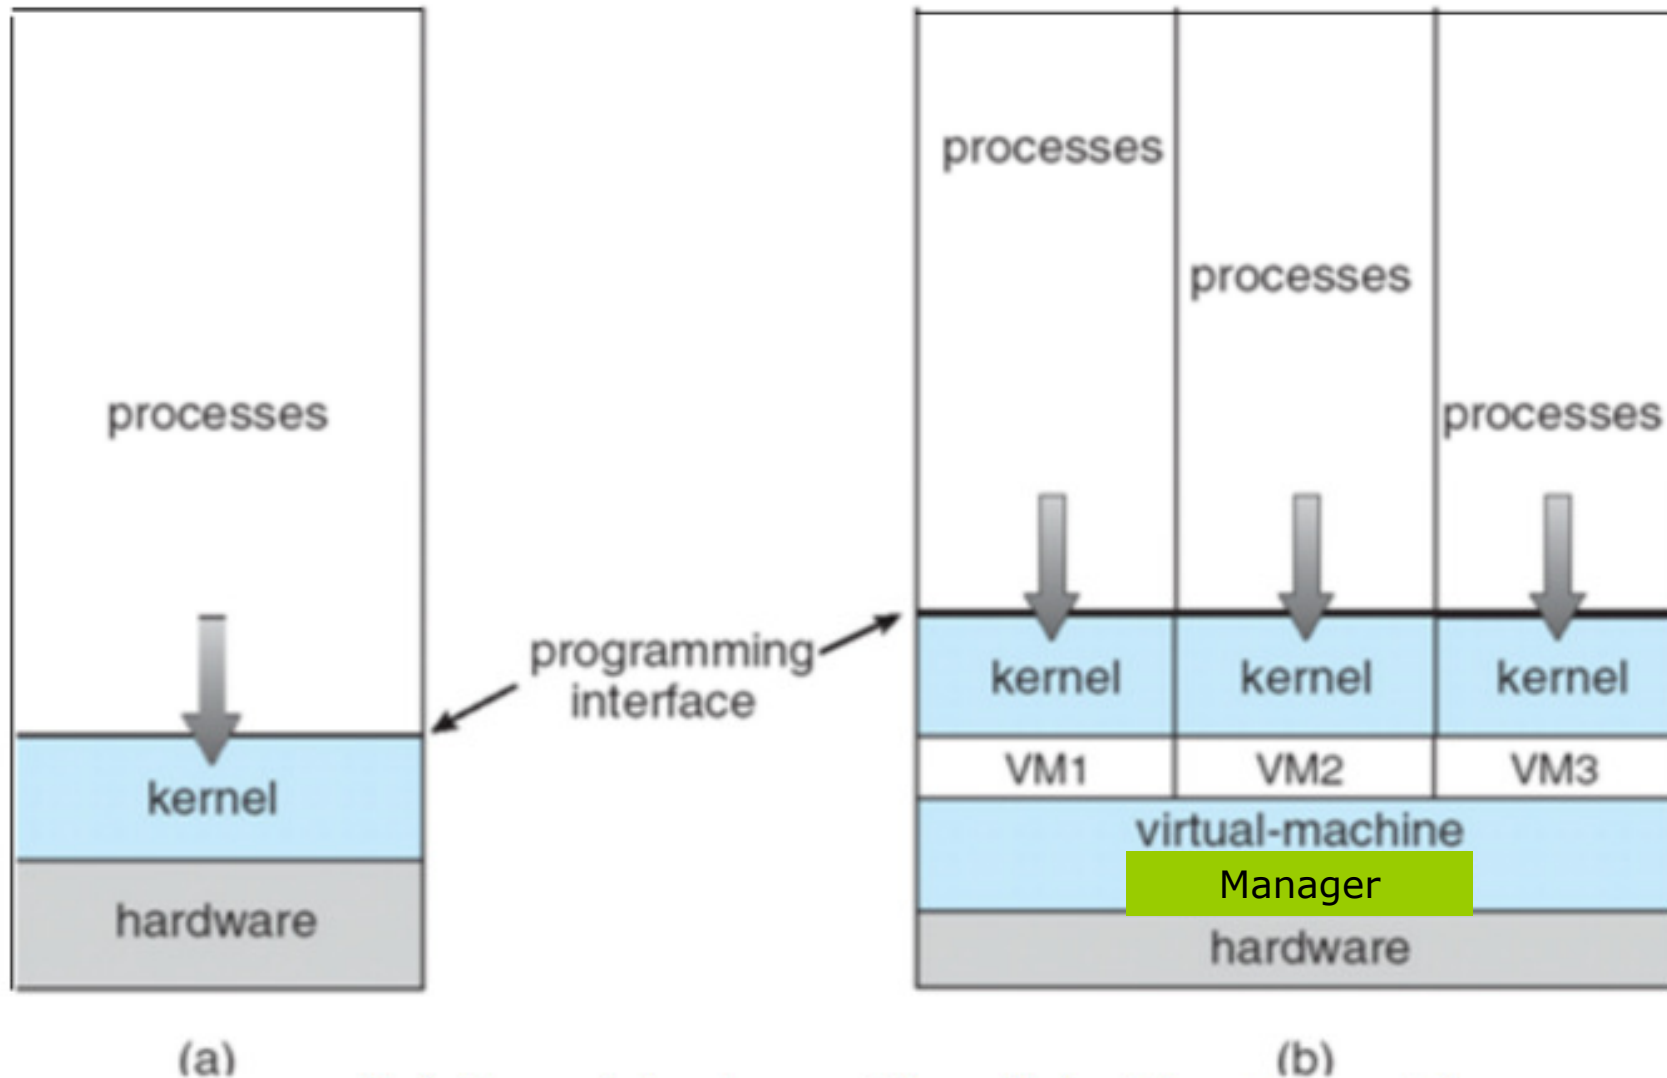
-OS itself running as application to another OS

-Emulation: Running an **operating system** on a hardware platform for which it was not originally engineered.





Virtual Machines

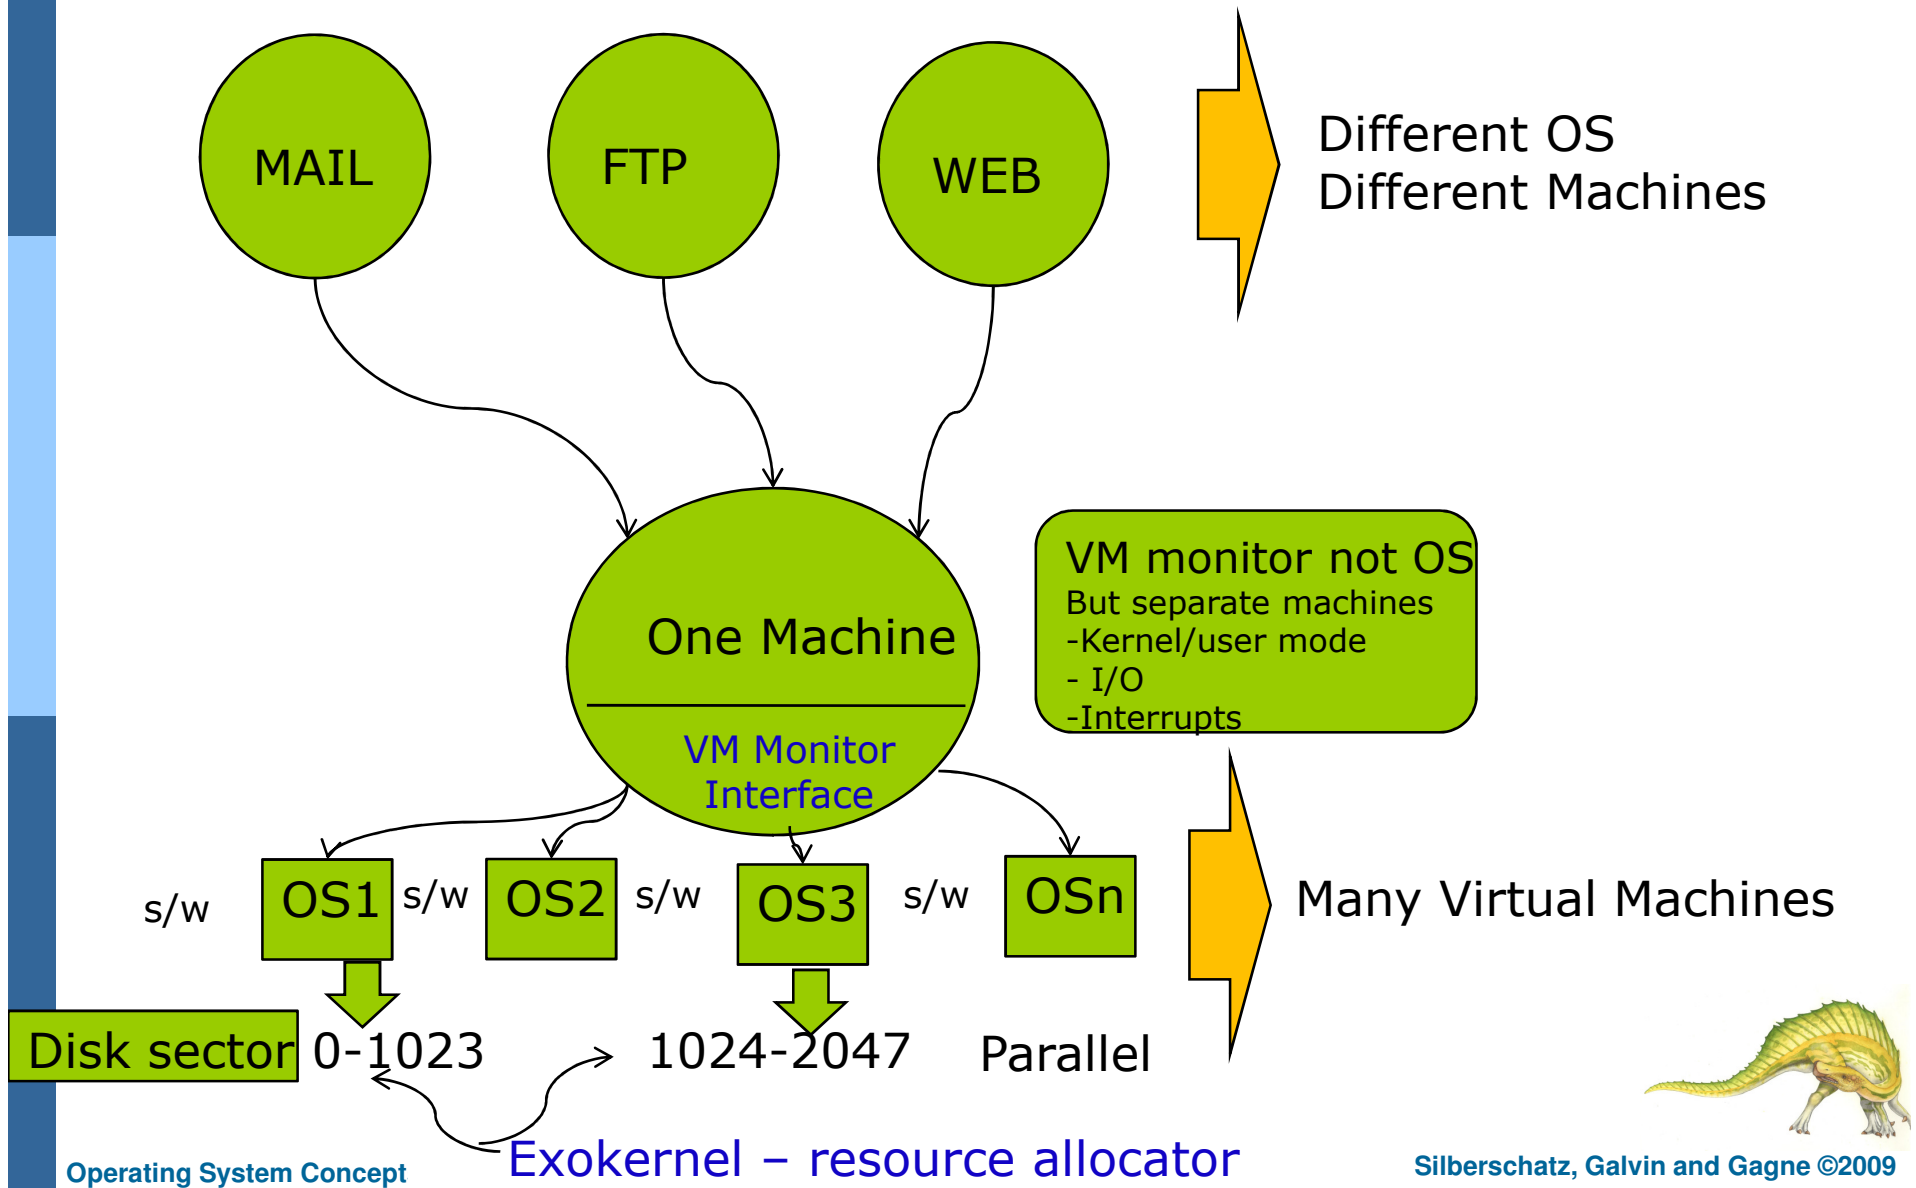


(a) Nonvirtual machine (b) virtual machine





Virtual Machines





Reading Materials

- Galvin : (Introduction) 1.1, 1.2, 1.4, 1.5, 1.11.6
- Galvin: (Operating System Structures) 2.1, 2.3, 2.4.1`, 2.5, 2.6, 2.7.1-2.7.3
- Galvin: (I/O Systems) 13.1-13.2,13.3.4
- Galvin: (Influential OS) 20.2



End of Lecture 1-4

