

The background of the slide is a light gray gradient. It is decorated with numerous water droplets of various sizes, some of which are in sharp focus while others are blurred, creating a sense of depth. The droplets are scattered across the slide, with a higher concentration in the top-left and bottom-right corners.

SCHEDULING

MAN

AUGUST 2015

SCHEDULING

- When more than one process is ready to run, but **only one CPU** is available, a choice is to make
- **Part** of **OS** that does it is **scheduler**
- The algorithm it uses is **scheduling algorithm**

SCHEDULING

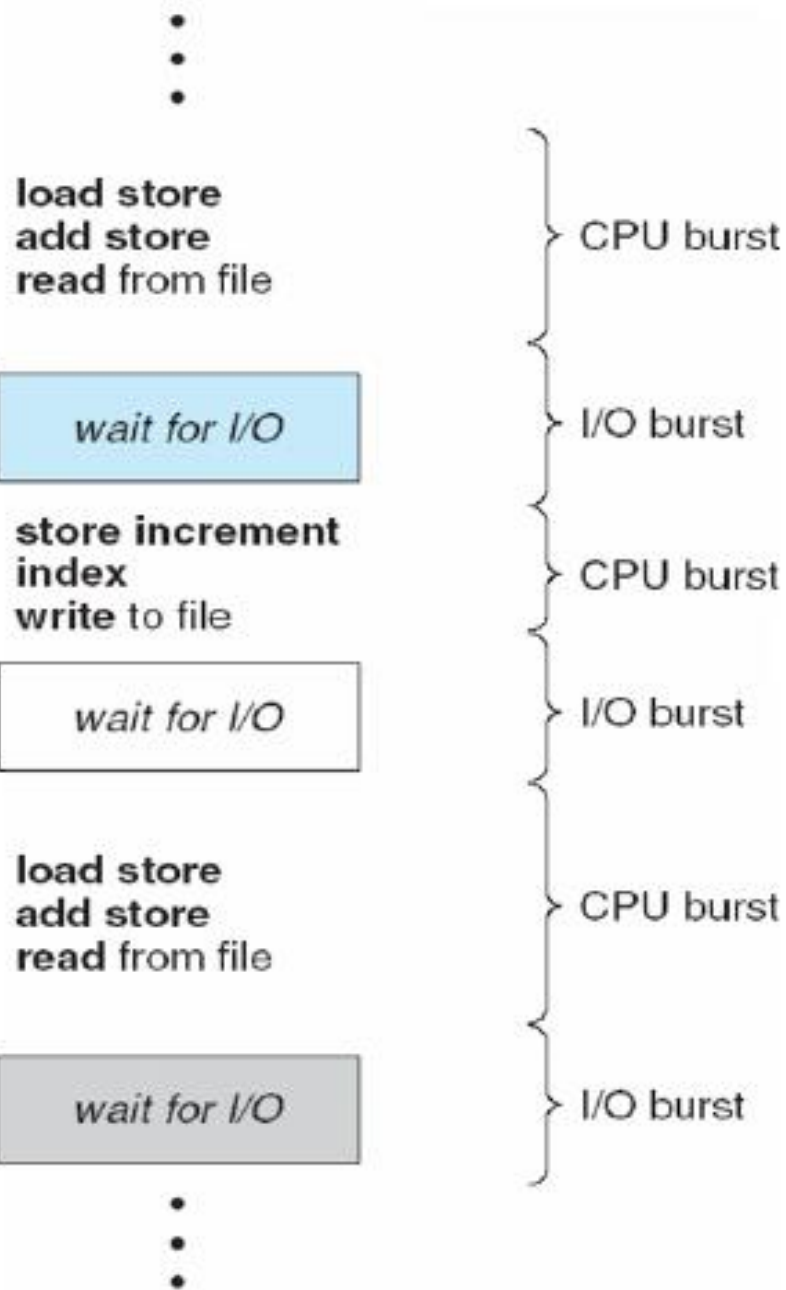
- Efficiency is needed as process switching is **costly**:
 - Switch from user mode to kernel mode
 - State of current process need to be saved
 - Memory map may be saved
 - A process is selected
 - MMU to be reloaded with memory map of new process
 - New process is started

IMPORTANCE OF SCHEDULING

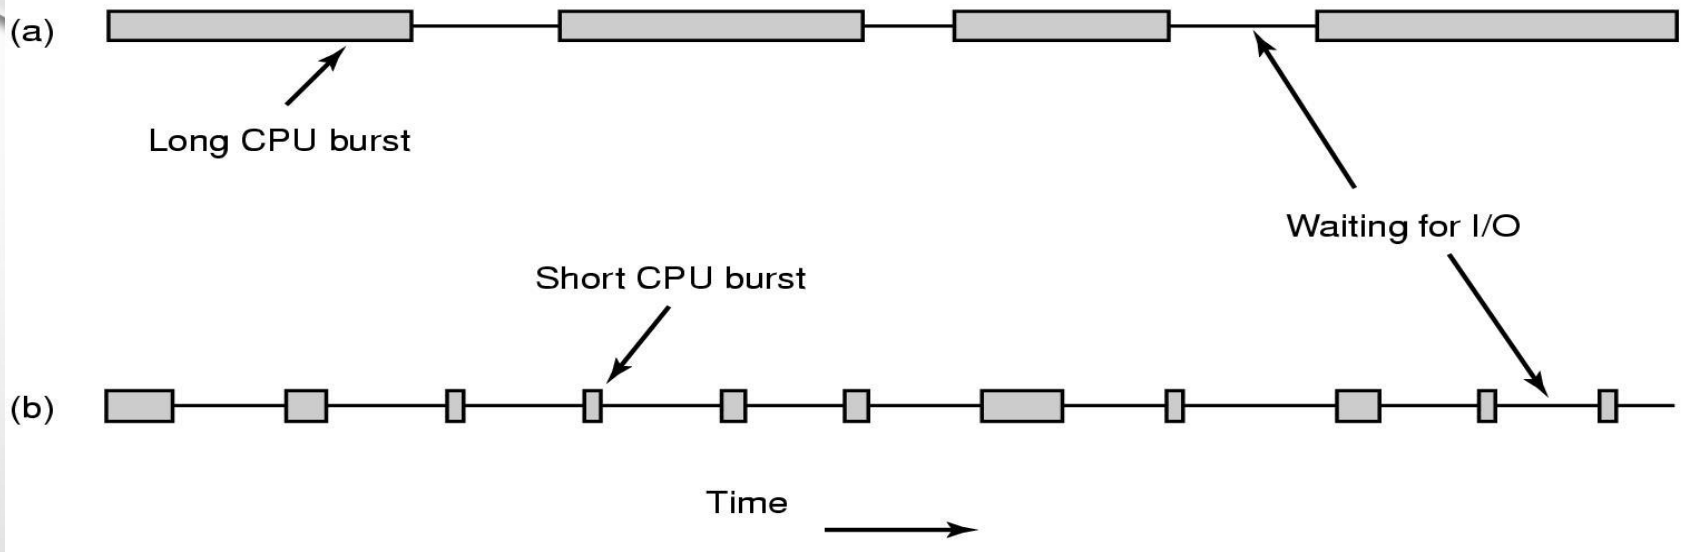
- **Good** scheduling algorithms can make a **big** difference
 - Resource utilization
 - Perceived performance & User satisfaction
 - Meeting other system goals (e.g., important tasks being taken care of immediately)

PROCESS BEHAVIOR

- Processes usually alternate **bursts** of *computing* with *I/O requests*.
- *CPU burst*: the **amount of time** the **process** uses the **processor** before it is no longer ready
- I/O in this sense is when a **process** enters the **blocked** state **waiting** for an external device to complete its work



PROCESS: COMPUTE AND I/O-BOUND



- a CPU-bound process (data encryption/decryption, multimedia encoding)
 - Spend **most** of the time **computing**
 - **Long** CPU bursts => infrequent I/O waits
- an I/O bound process (shell waiting for user commands)
 - Spend most of the time waiting for I/O
 - **Short** CPU bursts => frequent I/O waits
- Key factor is the **length of CPU burst** not the length of the I/O burst

PROCESS: COMPUTE AND I/O-BOUND

- As the CPUs get faster, processes tend to get more I/O bound: **WHY?**
- If a I/O bound process is ready, it should get a chance quickly.
 - Increase resource utilization

WHEN TO SCHEDULE

- When a new process is created:
 - Parent or child? Both are Ready
 - which one to run?
- When a process exits:
 - One of the ready processes should be run
- When a process blocks: Another process has to be selected to run
 - Blocking may occur for:
 - I/O
 - Semaphore

WHEN TO SCHEDULE

- When an I/O interrupt occurs:
 - In case of an interrupt of an I/O device having **completed** its work, some blocked process may now be ready
- If a h/w clock provides **periodic** interrupt: A scheduling decision can be made at each (or kth) clock interrupt

PREEMPTIVE & NON-PREEMPTIVE

Classification of **Scheduling Algorithm** depending on dealing with clock interrupt

- **Non-preemptive**: Picks a process to run and lets it run until it **blocks** or voluntarily releases the CPU. **In effect at each clock interrupt, no scheduling is done.**
- **Preemptive**: Picks a process and lets it run for a maximum of some fixed time. If still running, it is **suspended** and another is picked.
- Preemptive scheduling requires having a **clock interrupt** occur at the end of the time interval to give **control** of the CPU back to the **scheduler**

DIFFERENT SYSTEMS, DIFFERENT FOCUSES

All systems

- Fairness - giving each process a fair share of the CPU
- Policy enforcement - seeing that stated policy is carried out
- Balance - keeping all parts of the system busy

Batch systems

- Throughput - maximize jobs per hour
- Turnaround time - minimize time between submission and termination
- CPU utilization - keep the CPU busy all the time

Interactive systems

- Response time - respond to requests quickly
- Proportionality - meet users' expectations



Real-time systems

- Meeting deadlines - avoid losing data
- Predictability - avoid quality degradation in multimedia systems

BATCH SYSTEMS

- Users **submit** their job to the batch system
- Batch system starts user job when appropriate
- User gets notification that job is **done**
 - No interaction **in between**
- No users impatiently waiting at terminals for a **quick** response to a **short** request
- Used in business world such as Profit calculation at banks, claims processing at insurance companies...

BATCH SYSTEMS

- Common performance metrics
 - Throughput: number of jobs **completed** per hour 
 - Turnaround time: average time between the **submission** and **completion** of a job 
- Maximizing Throughput may not necessarily minimize Turnaround time

BATCH SYSTEMS

Algorithms used:

- Non-preemptive
- Preemptive algorithms with long time periods are often acceptable
 - Reduces process switches and improves performance

Representative algorithms:

1. First Come First Serve (FCFS)
2. Shortest Job First
3. Shortest Remaining Time First

FIRST COME FIRST SERVE (FCFS)

- Process that requests the CPU FIRST is allocated the CPU FIRST.
- Also called FIFO
- **non**-preemptive
- Used in Batch Systems
- Real life analogy?
 - Transaction at Sonali Bank
- Implementation
 - FIFO queues
 - A **new** process enters the **tail** of the queue
 - The **schedule** selects from the **head** of the queue.

FCFS EXAMPLE

Process	Duration	Order	Arrival Time
P1	24	1	0
P2	3	2	0
P3	4	3	0

The final schedule:



P1 turnaround: 24

P2 turnaround: 27

P3 turnaround: 31


The average turnaround:

$$(24+27+31)/3 = 27.33$$

FCFS EXAMPLE 2

Process	Duration	Order	Arrival Time
P1	24	3	0
P2	3	1	0
P3	4	2	0

The final schedule:



P1 turnaround: 31

P2 turnaround: 3

P3 turnaround: 7

The average turnaround:
 $(31+3+7)/3 = 13.67$

ADVANTAGE

- Easy to understand and implement
- Fair for equivalent processes

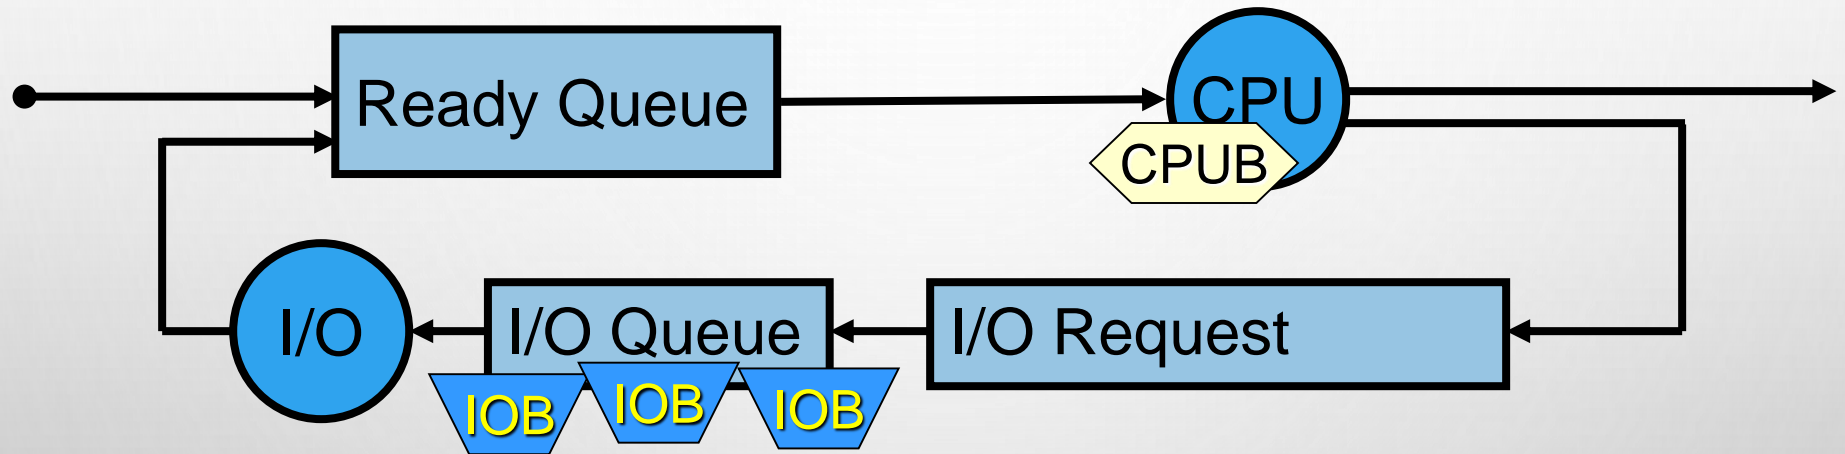
PROBLEMS WITH FCFS

- Non-preemptive
- Non optimal turnaround
- **Cannot utilize** resources in **parallel**:
 - Assume **1** process CPU bounded and **many** I/O bounded processes
 - result: **Convoy effect**,
 - **low** CPU **and** I/O Device utilization
 - Why?

CONVOY EFFECT

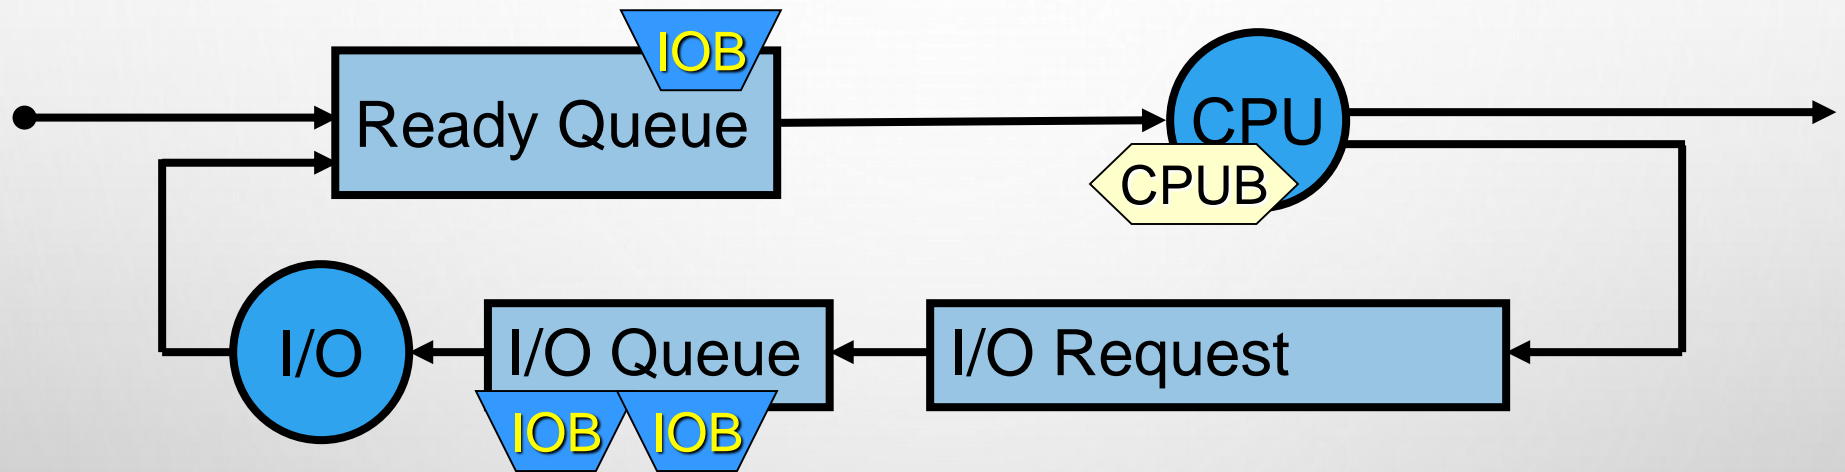
- When the CBP **uses** the CPU
 - IBPs **finish** their I/O and move into the ready queue, **waiting** for the CPU
 - the **I/O** devices are **idle**
- When the CBP finally relinquishes the CPU,
 - CBP moves to an I/O device
 - the IBPs pass through the CPU **quickly** and move back to the I/O queues
 - the CPU is **idle**
- The cycle **repeats** itself when the CBP gets back to the ready queue

CONVOY EFFECT



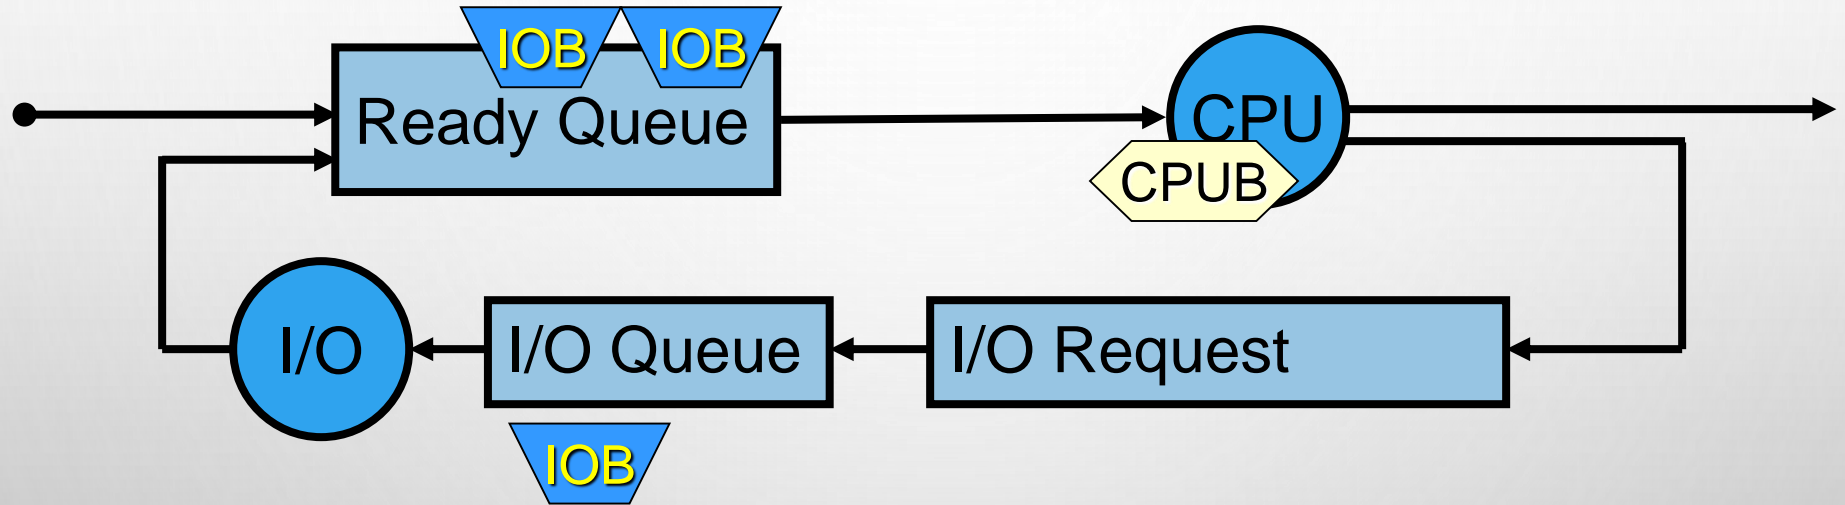
CPU is running CPUB

CONVOY EFFECT



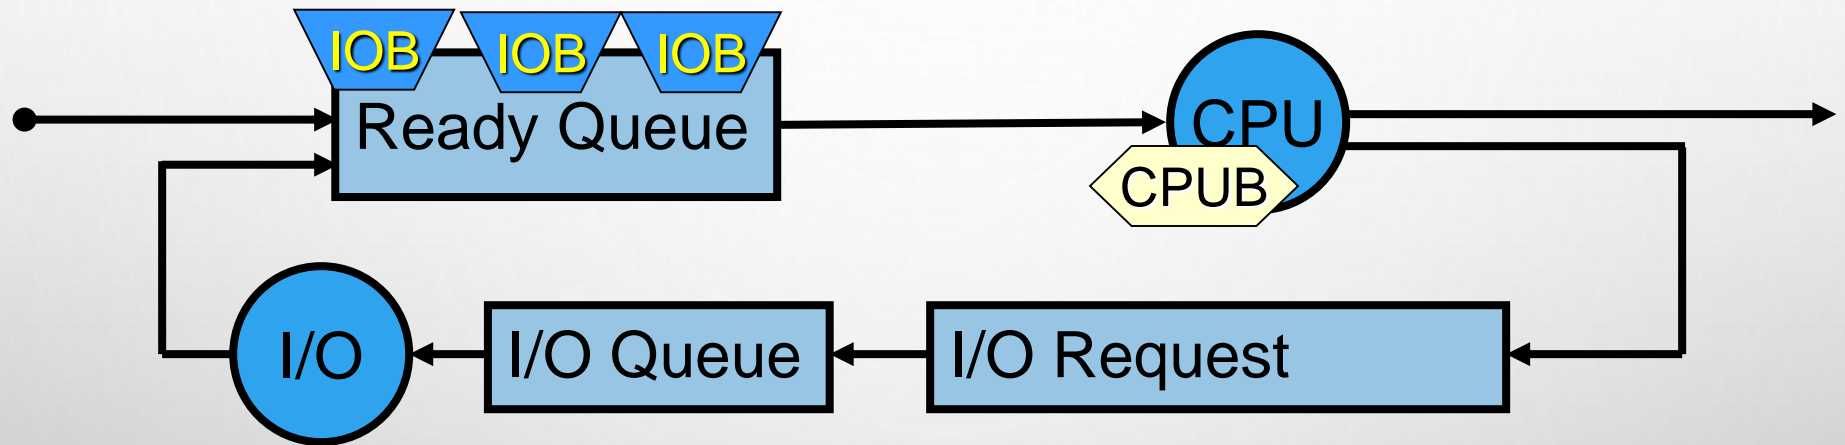
CPU is running CPUB

CONVOY EFFECT



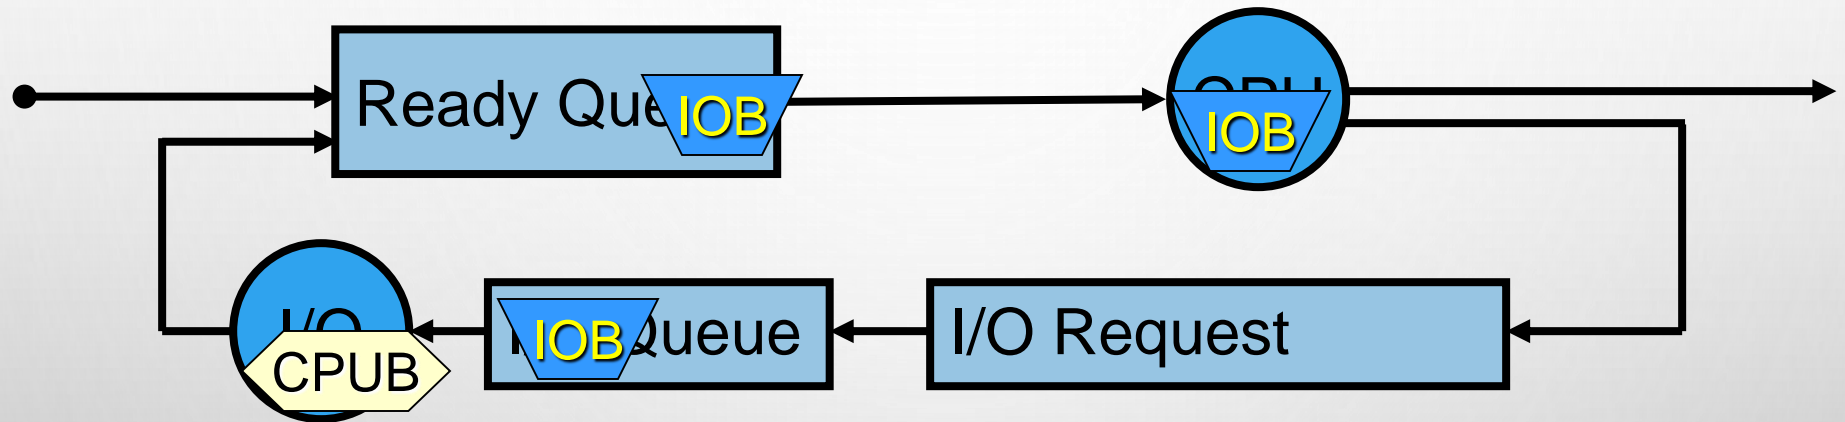
CPU is running CPUB

CONVOY EFFECT



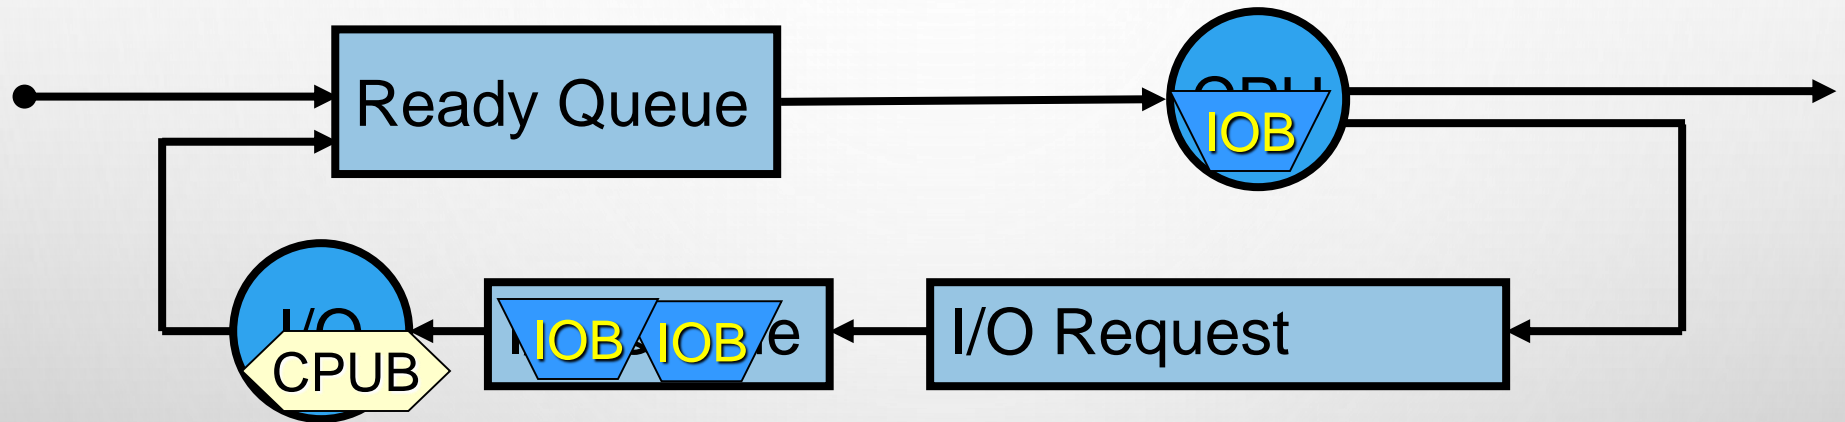
*CPU is running CPUB
I/O devices idle*

CONVOY EFFECT



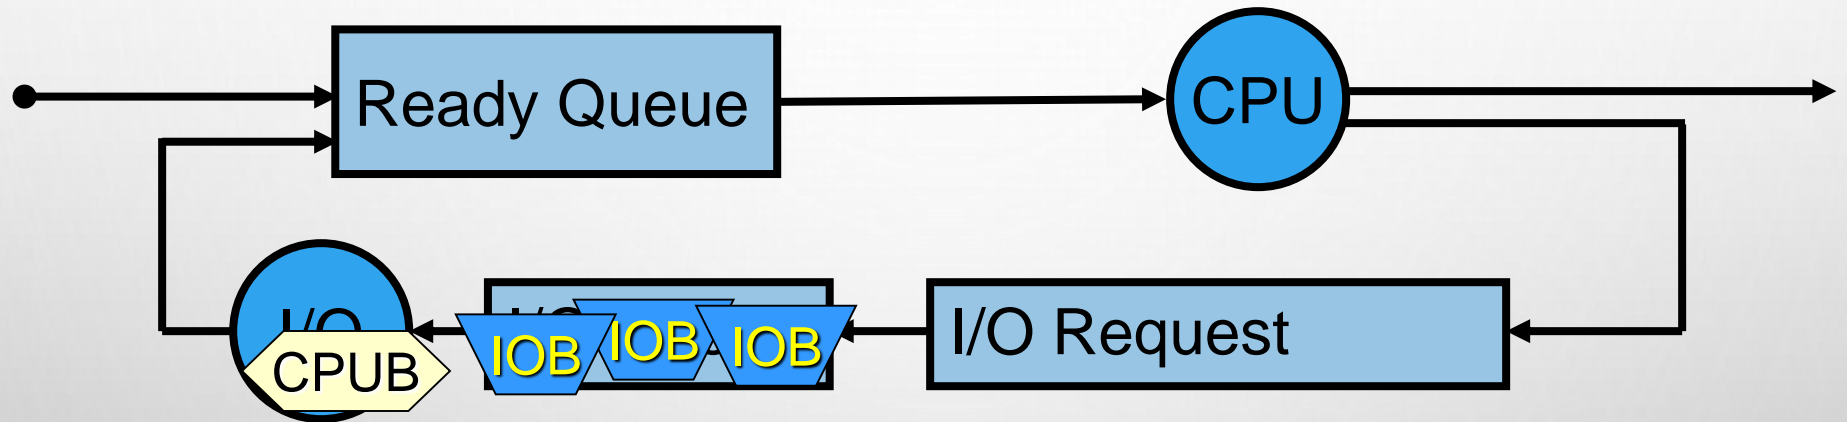
CPU moves to I/O device

CONVOY EFFECT



I/O Bound jobs take very small amount of CPU time and go for I/O

CONVOY EFFECT




CPU idle

SHORTEST JOB FIRST (SJF)

- Scheduling algorithm in **batch** systems
- Schedule the job with the shortest run time first
- Requirement: **the run time needs to be known in advance**
- SJF is **optimal** in terms of turnaround, if **all** jobs arrive at **same time**

SJF: EXAMPLE

Process	Duration	Order	Arrival Time
P1	6	1	0
P2	8	2	0
P3	7	3	0
P4	3	4	0



Do it yourself

P4 turnaround: 3
P1 turnaround: 9
P3 turnaround: 16
P2 turnaround: 24

Total execution time: 24
The average turnaround:
 $(3+9+16+24)/4 = 13$

COMPARING TO FCFS

Process	Duration	Order	Arrival Time
P1	6	1	0
P2	8	2	0
P3	7	3	0
P4	3	4	0



P1 turnaround: 6
P2 turnaround: 14
P3 turnaround: 21
P4 turnaround: 24

The total time is the same.
The average turnaround:
 $(6+14+21+24)/4 = 16.25$
(comparing to 13)

SJF IS NOT ALWAYS OPTIMAL

- SJF OPTIMAL ONLY IF ALL JOBS HAVE ARRIVED AT SCHEDULING TIME

Process	Duration	Order	Arrival Time
P1	10	1	0
P2	2	2	2



P1 turnaround: 10

P2 turnaround: 10


The average turnaround (AWT):
 $(10+10)/2 = 10$

PREEMPTIVE SJF

- Also called **Shortest Remaining Time Next**
 - Schedule the job with the shortest **remaining** time required to complete
 - When **new** job **arrives**, compare its **total** time with the **remaining** time of the running job
 - If the new job needs less time the current job is suspended and the new job started
- Requirement: the run time needs to be known in advance

PREEMPTIVE SJF: SAME EXAMPLE

Process	Duration	Order	Arrival Time
P1	10	1	0
P2	2	2	2



P1 turnaround: 12
P2 turnaround: 2

The average turnaround:
 $(2+12)/2 = 7$

PROBLEM WITH PREEMPTIVE SJF?

- Starvation
 - In some condition, a job is waiting for ever
 - Example: Preemptive SJF
 - Process A with run time of 1 hour arrives at time 0
 - But every 1 minute, a short process with run time of 1 minute arrives
 - Result of Preemptive SJF: A never gets to run

INTERACTIVE SYSTEM

- Example: Servers
 - Serve multiple remote users all of whom are in a big hurry
- Performance Criteria
 - Min response time:
 - amount of time it takes from when a request was submitted until the **first response** is produced, not output
 - respond to requests quickly

INTERACTIVE SYSTEM

- Algorithms used here usually preemptive
 - Time is **sliced** into quantum (time intervals)
 - Scheduling decision is also made at the beginning of each quantum
- Representative algorithms:
 - Round-robin
 - Priority-based
 - Shortest process time
 - Guaranteed Scheduling
 - Lottery Scheduling
 - Fair Sharing Scheduling


ROUND ROBIN

- **Round Robin (RR)**

- Often used for timesharing
- Each process is given a time slice called a *quantum*
- It is run for the quantum or until it blocks
- RR allocates the CPU uniformly (fairly) across participants from ready queue.

- Problem:

- Do not consider priority
- Context switch overhead




IMPLEMENTING ROUND ROBIN

- Keep the ready queue as a FIFO queue of processes.
- **New** processes are added to the **tail** of the ready queue.
- The scheduler
 - picks the **first** process from the ready queue
 - sets a timer to interrupt after 1 time quantum, and
 - Starts the process.
- When the quantum is over
 - The running process will be put at the **tail** of the ready queue.

RR WITH TIME QUANTUM = 20

<u>Process</u>	<u>Run Time</u>
P_1	53
P_2	17
P_3	68
P_4	24

- All processes arrive at time 0
- The **Gantt** chart is



- Higher average turnaround than SJF
- But better response time

RR: CHOICE OF TIME QUANTUM

- Performance depends on length of the timeslice
 - Context switching isn't a free operation.
 - If timeslice time is set too high
 - attempting to amortize context switch cost, you get FCFS.
 - i.e. processes will finish or block before their slice is up anyway
 - Poor response time
 - If it's set too low
 - you're spending all of your time context switching between threads.

PRIORITY SCHEDULING

- Each job is assigned a priority
- Select **highest** priority job to run next
- Rational: higher priority jobs are more important
 - Example: simulation vs. auto save a document
- Problems:
 - Low priority process may **starve**
- Solution:
 - Priority need to be **adjusted** depending on the situation

ASSIGN PRIORITY

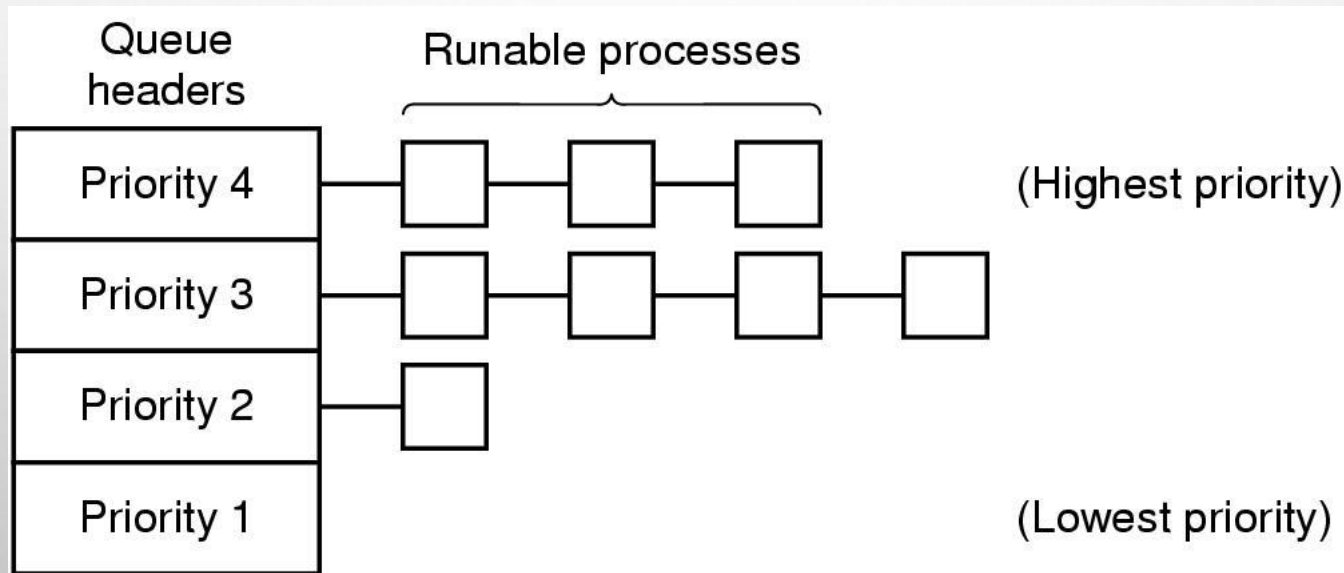
- Two approaches
 - Static (for system with well known and regular application behaviors)
 - Dynamic (otherwise)
- Priority may be based on:
 - Cost to user.
 - Importance of user
 - Percentage of CPU time used in last X hours

EXAMPLE: DYNAMIC PRIORITY ASSIGNMENT

- Whenever highly I/O bound processes wants the CPU it should be given the CPU immediately.
- Why?
- A simple algorithm for giving priority to I/O bound processes is to set the priority to $1/f$
 - f is the fraction of the last quantum used by a process
 - A process that used only 1 msec of its 50 msec quantum would get priority 50
 - A process that used 25 msec of its 50 msec quantum would get priority 2

PRIORITY CLASS

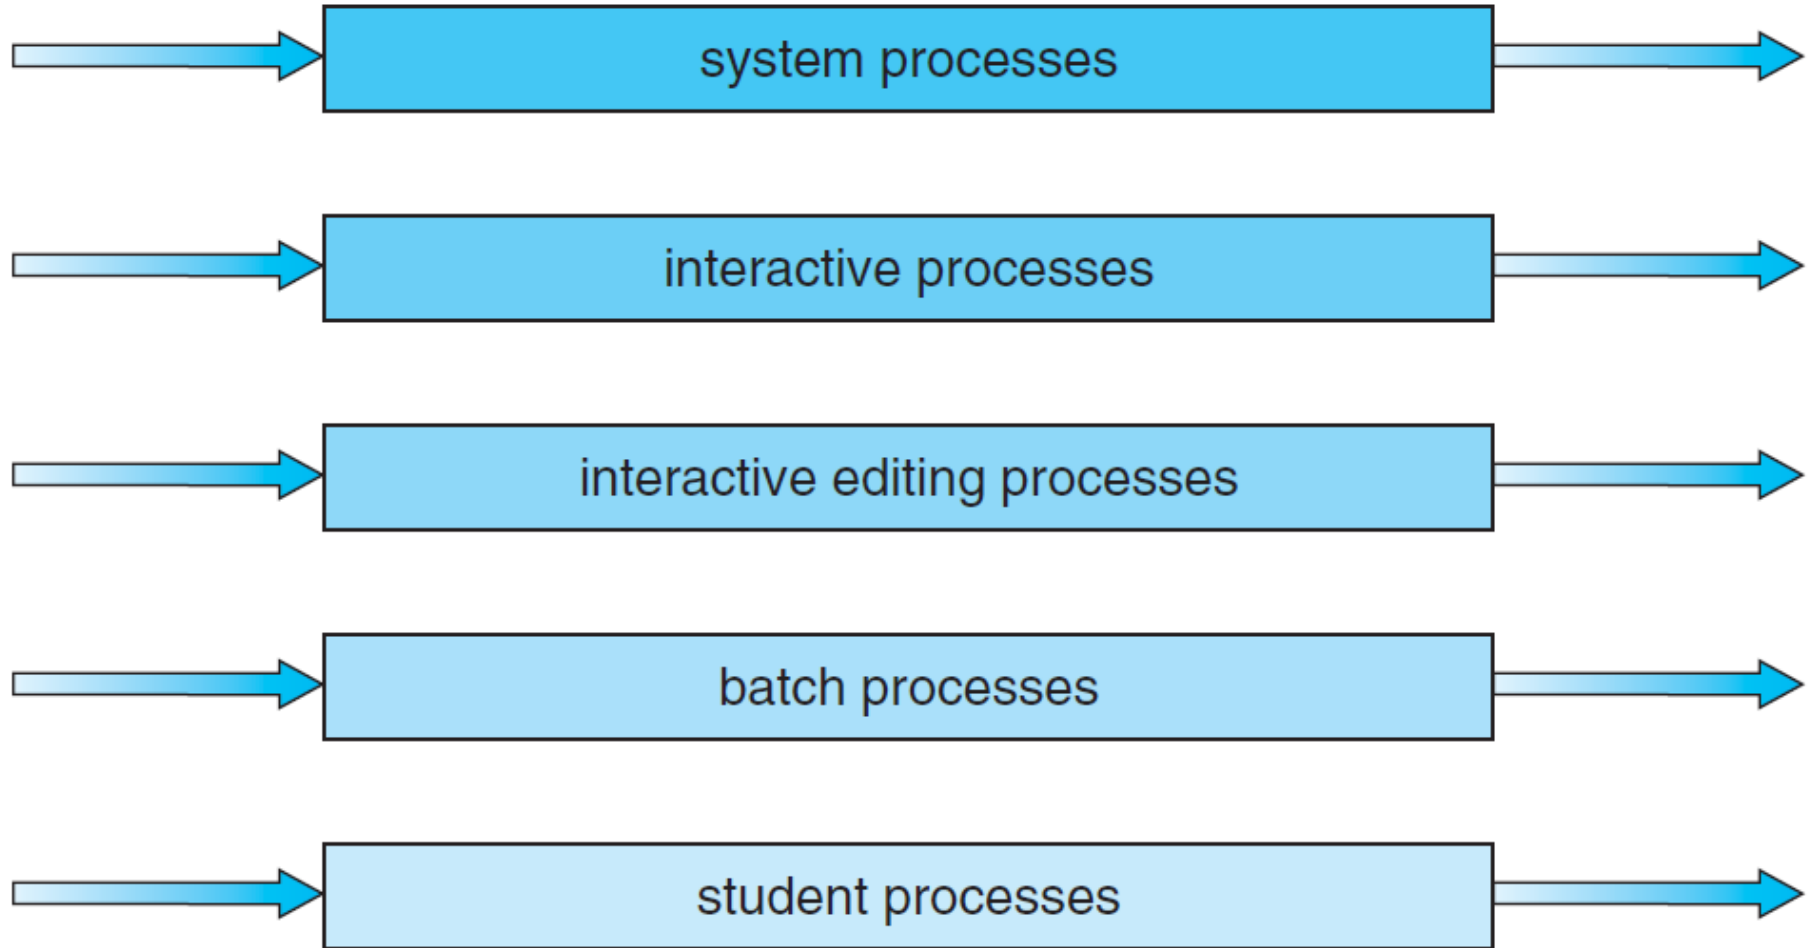
- It is often convenient to **group** processes into priority classes and use priority scheduling among the classes but RR within each class



- If priorities are not adjusted occasionally, lower priority classes may all **starve** to death

PRIORITY CLASS

highest priority



lowest priority

SHORTEST PROCESS NEXT


- Let's apply SJF for **interactive** processes
 - General pattern of a **interactive** process: CPU burst, I/O burst, ...
 - Let's regard the execution of each CPU burst as a separate "**job**"
 - Now we can minimize overall response time by running the process with shortest "**job**" first

SHORTEST PROCESS NEXT

- How to know the length of the next CPU burst?
- A possible answer: Exponential averaging
- Make **estimate** based on past behavior and run the process with the shortest estimated CPU burst
 - Let the **current estimated** CPU burst is τ_n
 - length of the nth CPU burst t_n
 - predicted value for the next CPU burst $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$


$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \cdots + (1 - \alpha)^j \alpha t_{n-j} + \cdots + (1 - \alpha)^{n+1} \tau_0.$$

EXPONENTIAL AVERAGING



CPU burst (t_i)	6	4	6	4	13	13	13	...	
"guess" (τ_i)	10	8	6	6	5	9	11	12	...

EXPONENTIAL AVERAGING



GUARANTEED SCHEDULING

- Make promises to users about performance & then meet those promises
- With n processes running, each one should get $1/n$ of the CPU cycles
- Calculate **ratio** for each process
 - $$\frac{\text{Amount of CPU time process has had since its creation}}{\text{Amount of CPU time process **should have** since creation}}$$
- Run the process with the **lowest** ratio until its ratio has moved above its closest competitor
- Problem:
 - Implementation is **difficult**

The background of the slide is a light gray gradient. In the top-left and bottom-right corners, there are several realistic water droplets of various sizes, rendered with soft shadows and highlights to give them a three-dimensional appearance.

THANKS FOR YOUR PATIENCE