

# INTERPROCESS COMMUNICATION

# INTERPROCESS COMMUNICATION

- Consider shell pipeline

- `cat chapter1 chapter2 chapter3 | grep tree`
- 2 processes
- Information sharing
- Order of execution

# INTERPROCESS COMMUNICATION

- Processes within a system may be **independent** or **cooperating**
- Cooperating process can affect or be affected by other processes
- Reasons for cooperating processes:
  - Information sharing
  - Computation speedup
  - Modularity
  - Convenience
- Cooperating processes require a **mechanism** to exchange data and information

# IPC ISSUES

1. How one process **passes** information to another ?
2. How to make sure that two or more processes do not get into each other's way when engaging in **critical** activities?
3. How to do proper **sequencing** when **dependencies** are present?
  - 1: easy for threads, for processes different approaches (e.g., message passing, shared memory)
  - 2 and 3: same problems and same solutions apply for threads and processes
    - **Mutual exclusion & Synchronization**

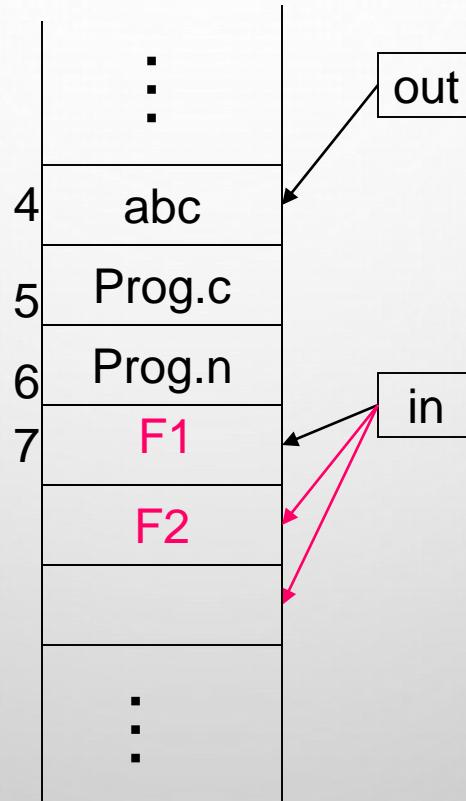
# SPOOLING EXAMPLE: CORRECT

Process 1

```
int next_free;
```

- ① `next_free = in;`
- ② Stores F1 into `next_free`;
- ③ `in=next_free+1`

Shared memory



Process 2

```
int next_free;
```

- ④ `next_free = in`
- ⑤ Stores F2 into `next_free`;
- ⑥ `in=next_free+1`

# SPOOLING EXAMPLE: RACES

Process 1

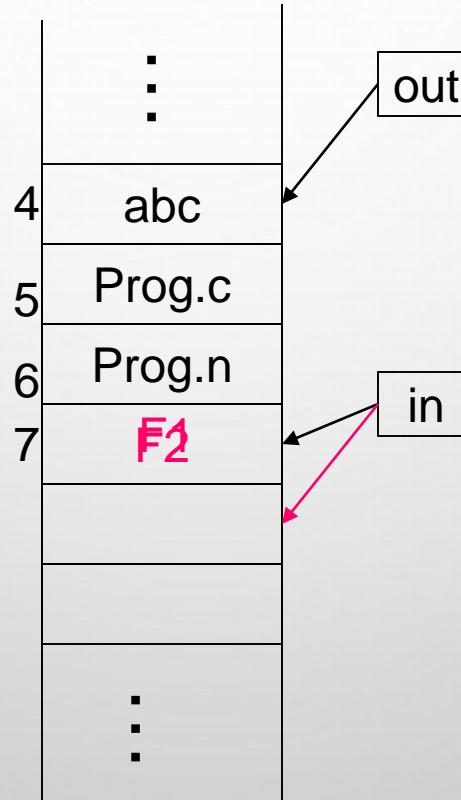
```
int next_free;
```

1 next\_free = in;

3 Stores F1 into  
next\_free;

4 in=next\_free+1

Shared memory



Process 2

```
int next_free;
```

2 next\_free = in  
/\* value: 7 \*/

5 Stores F2 into  
next\_free;

6 in=next\_free+1

# BETTER CODING?

- In previous code

```
for(;;){  
    int next_free = in;  
    slot[next_free] = file;  
    in = next_free+1;  
}
```

- What if we use one line of code?

```
for(;;){  
    slot[in++] = file  
}
```

# WHEN CAN PROCESS BE SWITCHED?

- After each **machine** instruction!
- `int++` is a C/C++ statement, translated into **three** machine instructions:
  - load mem, R
  - inc R
  - store R, mem
- Interrupt (and hence process switching) can happen in between.

# RACE CONDITION

- Two or more processes are reading or writing some **shared** data and the final result depends on who runs precisely when
- Very hard to Debug

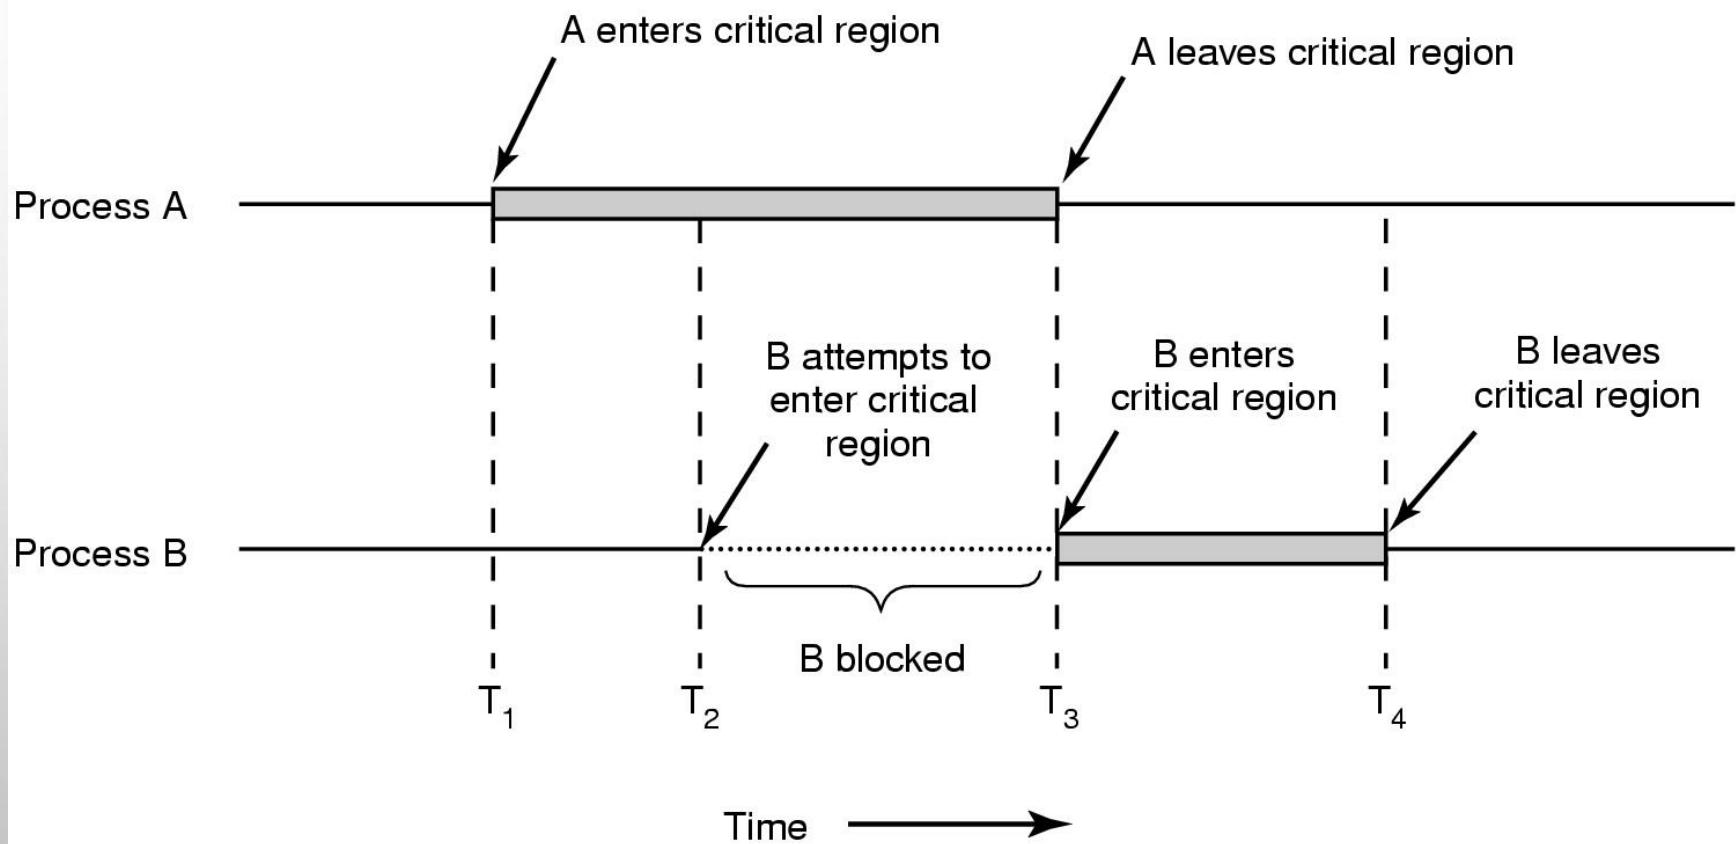
# CRITICAL REGION

- That **part** of the program that do critical things such as accessing shared memory
- Can lead to race condition

# SOLUTION REQUIREMENT

- 1) No two processes **simultaneously** in critical region
- 2) No assumptions made about speeds or numbers of CPUs
- 3) No process running **outside** its critical region may **block** another process
- 4) No process must wait forever to enter its critical region

# SOLUTION REQUIREMENT



# MUTUAL EXCLUSION WITH BUSY WAITING

- Possible Solutions
  - Disabling Interrupts
  - Lock Variables
  - Strict Alternation
  - Peterson's solution
  - TSL

# DISABLING INTERRUPTS

- How does it work?
  - Disable all interrupts just after entering a critical section
  - Re-enable them just before leaving it.
- Why does it work?
  - With interrupts disabled, no clock interrupts can occur
  - No switching can occur
- Problems:
  - What if the process forgets to enable the interrupts?
  - Multiprocessor? (disabling interrupts only affects one CPU)
- Only used **inside OS**

```
while (true) {  
    /* disable interrupts */;  
    /* critical section */;  
    /* enable interrupts */;  
    /* remainder */;  
}
```

# LOCK VARIABLES

```
int lock = 0;  
while (lock);  
lock = 1;  
//EnterCriticalSection;  
    access shared variable;  
//LeaveCriticalSection;  
lock = 0;
```

Does the above code work?

# STRICT ALTERNATION

```
while (TRUE) {  
    while (turn != 0)      /* loop */ ;  
    critical_region( );  
    turn = 1;  
    noncritical_region();  
}
```

(a)

**(a) Process 0**

Proposed solution to critical region problem

```
while (TRUE) {  
    while (turn != 1)      /* loop */ ;  
    critical_region( );  
    turn = 0;  
    noncritical_region();  
}
```

(b)

**(b) Process 1**

# PROBLEM

- Busy waiting: Continuously testing a variable until some value appear
  - Wastes CPU time
- Violates condition 3
  - When one process is much slower than the other

# PETERSON'S SOLUTION

- Consists of 2 procedures
- Each process has to call
  - enter\_region with its own process # before entering its C.R.
  - And Leave\_region after leaving C.R.

do {

enter\_region(process#)

critical section

leave\_region(process#)

remainder section

} while (TRUE);

# PETERSON'S SOLUTION (FOR 2 PROCESSES)

```
#define FALSE 0
#define TRUE 1
#define N      2           /* number of processes */

int turn;                  /* whose turn is it? */
int interested[N];         /* all values initially 0 (FALSE) */

void enter_region(int process); /* process is 0 or 1 */
{
    int other;               /* number of the other process */

    other = 1 - process;    /* the opposite of process */
    interested[process] = TRUE; /* show that you are interested */
    turn = process;          /* set flag */
    while (turn == process && interested[other] == TRUE) /* null statement */ ;
}

void leave_region(int process) /* process: who is leaving */
{
    interested[process] = FALSE; /* indicate departure from critical region */
}
```

# PETERSON'S SOLUTION: ANALYSIS(1)

- Let Process 1 is not interested and Process 0 calls enter\_region with 0
- So, turn = 0 and interested[0] = true and Process 0 is in CR
- Now if Process 1 calls enter\_region, it will hang there until interested[0] is false. Which only happens when Process 0 calls leave\_region i.e. leaves the C.R.

# PETERSON'S SOLUTION: ANALYSIS(2)

- Let both processes call `enter_region` **simultaneously**
- Say `turn = 1.` (i.e. Process 1 stores **last**)
- Process 0 enters critical region: `while (turn == 0 && ...)` returns **false** since `turn = 1.`
- Process 1 loops until process 0 exits: `while (turn == 1 && interested[0] == true)` returns `true`.
- Done!!

# TSL

- Requires hardware support
- TSL instruction: test and set lock
  - Reads content of *lock* into a Register
  - Stores a nonzero value at *lock*.
- CPU executing TSL locks the memory bus prohibiting other CPUs from accessing memory

*TSL REGISTER,LOCK*

Indivisible/Atomic

enter\_region:

TSL REGISTER,LOCK

| copy lock to register and set lock to 1

CMP REGISTER,#0

| was lock zero?

JNE enter\_region

| if it was non zero, lock was set, so loop

RET | return to caller; critical region entered

leave\_region:

MOVE LOCK,#0

| store a 0 in lock

RET | return to caller

# BUSY WAITING: PROBLEMS

- Waste CPU time since it sits on a tight loop
- May have unexpected effects:
  - Priority Inversion Problem

Example:

- 2 Cooperating Processes: H (high priority) and L (low priority)
- Scheduling rule: H is run whenever it is ready
- Let L in C. R. and H is ready and wants to enter C.R.
- Since H is ready it is given the CPU and it starts busy waiting
- L will never gets the chance to leave its C.R.
- H loops forever
- [http://research.microsoft.com/en-us/um/people/mbj/Mars\\_Pathfinder/Mars\\_Pathfinder.html](http://research.microsoft.com/en-us/um/people/mbj/Mars_Pathfinder/Mars_Pathfinder.html)

# SLEEP & WAKEUP

- When a process has to **wait**, change its **state** to **BLOCKED/WAITING**
- Switched to **READY** state, when it is OK to retry entering the critical section
- Sleep is a **system call** that causes the caller to block
  - be suspended until another process wakes it up
- Wakeup system call has one parameter, the process to be awakened.

# PRODUCER CONSUMER PROBLEM

- Also called bounded-buffer problem
- Two ( $m+n$ ) processes share a **common** buffer
- One ( $m$ ) of them is (are) **producer**(s): put(s) information in the buffer
- One ( $n$ ) of them is (are) **consumer**(s): take(s) information out of the buffer
- Trouble and solution
  - Producer wants to put but buffer **full**- Go to **sleep** and **wake up** when consumer takes one or more
  - Consumer wants to take but buffer **empty**- go to sleep and wake up when producer puts one or more

# SLEEP AND WAKEUP

```
#define N 100                                /* number of slots in the buffer */
int count = 0;                                /* number of items in the buffer */

void producer(void)
{
    int item;

    while (TRUE) {
        item = produce_item();                /* repeat forever */
        if (count == N) sleep();              /* generate next item */
        insert_item(item);                  /* if buffer is full, go to sleep */
        count = count + 1;                  /* put item in buffer */
        if (count == 1) wakeup(consumer);    /* increment count of items in buffer */
                                            /* was buffer empty? */
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {
        if (count == 0) sleep();            /* repeat forever */
        item = remove_item();             /* if buffer is empty, got to sleep */
        count = count - 1;               /* take item out of buffer */
        if (count == N - 1) wakeup(producer); /* decrement count of items in buffer */
                                            /* was buffer full? */
        consume_item(item);              /* print item */
    }
}
```

# SLEEP AND WAKEUP

```
#define N 100                                /* number of slots in the buffer */
int count = 0;                                /* number of items in the buffer */

void producer(void)
{
    int item;

    while (TRUE) {
        item = produce_item();
        if (count == N) sleep();
        insert_item(item);
        count = count + 1;
        if (count == 1) wakeup(consumer);
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {
        if (count == 0) sleep();
        item = remove_item();
        count = count - 1;
        if (count == N - 1) wakeup(producer);
        consume_item(item);
    }
}
```

# SLEEP AND WAKEUP: RACE CONDITION

- **Race condition**
- Unconstrained access to *count*
  - CPU is given to P just after C has **read** count to be 0 but not yet gone to sleep.
  - P calls wakeup
  - Result is **lost** wake-up signal
  - Both will sleep forever



# SEMAPHORES

- A new variable type
- A kind of **generalized** lock
  - First defined by Dijkstra in late 60s
  - Main synchronization **primitive** used in original UNIX
- Semaphores are like integers, except
  - No negative values
  - Only operations allowed are *up* and *down* – can't read or write value, except to set it initially

# SEMAPHORES

- Operation “down”:
  - if value > 0; value-- and then continue.
  - if value = 0; process is put to sleep without completing the down for the moment
    - Checking the value, changing it, and possibly going to sleep, is all done as an **atomic** action.
- Operation “up”:
  - increments the value of the semaphore addressed.
  - If one or more process were sleeping on that semaphore, one of them is chosen by the system (e.g. at **random**) and is allowed to complete its *down*
    - The operation of incrementing the semaphore and waking up one process is also **indivisible**
  - No process ever blocks doing an *up*.

# SEMAPHORES

- Operations must be **atomic**
  - Two *down*'s together can't decrement value below zero
  - Similarly, process going to sleep in *down* won't miss wakeup from *up* – even if they both happen at same time

# SEMAPHORES

- **Counting semaphore.**
  - The value can range over an unrestricted domain
- **Binary semaphore**
  - The value can range only between 0 and 1.
  - On some systems, binary semaphores are known as **mutex** locks as they provide mutual exclusion

# SEMAPHORES USAGE

1. Mutual exclusion
2. Controlling access to **limited** resource
3. Synchronization

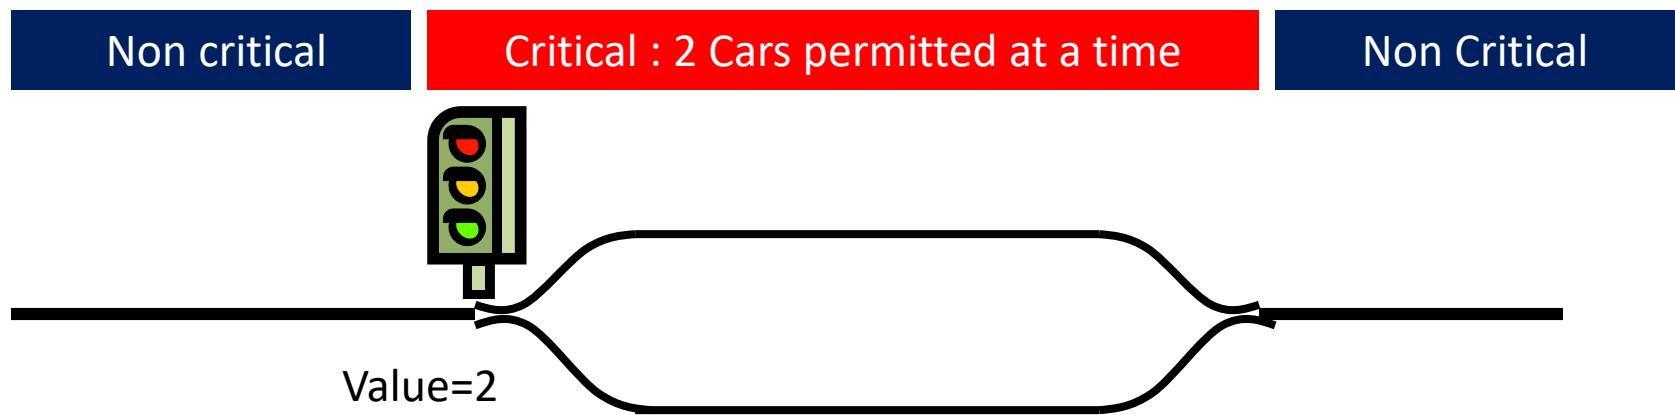
# Mutual exclusion

- How to ensure that only one process can enter its C.R.?
- Binary semaphore **initialized to 1**
- Shared by all collaborating processes
- If each process does a *down* just before entering CR and an *up* just after leaving then mutual exclusion is guaranteed

```
do  {
    down (mutex) ;
        // critical section
    up (mutex) ;
        // remainder section
} while (TRUE);
```

# Controlling access to a resource

- What if we want maximum  $m$  process/thread can use a resource simultaneously ?
- Counting semaphore **initialized to the number of available resources**
- Semaphore from railway analogy
  - Here is a semaphore **initialized to 2** for resource control:



# Synchronization

- How to resolve dependency among processes
- Binary semaphore **initialized to 0**
- consider 2 concurrently running processes:
  - P1 with a statement S1 and
  - P2 with a statement S2.
  - Suppose we require that S2 be executed only after S1 has completed.

P1

```
S1;  
up(synch);
```

P2

```
down(synch);  
S2;
```

# PRODUCER & CONSUMER

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;
void producer(void)
{
    int item;

    while (TRUE) {
        item = produce_item();
        down(&empty);
        down(&mutex);
        insert_item(item);
        up(&mutex);
        up(&full);
    }
}
```

```
/* number of slots in the buffer */
/* semaphores are a special kind of int */
/* controls access to critical region */
/* counts empty buffer slots */
/* counts full buffer slots */
void consumer(void)
{
    int item;

    while (TRUE) {
        down(&full);
        down(&mutex);
        item = remove_item();
        up(&mutex);
        up(&empty);
        consume_item(item);
    }
}
```

# SEMAPHORES IN PRODUCER CONSUMER PROBLEM: ANALYSIS

- 3 semaphores are used
  - *full* (initially 0) for counting **occupied** slots
  - *Empty* (initially  $N$ ) for counting **empty** slots
  - *mutex* (initially 1) to make sure that Producer and Consumer do not access the buffer at the same time
- Here 2 uses of semaphores
  - Mutual exclusion (mutex)
  - Synchronization (full and empty)
    - To guarantee that certain event sequences do or do not occur

**Block on:**

Producer: insert in **full** buffer

**Unblock on:**

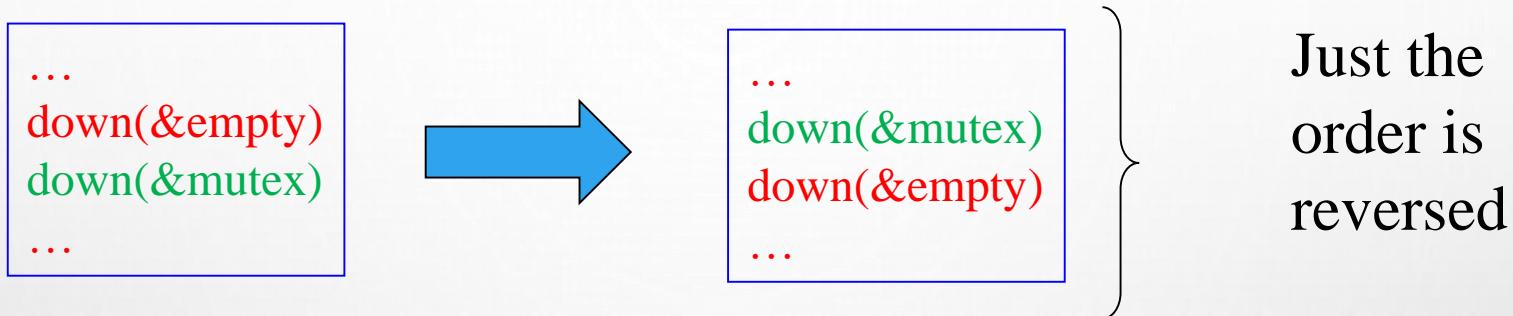
Consumer: item **inserted**

Consumer: remove from **empty** buffer

Producer: item **removed**

# SEMAPHORES: “BE CAREFUL”

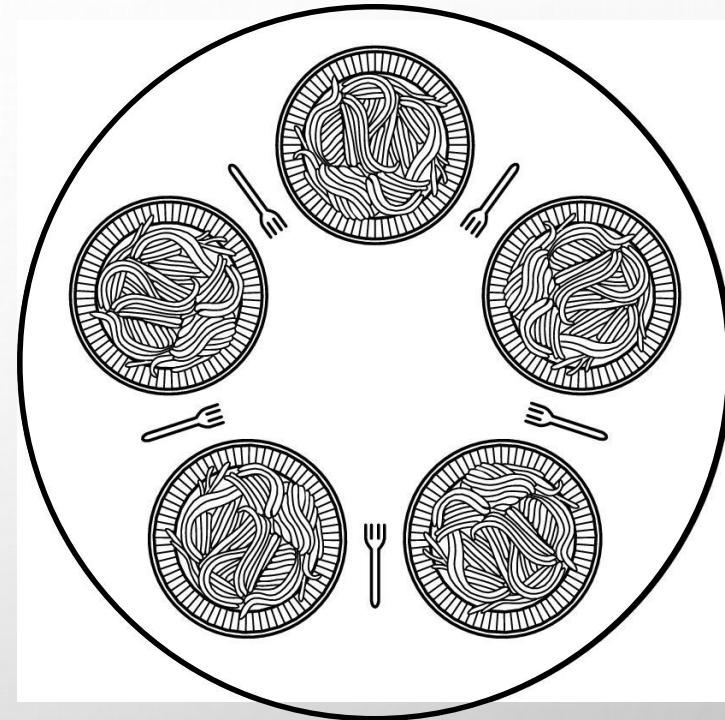
- Suppose the following is done in **Producer's** code



- If buffer **full** Producer would block due to `down(&empty)` with `mutex = 0`.
- So now if Consumer tries to access the buffer, it would block too due to its `down(&mutex)`.
- Both processes would stay blocked **forever**: **DEADLOCK**

# DINING PHILOSOPHERS

- Philosophers spend their lives **thinking** and **eating**
- Don't interact with their neighbors
- When get hungry try to pick up 2 chopsticks (one at a time in either order) to eat
- Need both to eat, then release both when done
- How to program the scenario avoiding all concurrency problems?



# DINING PHILOSOPHERS: A SOLUTION

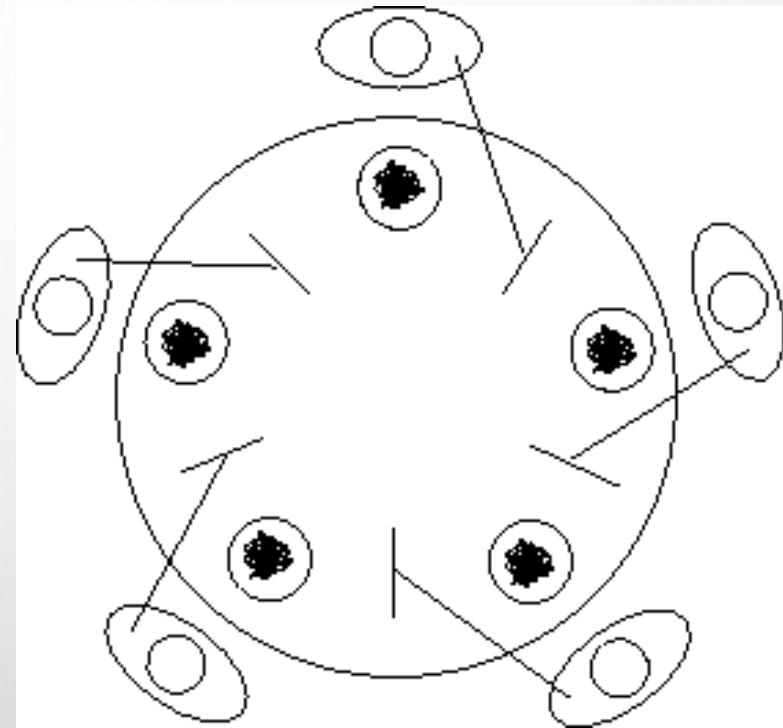
```
#define N 5                                     /* number of philosophers */

void philosopher(int i)                         /* i: philosopher number, from 0 to 4 */

{
    while (TRUE) {
        think();                                /* philosopher is thinking */
        take_fork(i);                            /* take left fork */
        take_fork((i+1) % N);                   /* take right fork; % is modulo operator */
        eat();                                   /* yum-yum, spaghetti */
        put_fork(i);                            /* put left fork back on the table */
        put_fork((i+1) % N);                   /* put right fork back on the table */
    }
}
```

# DINING PHILOSOPHERS: PROBLEMS WITH PREVIOUS SOLUTION

- Deadlock may happen
- Does this solution prevents any such thing from happening ?
  - Everyone takes the **left** fork simultaneously



# DINING PHILOSOPHERS: PROBLEMS WITH PREVIOUS SOLUTION

Tentative Solution:

- After taking left fork, check whether right fork is available.
- If not, then return left one, **wait for some time** and repeat again.

Problem:

- All of them start and do the algorithm synchronously and simultaneously:
- **STARVATION** (A situation in which all the programs run indefinitely but fail to make any progress)
- Solution: **Random** wait; but what if the most unlikely of **same** random number happens?

# ANOTHER ATTEMPT, SUCCESSFUL!

```
void philosopher(int i)
{
    while (true)
    {
        think();
        down(&mutex);
        take_fork(i);
        take_fork((i+1)%N);
        eat();
        put_fork(i);
        put_fork((i+1)%N);
        up(&mutex);
    }
}
```

- Theoretically solution is OK - no deadlock, no starvation.
- Practically with a performance bug:
  - Only **one** philosopher can be eating at any instant: absence of parallelism

# FINAL SOLUTION PART 1

```
#define N          5           /* number of philosophers */
#define LEFT        (i+N-1)%N   /* number of i's left neighbor */
#define RIGHT       (i+1)%N    /* number of i's right neighbor */
#define THINKING    0           /* philosopher is thinking */
#define HUNGRY      1           /* philosopher is trying to get forks */
#define EATING      2           /* philosopher is eating */
typedef int semaphore;
int state[N];
semaphore mutex = 1;
semaphore s[N];

void philosopher(int i)
{
    while (TRUE) {
        think();
        take_forks(i);
        eat();
        put_forks(i);
    }
}
```

/\* i: philosopher number, from 0 to N-1 \*/

```
/* repeat forever */
/* philosopher is thinking */
/* acquire two forks or block */
/* yum-yum, spaghetti */
/* put both forks back on table */
```

# FINAL SOLUTION PART 2

```
void take_forks(int i)                                /* i: philosopher number, from 0 to N-1 */  
{  
    down(&mutex);  
    state[i] = HUNGRY;  
    test(i);  
    up(&mutex);  
    down(&s[i]);  
}  
  
void put_forks(i)                                     /* i: philosopher number, from 0 to N-1 */  
{  
    down(&mutex);  
    state[i] = THINKING;  
    test(LEFT);  
    test(RIGHT);  
    up(&mutex);  
}  
  
void test(i)                                         /* i: philosopher number, from 0 to N-1 */  
{  
    if (state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {  
        state[i] = EATING;  
        up(&s[i]);  
    }  
}
```

# THE READERS AND WRITERS PROBLEM

- Dining Philosopher Problem: Models processes that are competing for **exclusive** access to a **limited** resource
- Readers Writers Problem: Models access to a **database**
- Example: An airline reservation system- many competing process wishing to read and write-
  - Multiple readers simultaneously- acceptable
  - Multiple writers simultaneously- not acceptable
  - Reading, while write is writing- not acceptable

```

typedef int semaphore;
semaphore mutex = 1;
semaphore db = 1;
int rc = 0;

void reader(void)
{
    while (TRUE) {
        down(&mutex);
        rc = rc + 1;
        if (rc == 1) down(&db);
        up(&mutex);
        read_data_base();
        down(&mutex);
        rc = rc - 1;
        if (rc == 0) up(&db);
        up(&mutex);
        use_data_read();
    }
}

void writer(void)
{
    while (TRUE) {
        think_up_data();
        down(&db);
        write_data_base();
        up(&db);
    }
}

/* use your imagination */
/* controls access to 'rc' */
/* controls access to the database */
/* # of processes reading or wanting to */

/* repeat forever */
/* get exclusive access to 'rc' */
/* one reader more now */
/* if this is the first reader ... */
/* release exclusive access to 'rc' */
/* access the data */
/* get exclusive access to 'rc' */
/* one reader fewer now */
/* if this is the last reader ... */
/* release exclusive access to 'rc' */
/* noncritical region */

```

# ISSUE REGARDING THE SOLUTION

- Inherent **priority** to the readers
- Say a new reader arrives every 2 seconds and each reader takes 5 seconds to do its work. What will happen to a writer?
- Issue regarding second variation
  - Writer don't have to wait for readers that came along **after** it
  - Less concurrency, lower performance

**THANKS FOR YOUR  
PATIENCE**