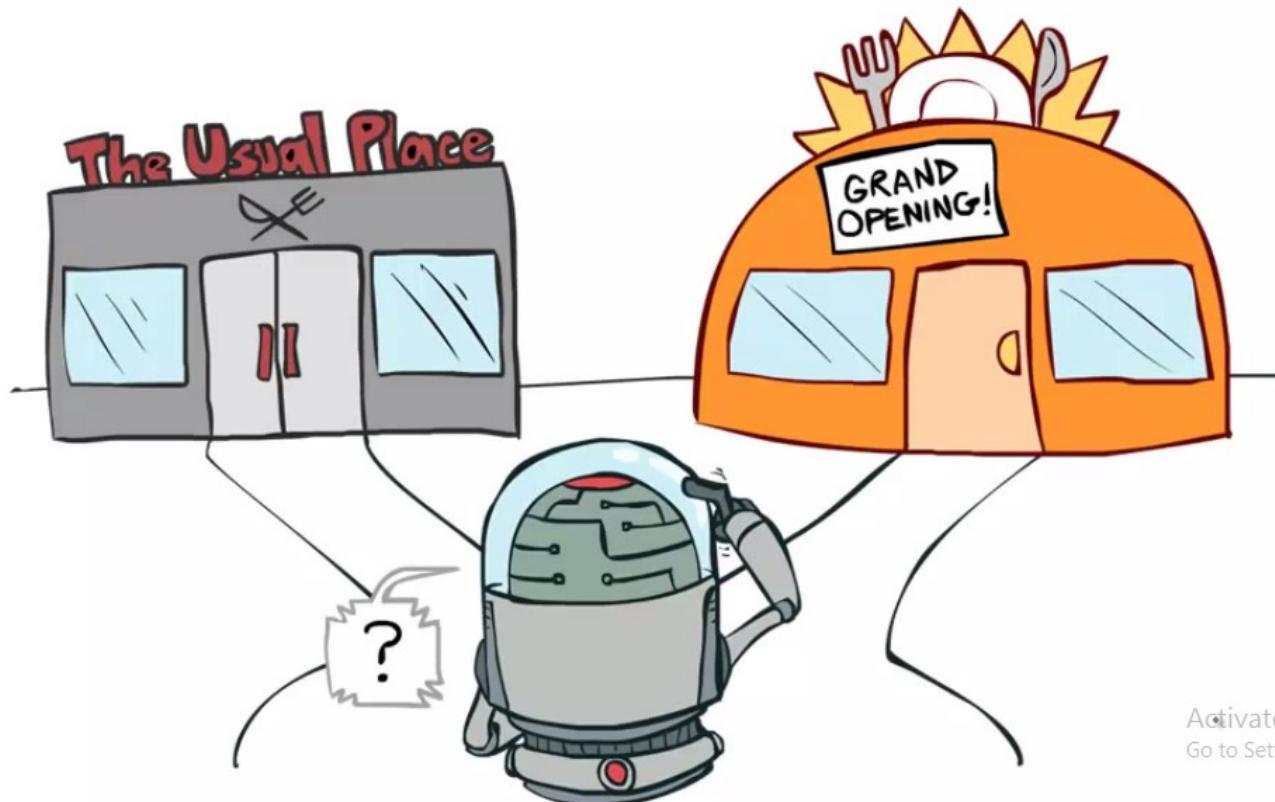


# Lecture 17

Amit Kumar Das

Senior Lecturer,  
Department of Computer Science &  
Engineering,  
East West University  
Dhaka, Bangladesh.

# Exploration vs. Exploitation



Activate Windows  
Go to Settings to activate Windows.

# Exploration vs. exploitation

- **Exploration:** take a new action with unknown consequences
  - Pros:
    - Get a more accurate model of the environment
    - Discover higher-reward states than the ones found so far
  - Cons:
    - When you're exploring, you're not maximizing your utility
    - Something bad might happen
- **Exploitation:** go with the best strategy found so far
  - Pros:
    - Maximize reward as reflected in the current utility estimates
    - Avoid bad stuff
  - Cons:
    - Might also prevent you from discovering the true optimal strategy

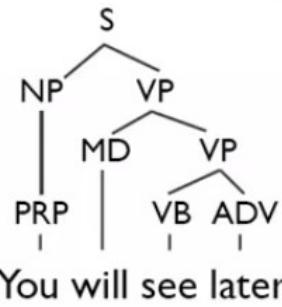
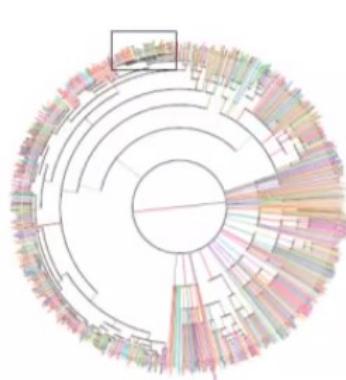
# How to Explore?

- Several schemes for forcing exploration
  - Simplest: random actions ( $\epsilon$ -greedy)
    - Every time step, flip a coin
    - With (small) probability  $\epsilon$ , act randomly
    - With (large) probability  $1-\epsilon$ , act on current policy
  - Problems with random actions?
    - You do eventually explore the space, but keep thrashing around once learning is done
    - One solution: lower  $\epsilon$  over time
    - Another solution: exploration functions



Activate Windows  
Go to Settings to activate Windows.

# Now: Advanced Applications

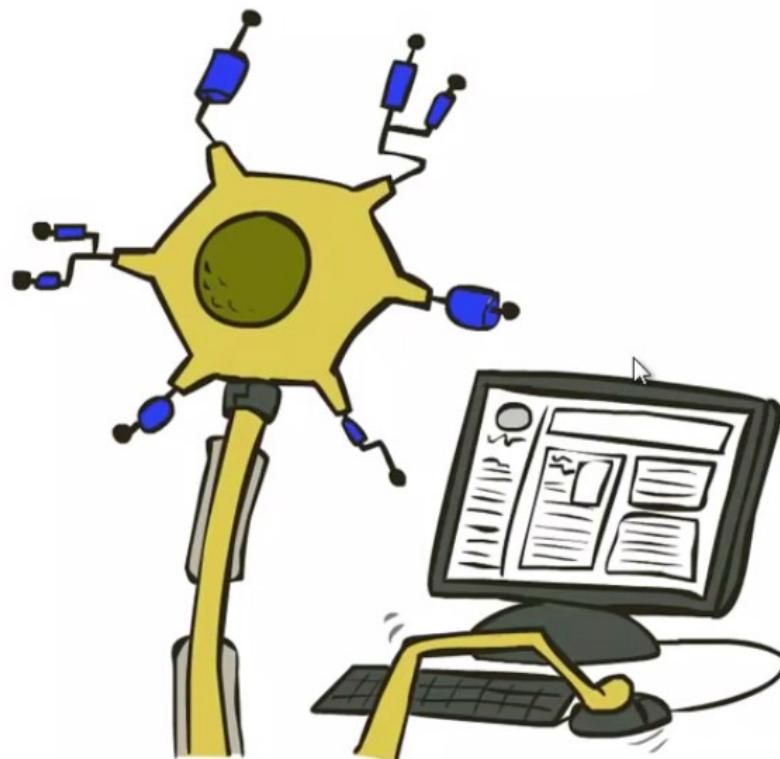


Después lo verás



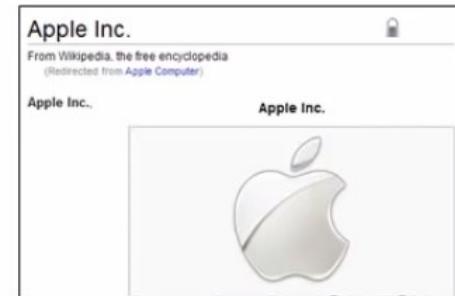
# Web Search

---



# Extension: Web Search

- Information retrieval:
    - Given information needs, produce information
    - Includes, e.g. web search, question answering, and classic IR
  - Web search: not exactly classification, but rather ranking
- $x = \text{“Apple Computers”}$



# Feature-Based Ranking

$x$  = “Apple Computer”

$f(x,$



) = [0.3 5 0 0 ...]

$f(x,$



) = [0.8 4 2 1 ...]

# What is NLP?



# What is NLP?



# What is NLP?



- Fundamental goal: analyze and process human language, broadly, robustly, accurately...
- End systems that we want to build:
  - Ambitious: speech recognition, machine translation, information extraction, dialog interfaces, question answering...
  - Modest: spelling correction, text categorization...

# Problem: Ambiguities

- Headlines:
  - Enraged Cow Injures Farmer With Ax

# Problem: Ambiguities

- Headlines:
  - Enraged Cow Injures Farmer With Ax



# Machine Translation



# Machine Translation

## "Il est impossible aux journalistes de rentrer dans les régions tibétaines"

Bruno Philip, correspondant du "Monde" en Chine, estime que les journalistes de l'AFP qui ont été expulsés de la province tibétaine du Qinghai "n'étaient pas dans l'illégalité".

**Les faits** Le dalaï-lama dénonce l'"enfer" imposé au Tibet depuis sa fuite, en 1959

**Vidéo** Anniversaire de la rébellion tibétaine: La Chine sur son gardien



## "It is impossible for journalists to enter Tibetan areas"

Philip Bruno, correspondent for "World" in China, said that journalists of the AFP who have been deported from the Tibetan province of Qinghai "were not illegal."

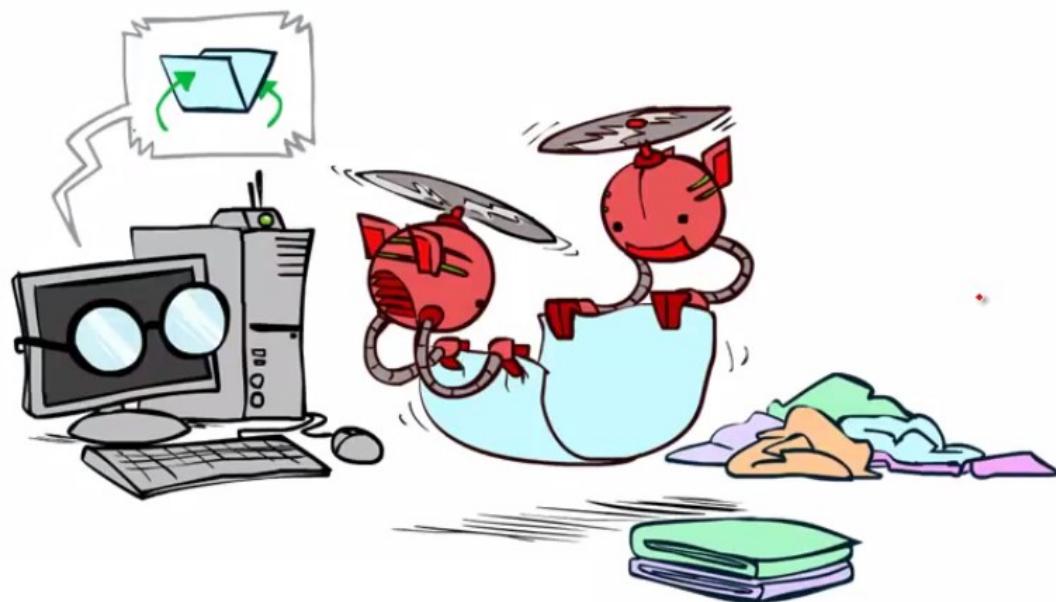
**Facts** The Dalai Lama denounces the "hell" imposed since he fled Tibet in 1959

**Video** Anniversary of the Tibetan rebellion: China on guard



- Translate text from one language to another
- Recombines fragments of example translations
- Challenges:
  - What fragments? [learning to translate]
  - How to make efficient? [fast translation search]

## Advanced Applications: Computer Vision and Robotics \*



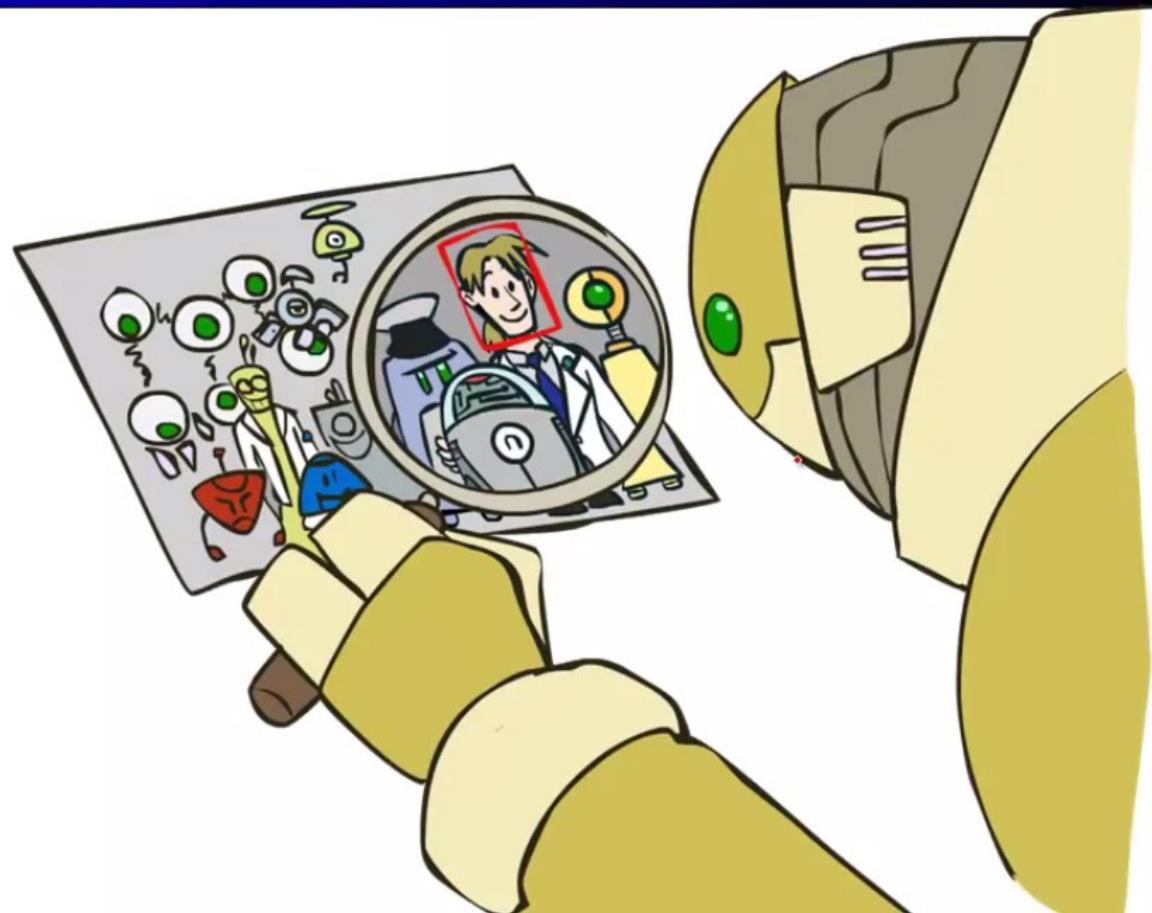
# Computer Vision

---



# Object Detection

---

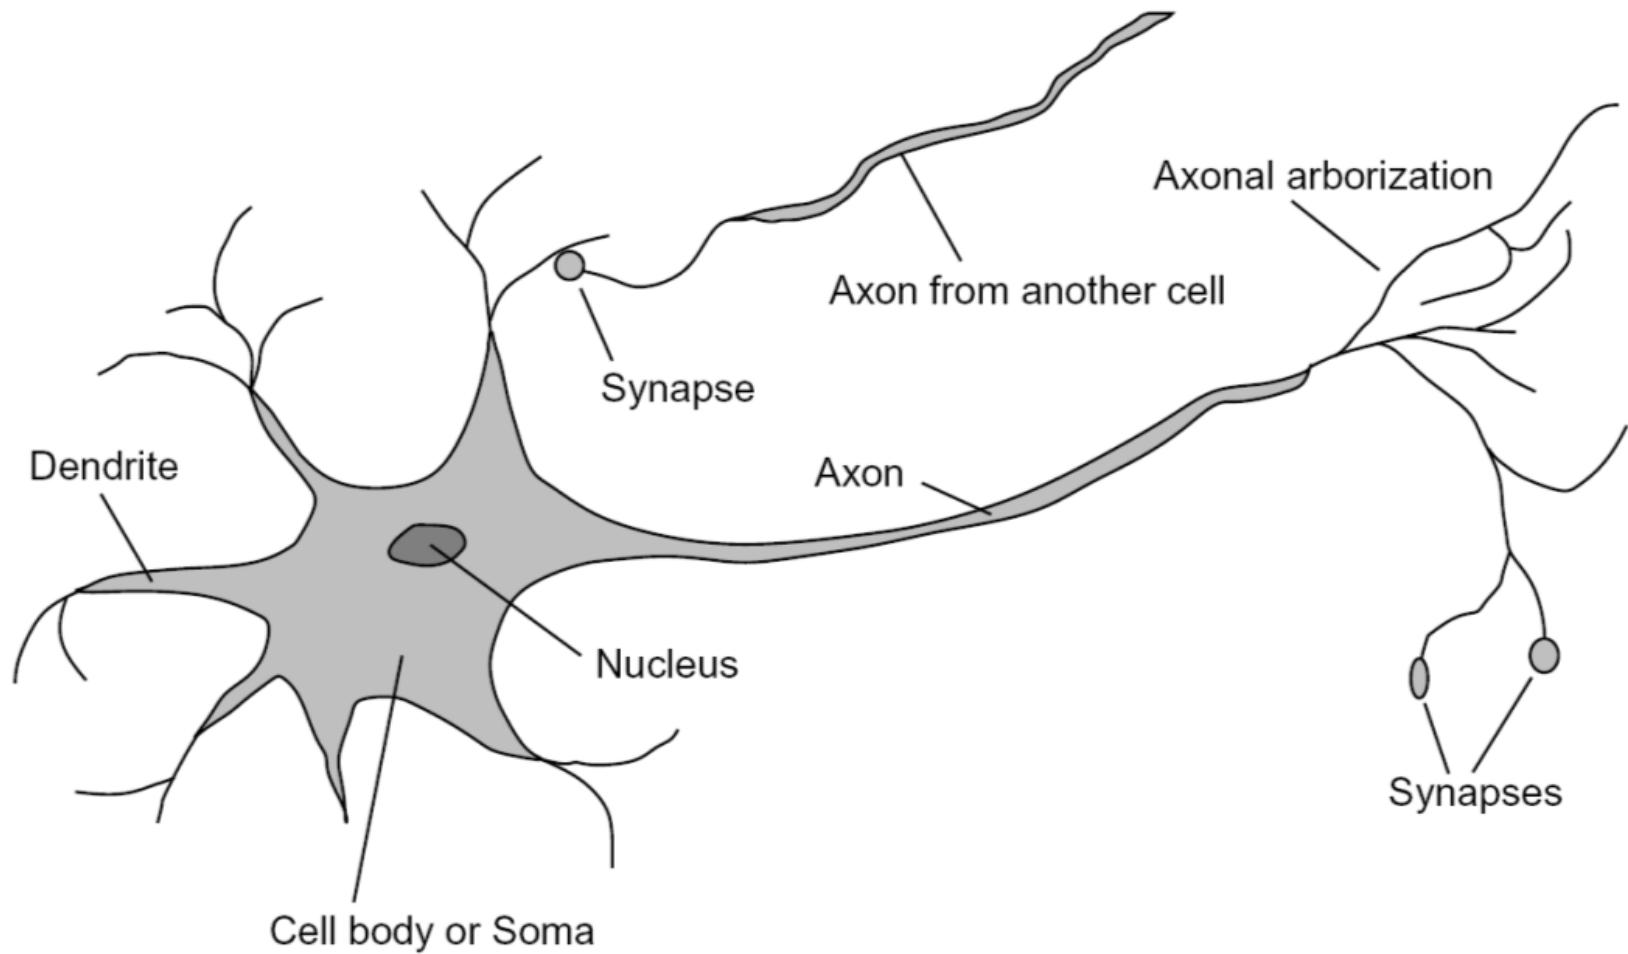


# Object Detection Approach 1: HOG + SVM

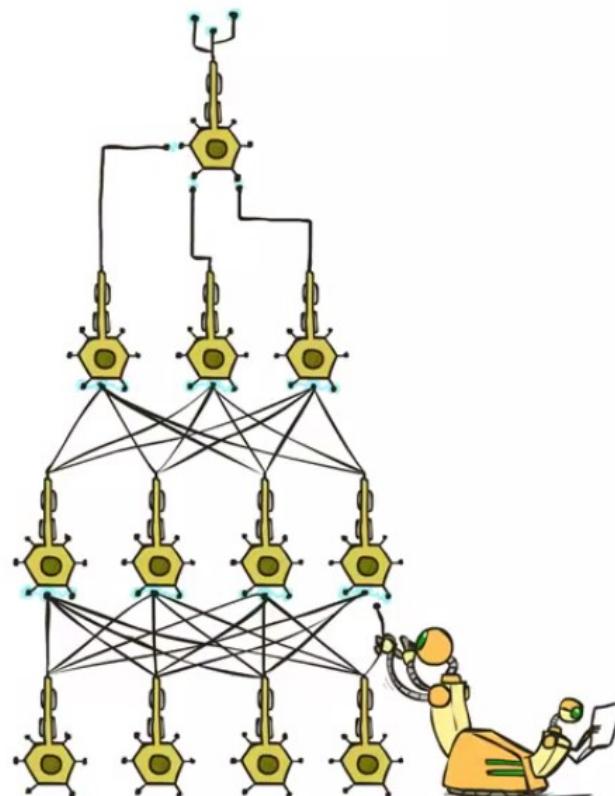
---



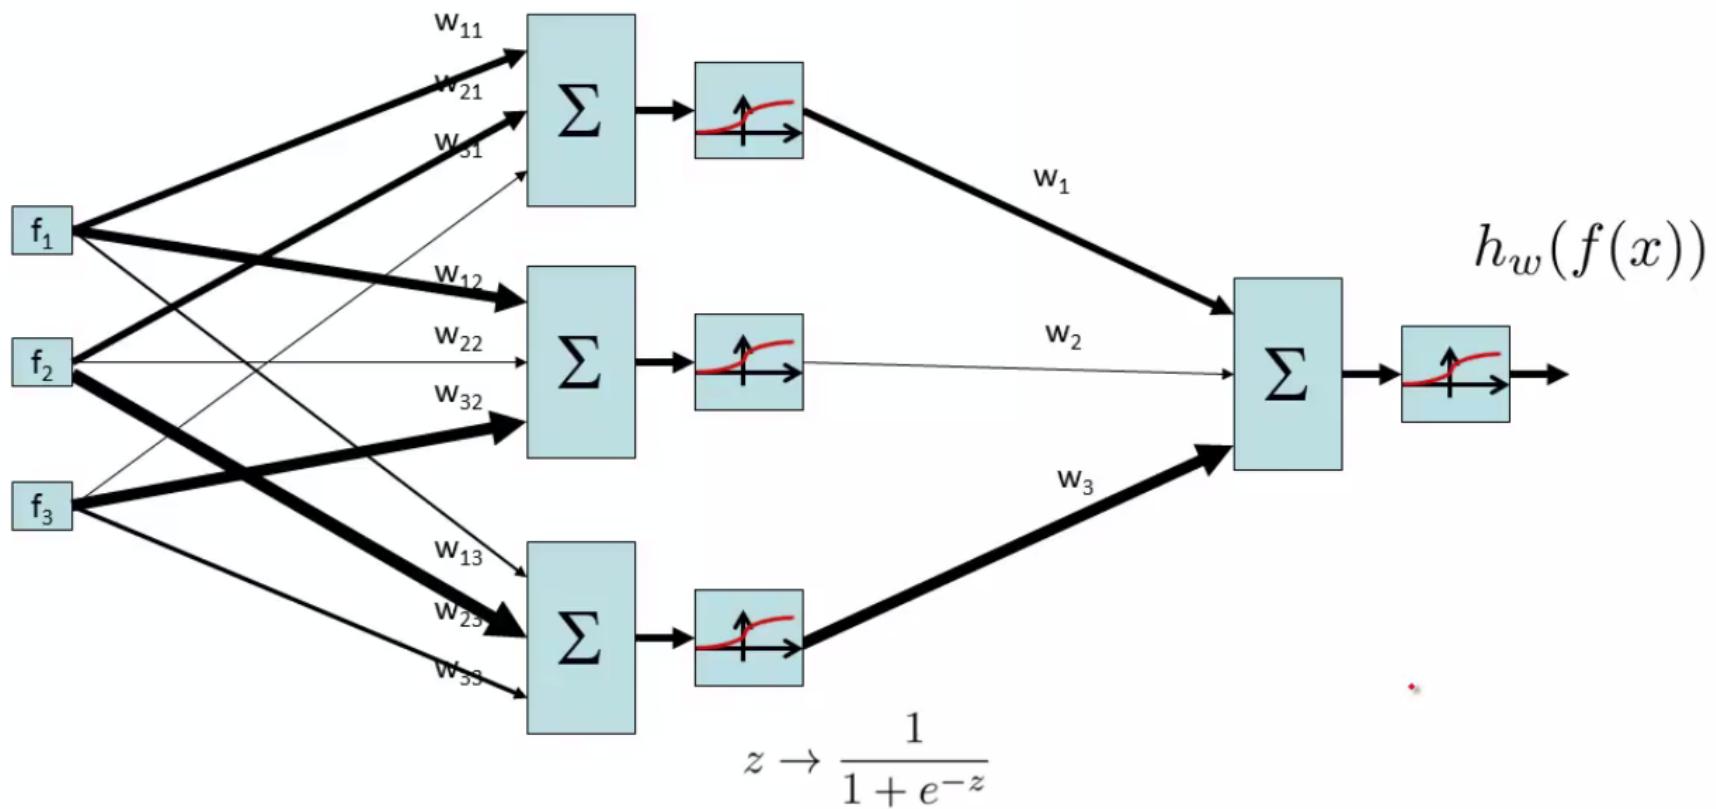
# Loose inspiration: Human neurons



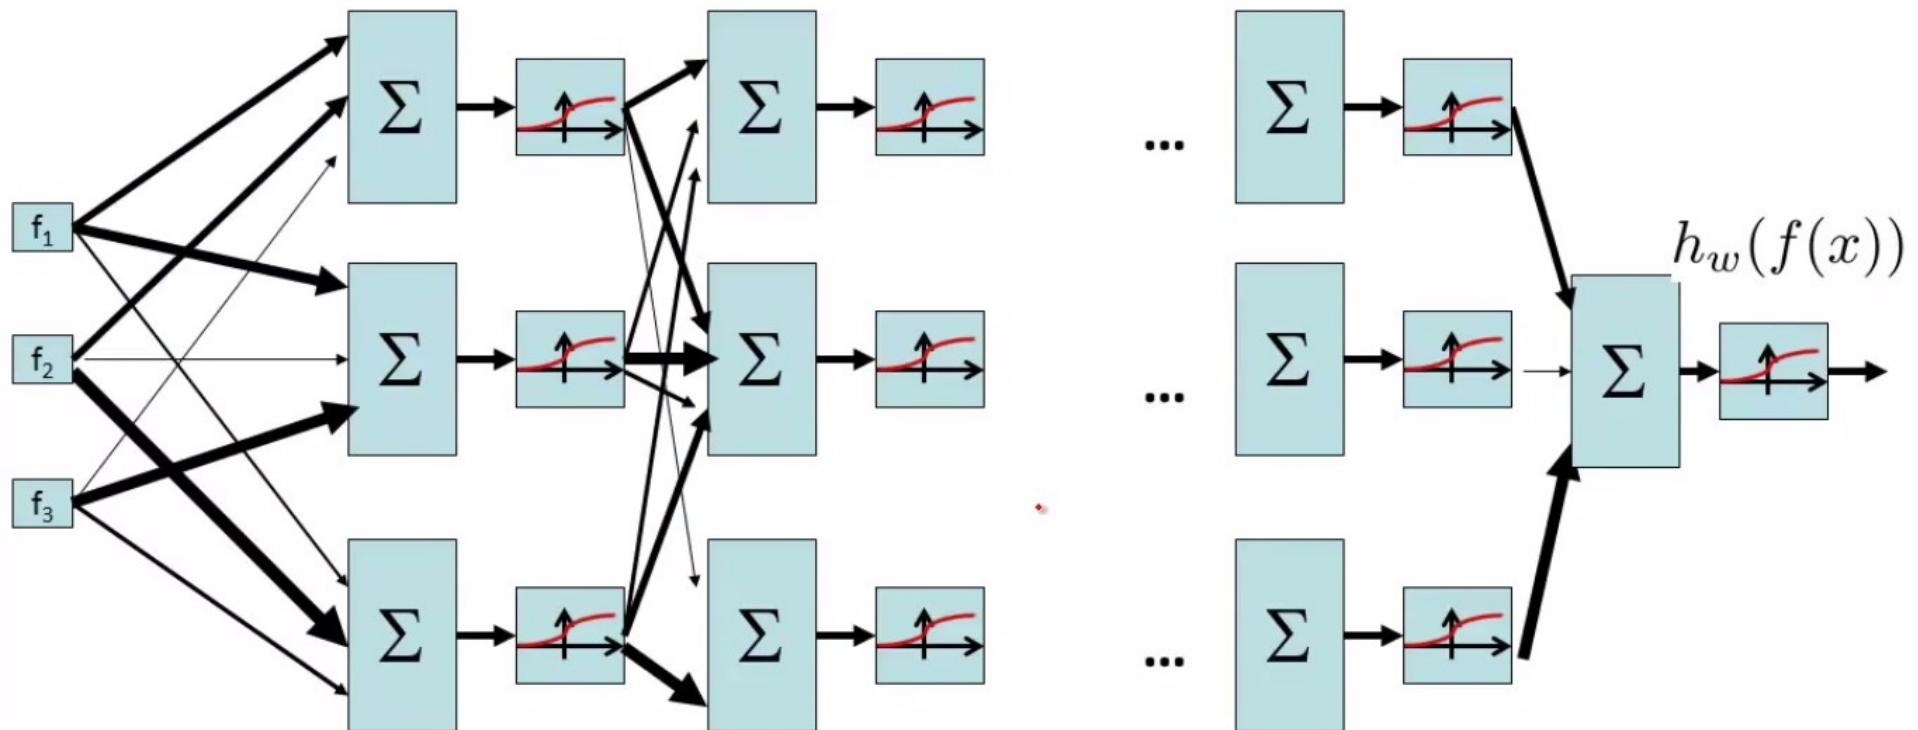
# Object Detection Approach 2: Deep Learning



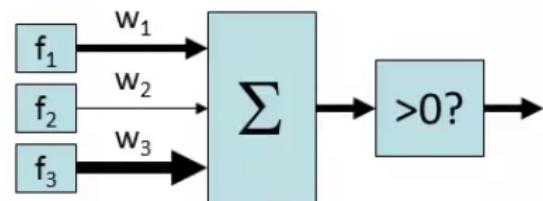
# Two-Layer Neural Network



# N-Layer Neural Network



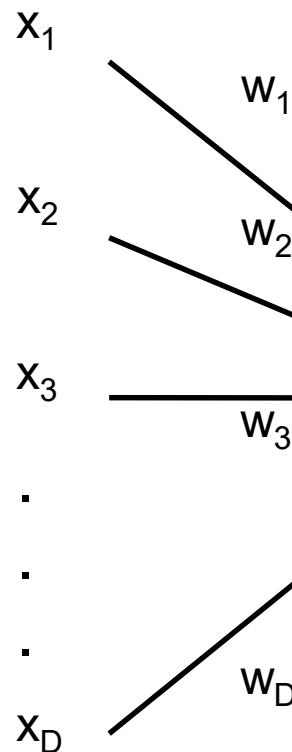
# Perceptron



# Perceptron

Input

Weights



Output:  $\text{sgn}(w \cdot x + b)$

Can incorporate bias as component of the weight vector by always including a feature with value set to 1

# Perceptron

- Current weight vector:  $w_1 = (1, 2, -2)$ ,  $w_2 = (3, -2, -1)$ ,  $w_3 = (-1, 2, 4)$
- Next training data point:  $f(x) = (1, -0.5, 3)$ ,  $y = 2$
- Perceptron update:

# Perceptron

- Current weight vector:  $w_1 = (1, 2, -2)$ ,  $w_2 = (3, -2, -1)$ ,  $w_3 = (-1, 2, 4)$
- Next training data point:  $f(x) = (1, -0.5, 3)$ ,  $y = 2$
- Perceptron update:

$$\left. \begin{array}{l} w_1 \cdot f(x) = 1 \cdot 1 + 2 \cdot (-0.5) + (-2) \cdot 3 = -6 \\ w_2 \cdot f(x) = 3 \cdot 1 - 2 \cdot (-0.5) - 1 \cdot 3 = 1 \\ w_3 \cdot f(x) = -1 \cdot 1 + 2 \cdot (-0.5) + 4 \cdot 3 = 10 \end{array} \right\}$$

# Perceptron

- Current weight vector:  $w_1 = (1, 2, -2)$ ,  $w_2 = (3, -2, -1)$ ,  $w_3 = (-1, 2, 4)$
- Next training data point:  $f(x) = (1, -0.5, 3)$ ,  $y = 2$
- Perceptron update:

$$\left. \begin{array}{l} w_1 \cdot f(x) = 1 \cdot 1 + 2 \cdot (-0.5) + (-2) \cdot 3 = -6 \\ w_2 \cdot f(x) = 3 \cdot 1 - 2 \cdot (-0.5) - 1 \cdot 3 = 1 \\ w_3 \cdot f(x) = -1 \cdot 1 + 2 \cdot (-0.5) + 4 \cdot 3 = 10 \end{array} \right\}$$

→ prediction:  $y = 3$

# Perceptron

- Current weight vector:  $w_1 = (1, 2, -2), w_2 = (3, -2, -1), w_3 = (-1, 2, 4)$
- Next training data point:  $f(x) = (1, -0.5, 3), y^* = 2$
- Perceptron update:

$$w_1 \cdot f(x) = 1 \cdot 1 + 2 \cdot (-0.5) + (-2) \cdot 3 = -6 \quad \}$$

$$w_2 \cdot f(x) = 3 \cdot 1 - 2 \cdot (-0.5) - 1 \cdot 3 = 1 \quad \}$$

$$w_3 \cdot f(x) = -1 \cdot 1 + 2 \cdot (-0.5) + 4 \cdot 3 = 10 \quad \}$$

$\rightarrow$  prediction:  $y = 3$        $w_2 \leftarrow w_2 + f(x)$   
 $(3, -2, -1) + (1, -0.5, 3) = (4, -2.5, 2)$

# Perceptron training algorithm

- Initialize weights
- Cycle through training examples in multiple passes (*epochs*)
- For each training example:
  - If classified correctly, do nothing
  - If classified incorrectly, update weights

# Perceptron update rule

- For each training instance  $\mathbf{x}$  with label  $y$ :
  - Classify with current weights:  $y' = \text{sgn}(\mathbf{w} \cdot \mathbf{x})$
  - Update weights:  $\mathbf{w} \leftarrow \mathbf{w} + \alpha(y - y')\mathbf{x}$
  - $\alpha$  is a learning rate that should decay as a function of epoch  $t$ , e.g.,  $1000/(1000+t)$
  - What happens if  $y'$  is correct?
  - Otherwise, consider what happens to individual weights  $w_i \leftarrow w_i + \alpha(y - y')x_i$ 
    - If  $y = 1$  and  $y' = -1$ ,  $w_i$  will be increased if  $x_i$  is positive or decreased if  $x_i$  is negative  $\rightarrow \mathbf{w} \cdot \mathbf{x}$  will get bigger
    - If  $y = -1$  and  $y' = 1$ ,  $w_i$  will be decreased if  $x_i$  is positive or increased if  $x_i$  is negative  $\rightarrow \mathbf{w} \cdot \mathbf{x}$  will get smaller

# Convergence of perceptron update rule

- **Linearly separable data:** converges to a perfect solution
- **Non-separable data:** converges to a minimum-error solution assuming learning rate decays as  $O(1/t)$  and examples are presented in random sequence

# Implementation details

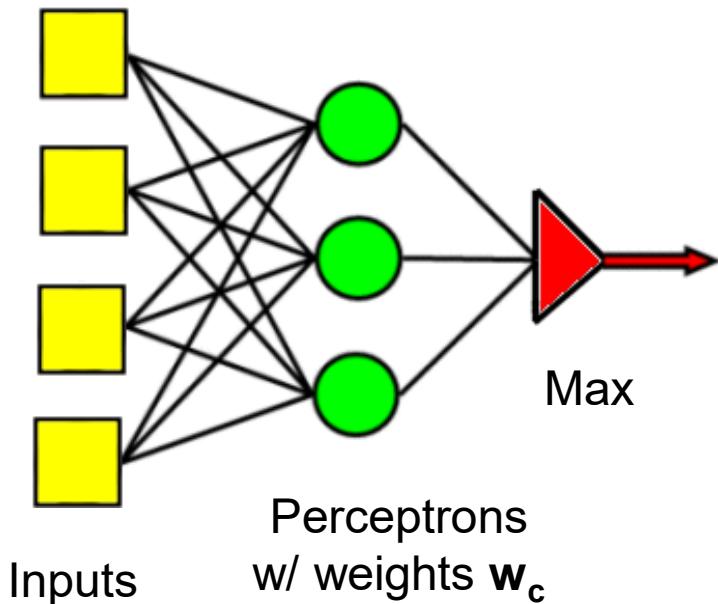
- Bias (add feature dimension with value fixed to 1) vs. no bias
- Initialization of weights: all zeros vs. random
- Learning rate decay function
- Number of epochs (passes through the training data)
- Order of cycling through training examples (random)

# Multi-class perceptrons

- *One-vs-others* framework: Need to keep a weight vector  $\mathbf{w}_c$  for each class  $c$
- Decision rule:  $c = \operatorname{argmax}_c \mathbf{w}_c \cdot \mathbf{x}$
- Update rule: suppose example from class  $c$  gets misclassified as  $c'$ 
  - Update for  $c$ :  $\mathbf{w}_c \leftarrow \mathbf{w}_c + \alpha \mathbf{x}$
  - Update for  $c'$ :  $\mathbf{w}_{c'} \leftarrow \mathbf{w}_{c'} - \alpha \mathbf{x}$

# Review: Multi-class perceptrons

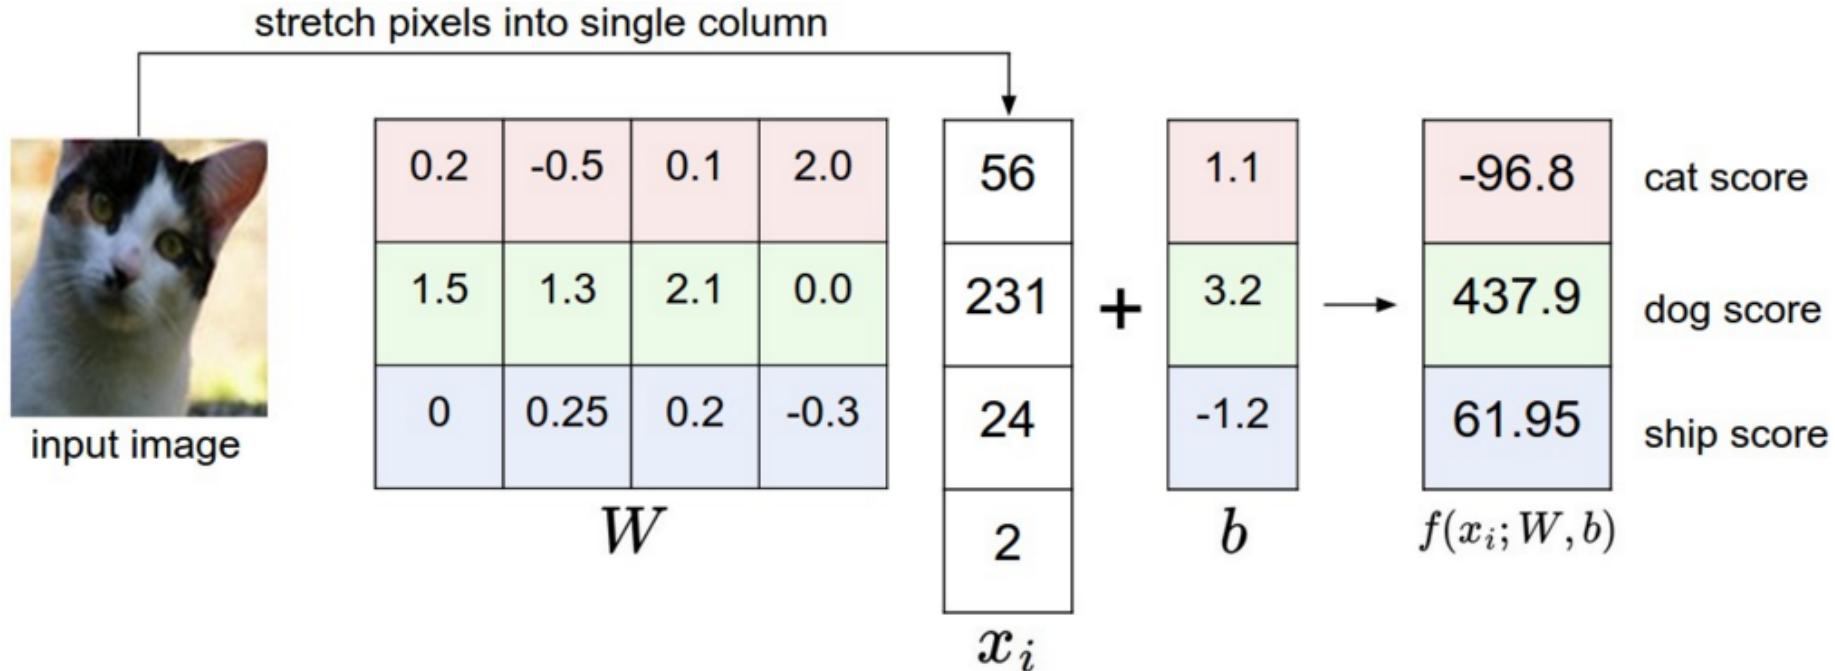
- One-vs-others framework: Need to keep a weight vector  $\mathbf{w}_c$  for each class c
- Decision rule:  $c = \operatorname{argmax}_c \mathbf{w}_c \cdot \mathbf{x}$



Softmax:

$$P(c | \mathbf{x}) = \frac{\exp(\mathbf{w}_c \cdot \mathbf{x})}{\sum_{k=1}^C \exp(\mathbf{w}_k \cdot \mathbf{x})}$$

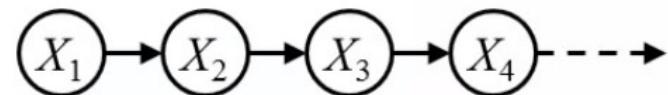
# Visualizing linear classifiers



Source: Andrej Karpathy, <http://cs231n.github.io/linear-classify/>

# Markov Models

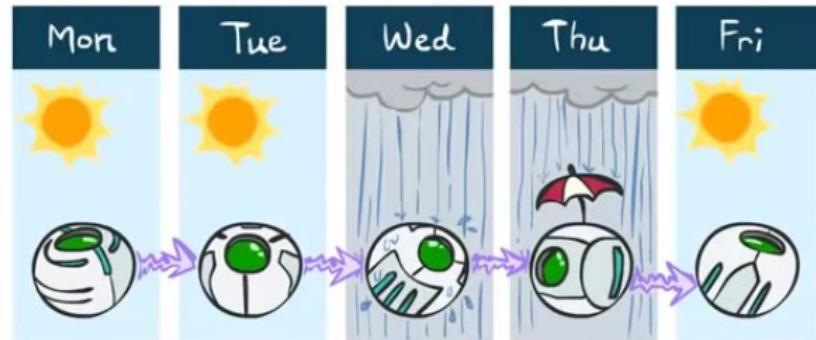
- Value of  $X$  at a given time is called the **state**



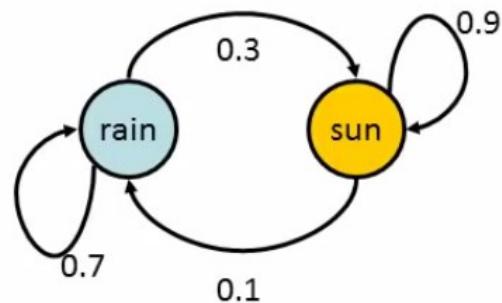
# Example Markov Chain: Weather

- States:  $X = \{\text{rain, sun}\}$
- Initial distribution: 1.0 sun
- CPT  $P(X_t | X_{t-1})$ : 

$X_{t-1}$	$X_t$	$P(X_t   X_{t-1})$
sun	sun	0.9
sun	rain	0.1
rain	sun	0.3
rain	rain	0.7



Two new ways of representing the same CPT

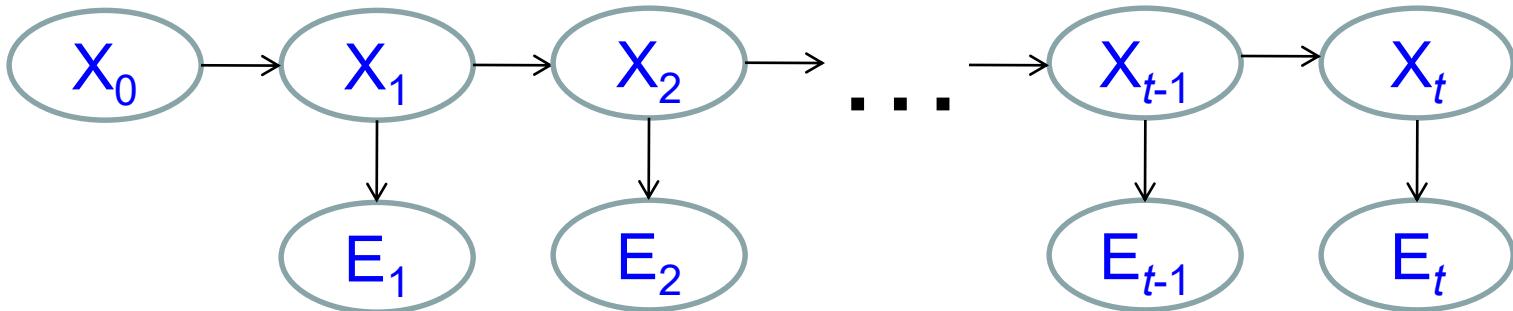


# Probabilistic reasoning over time

- So far, we've mostly dealt with *episodic* environments
  - Exceptions: games with multiple moves, planning
- In particular, the Bayesian networks we've seen so far describe static situations
  - Each random variable gets a single fixed value in a single problem instance
- Now we consider the problem of describing probabilistic environments that evolve over time
  - Examples: robot localization, tracking, speech, ...

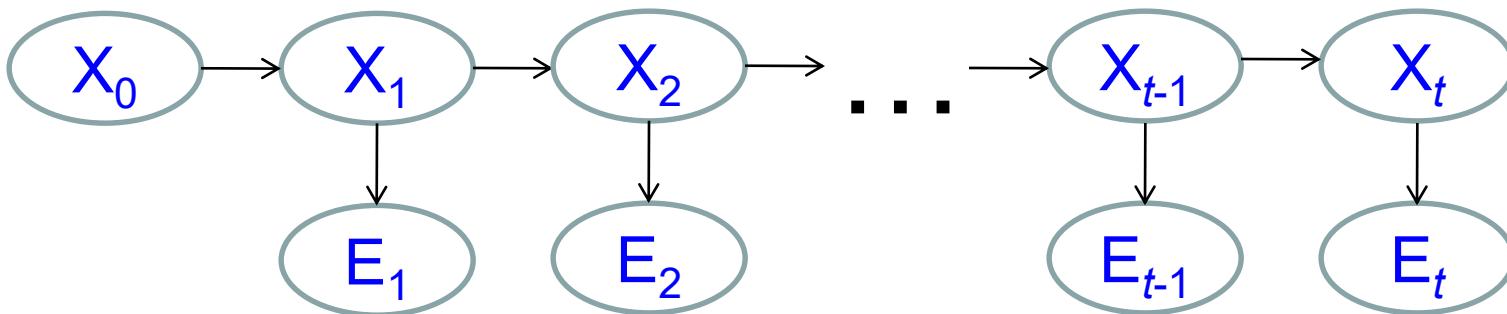
# Hidden Markov Models

- At each time slice  $t$ , the state of the world is described by an unobservable variable  $X_t$  and an observable *evidence* variable  $E_t$
- **Transition model:** distribution over the current state given the whole past history:  
 $P(X_t | X_0, \dots, X_{t-1}) = P(X_t | X_{0:t-1})$
- **Observation model:**  $P(E_t | X_{0:t}, E_{1:t-1})$



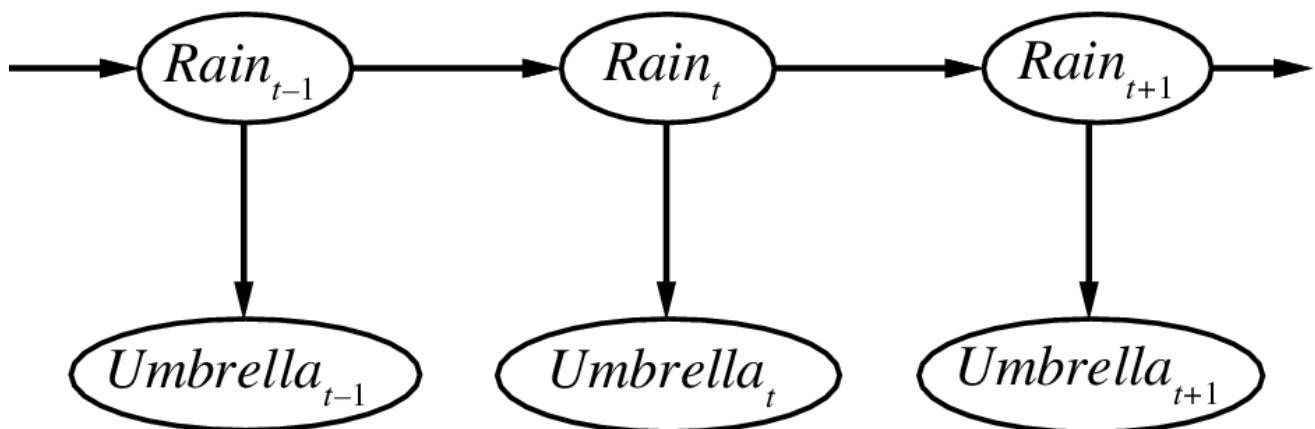
# Hidden Markov Models

- **Markov assumption** (first order)
  - The current state is conditionally independent of all the other states given the state in the previous time step
  - What does  $P(X_t | X_{0:t-1})$  simplify to?  
$$P(X_t | X_{0:t-1}) = P(X_t | X_{t-1})$$
- Markov assumption for observations
  - The evidence at time  $t$  depends only on the state at time  $t$
  - What does  $P(E_t | X_{0:t}, E_{1:t-1})$  simplify to?  
$$P(E_t | X_{0:t}, E_{1:t-1}) = P(E_t | X_t)$$



# Example

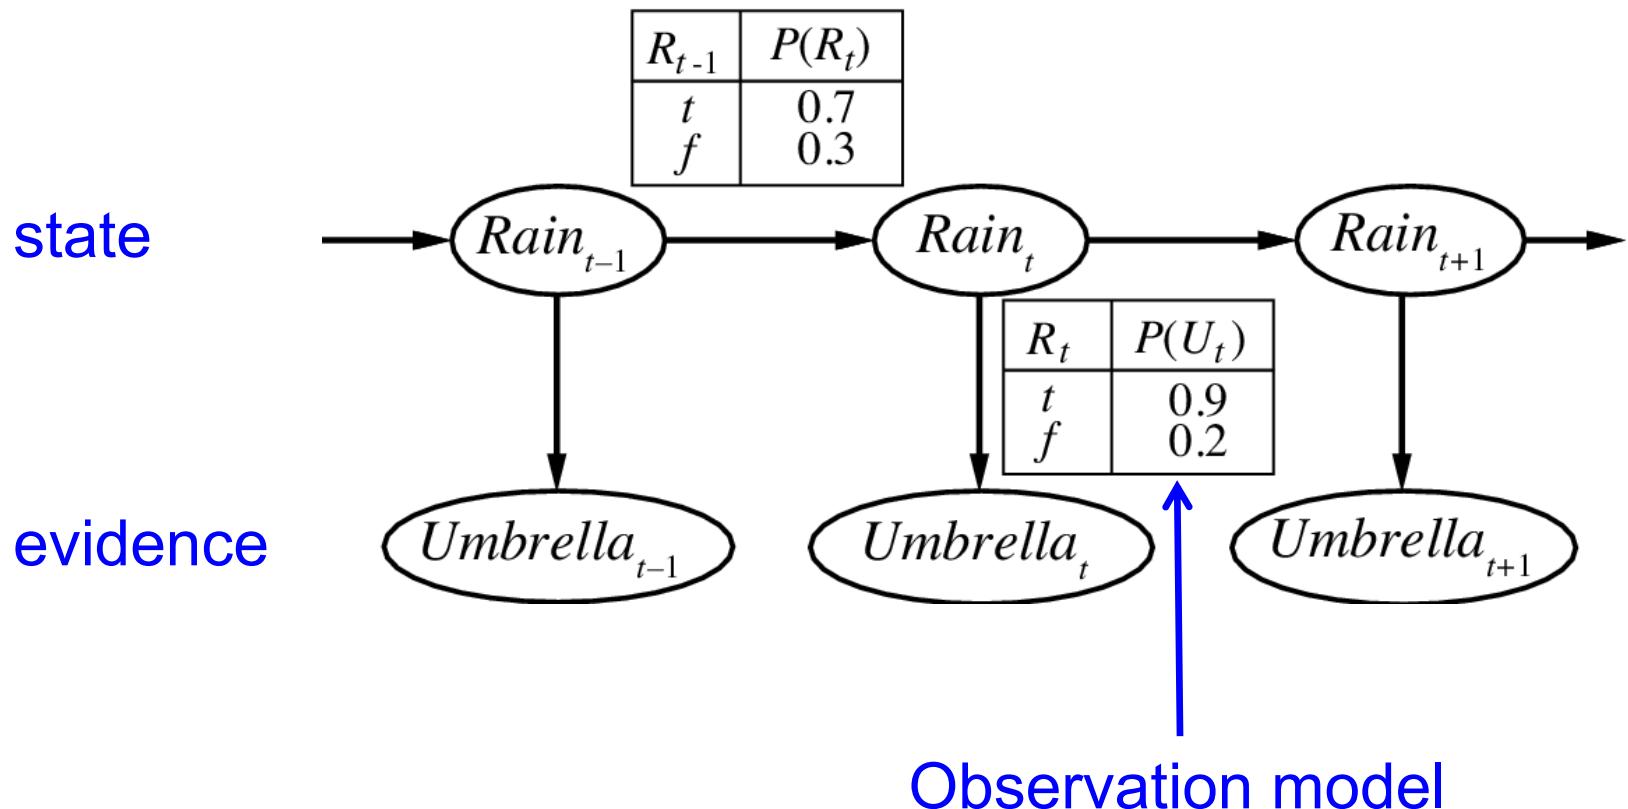
state



evidence

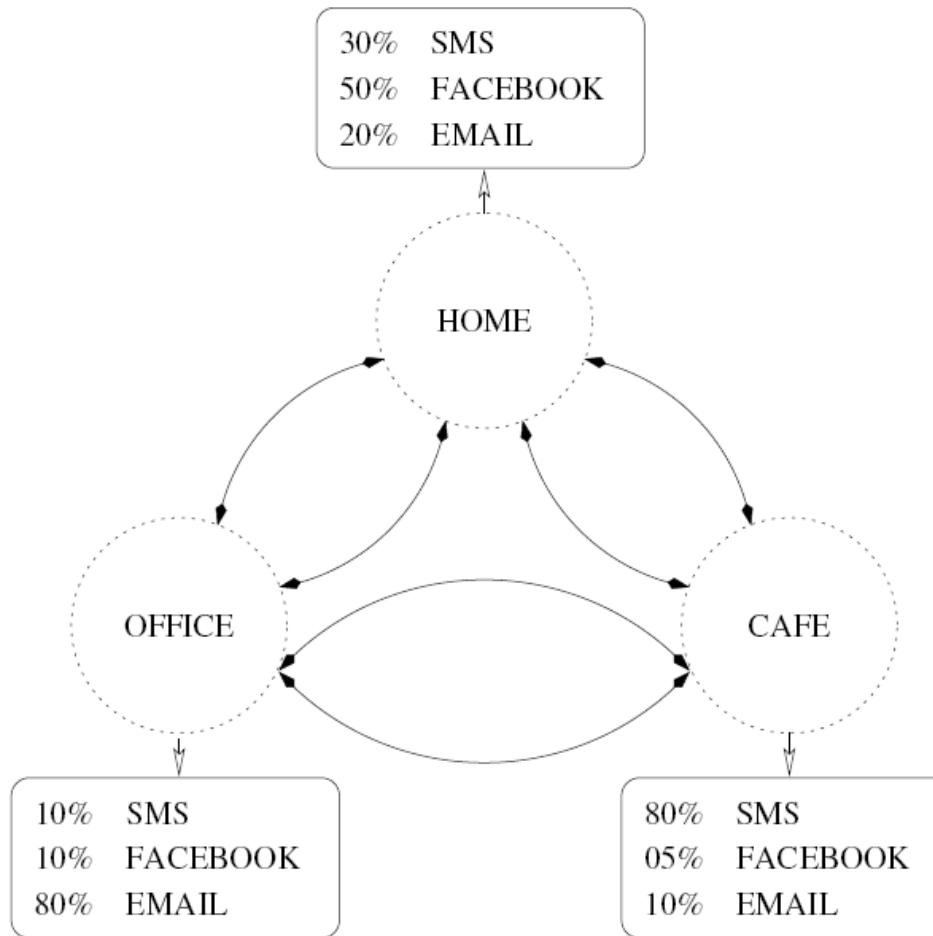
# Example

Transition model



# Another example

- **States:**  $X = \{\text{home, office, cafe}\}$
- **Observations:**  $E = \{\text{sms, facebook, email}\}$



Transition Probabilities

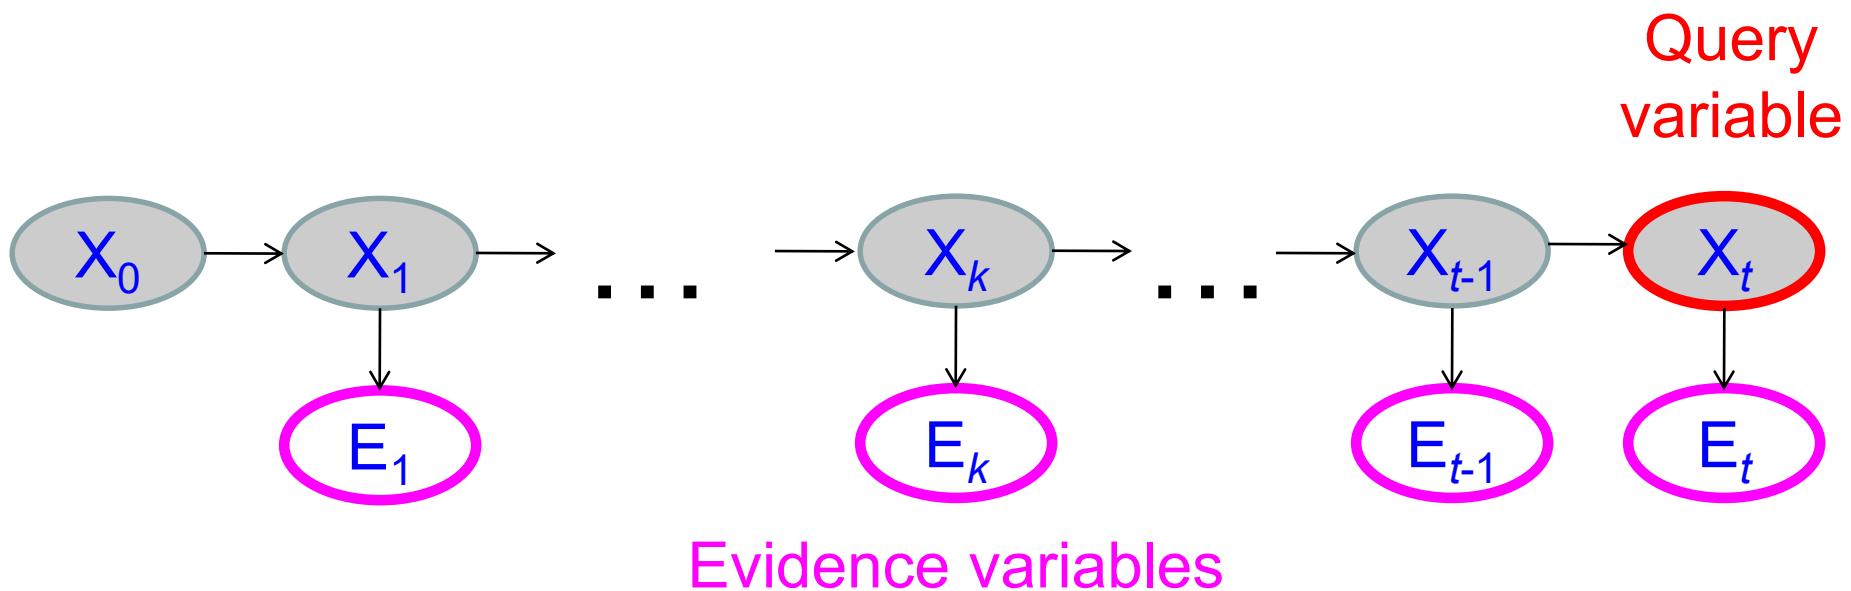
	home	office	cafe
home	0.2	0.6	0.2
office	0.5	0.2	0.3
cafe	0.2	0.8	0.0

Emission Probabilities

	sms	facebook	email
home	0.3	0.5	0.2
office	0.1	0.1	0.8
cafe	0.8	0.1	0.1

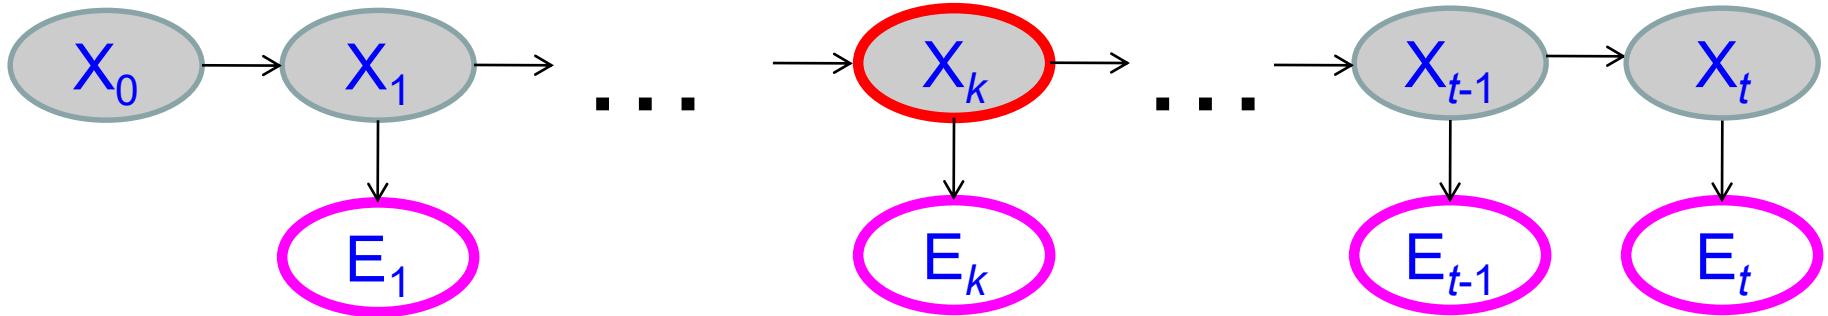
# HMM inference tasks

- **Filtering:** what is the distribution over the current state  $X_t$  given all the evidence so far,  $e_{1:t}$ ?
  - The forward algorithm



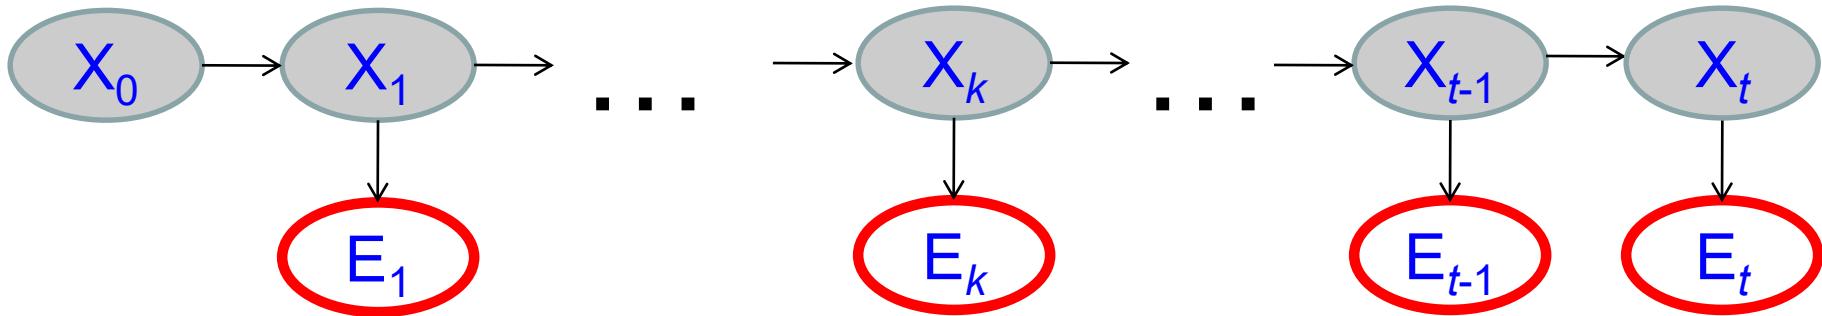
# HMM inference tasks

- **Filtering:** what is the distribution over the current state  $X_t$  given all the evidence so far,  $e_{1:t}$  ?
- **Smoothing:** what is the distribution of some state  $X_k$  given the entire observation sequence  $e_{1:t}$ ?
  - The forward-backward algorithm



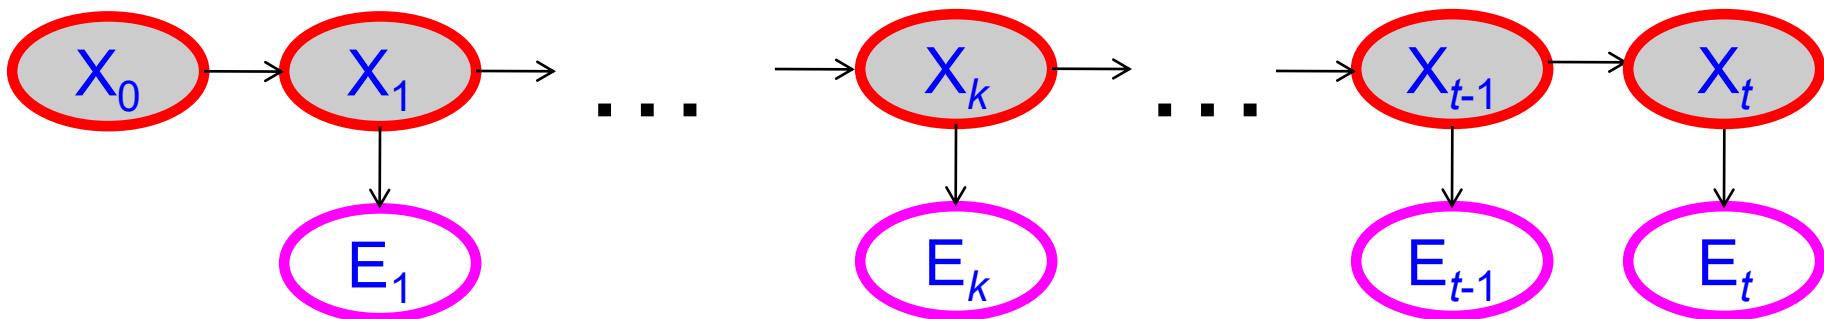
# HMM inference tasks

- **Filtering:** what is the distribution over the current state  $X_t$  given all the evidence so far,  $e_{1:t}$  ?
- **Smoothing:** what is the distribution of some state  $X_k$  given the entire observation sequence  $e_{1:t}$ ?
- **Evaluation:** compute the probability of a given observation sequence  $e_{1:t}$



# HMM inference tasks

- **Filtering:** what is the distribution over the current state  $X_t$  given all the evidence so far,  $e_{1:t}$
- **Smoothing:** what is the distribution of some state  $X_k$  given the entire observation sequence  $e_{1:t}$ ?
- **Evaluation:** compute the probability of a given observation sequence  $e_{1:t}$
- **Decoding:** what is the most likely state sequence  $X_{0:t}$  given the observation sequence  $e_{1:t}$ ?
  - The Viterbi algorithm



# HMM Learning and Inference

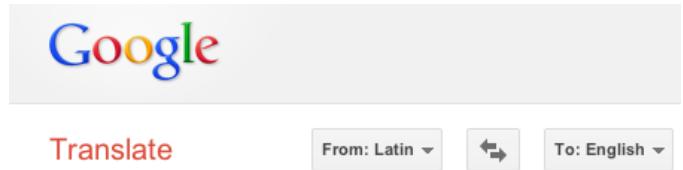
- Inference tasks
  - **Filtering:** what is the distribution over the current state  $X_t$  given all the evidence so far,  $e_{1:t}$
  - **Smoothing:** what is the distribution of some state  $X_k$  given the entire observation sequence  $e_{1:t}$ ?
  - **Evaluation:** compute the probability of a given observation sequence  $e_{1:t}$
  - **Decoding:** what is the most likely state sequence  $X_{0:t}$  given the observation sequence  $e_{1:t}$ ?
- Learning
  - Given a training sample of sequences, learn the model parameters (transition and emission probabilities)
    - EM algorithm

# Applications of HMMs

- Speech recognition HMMs:
  - Observations are acoustic signals (continuous valued)
  - States are specific positions in specific words (so, tens of thousands)



- Machine translation HMMs:
  - Observations are words (tens of thousands)
  - States are translation options

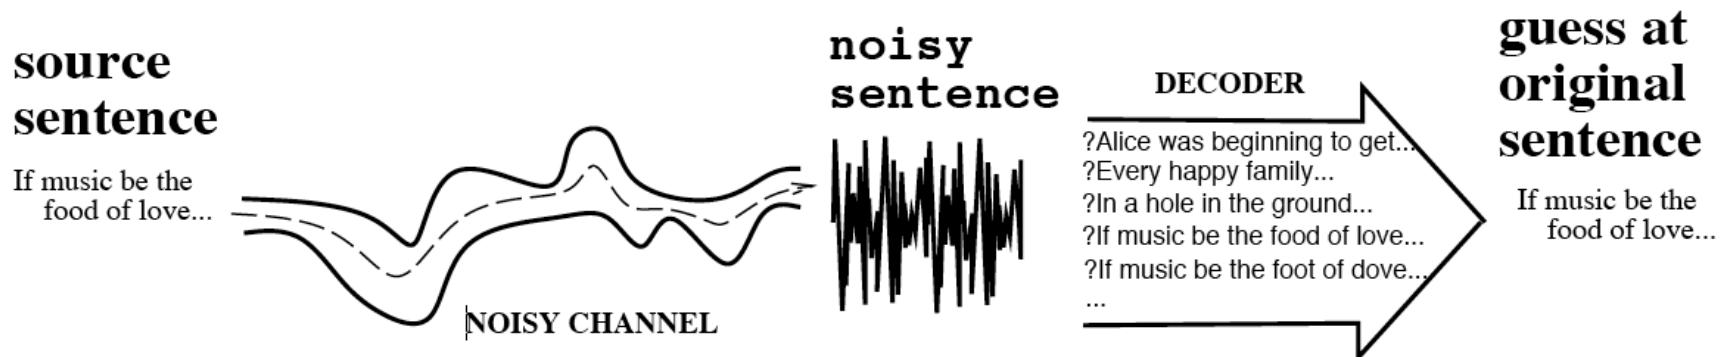


- Robot tracking:
  - Observations are range readings (continuous)
  - States are positions on a map (continuous)

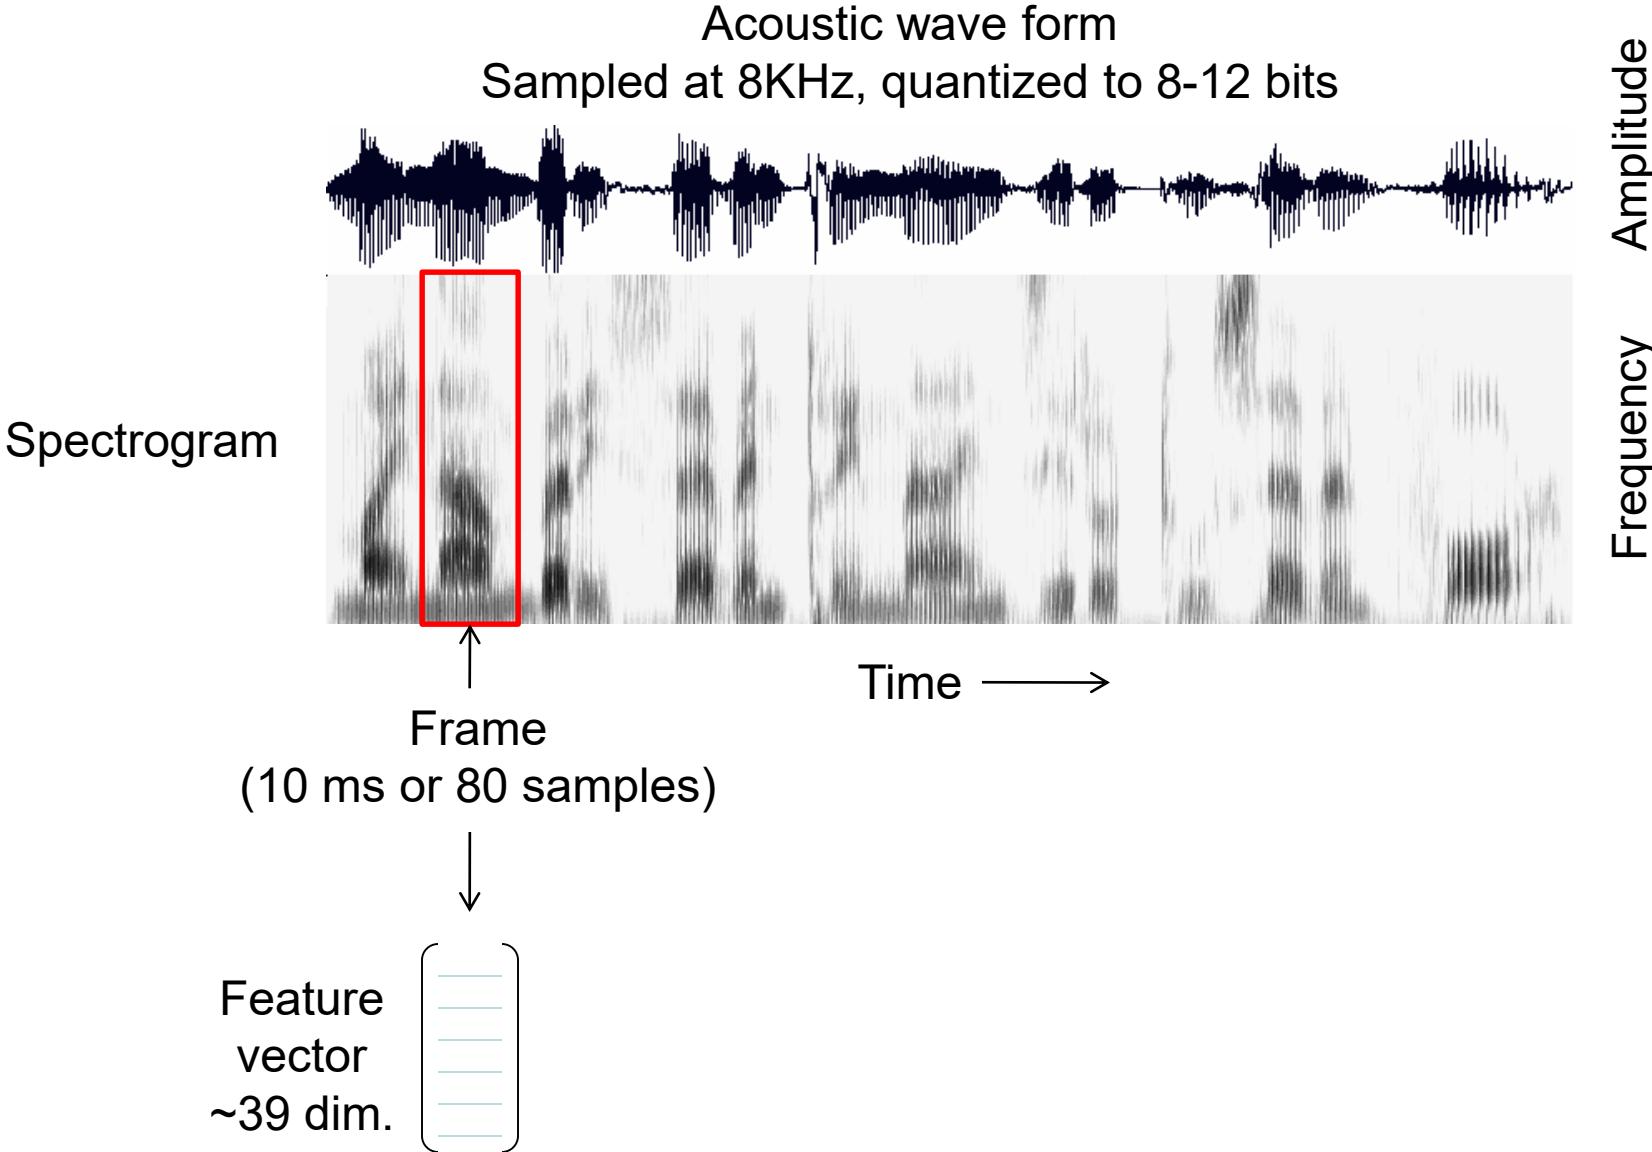


# Application of HMMs: Speech recognition

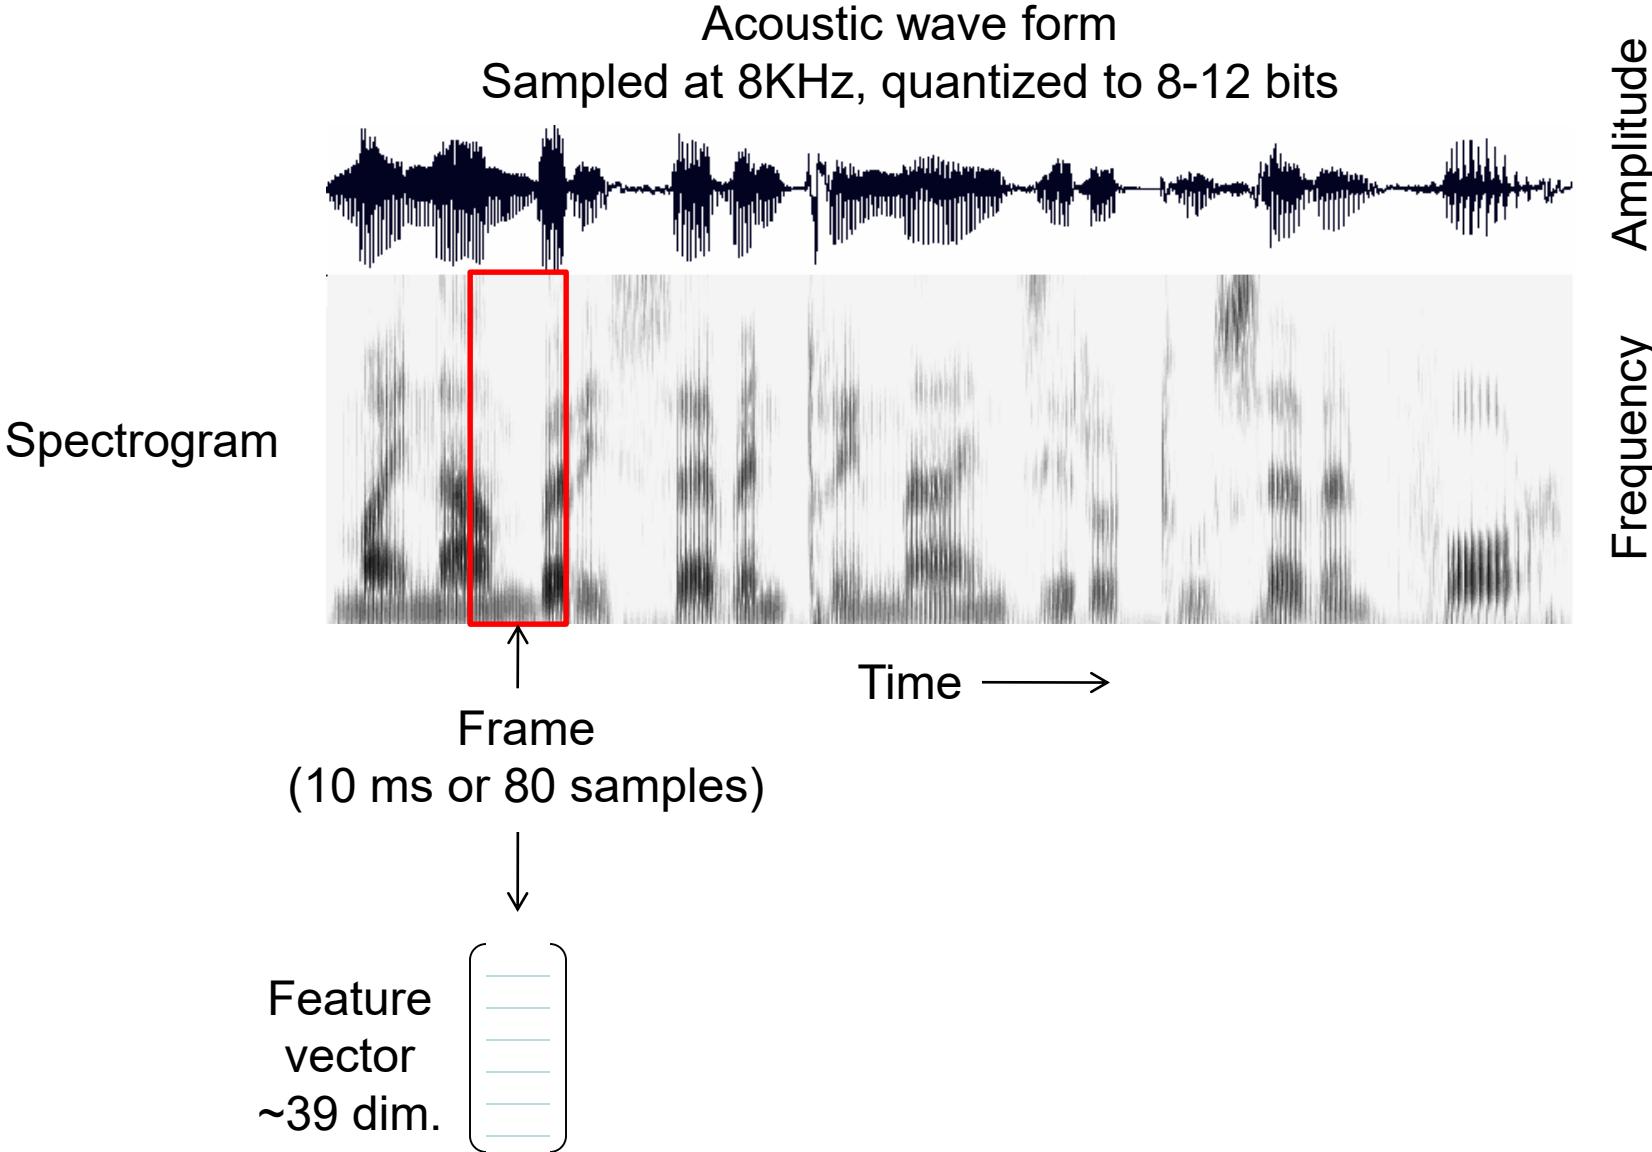
- “Noisy channel” model of speech



# Speech feature extraction



# Speech feature extraction



# Phonetic model

- **Phones:** speech sounds
- **Phonemes:** groups of speech sounds that have a unique meaning/function in a language (e.g., there are several different ways to pronounce “t”)

# Phonetic model

IPA	ARPAbet		IPA	ARPAbet
Symbol	Symbol	Word	Transcription	Transcription
[p]	[p]	parsley	['parsli]	[p a a r s l i y]
[t]	[t]	tarragon	['tærəgən]	[t ae r ax g aa n]
[k]	[k]	catnip	['kætnip]	[k ae t n ix p]
[b]	[b]	bay	[beɪ]	[b ey]
[d]	[d]	dill	[dɪl]	[d ih l]
[g]	[g]	garlic	['garlik]	[g aa r l ix k]
[m]	[m]	mint	[mɪnt]	[m ih n t]
[n]	[n]	nutmeg	['nʌtmeg]	[n ah t m eh g]
[ŋ]	[ŋ]	ginseng	['dʒɪnsɪŋ]	[jh ih n s ix ng]
[f]	[f]	fennel	['fenl]	[f eh n el]
[v]	[v]	clove	[klouv]	[k l ow v]
[θ]	[th]	thistle	['θɪsl]	[th ih s el]
[ð]	[dh]	heather	['heðə]	[h eh dh axr]
[s]	[s]	sage	[seɪdʒ]	[s ey jh]
[z]	[z]	hazelnut	['herzlnat]	[h ey z el n ah t]
[ʃ]	[sh]	squash	[skwɑʃ]	[s k w a sh]
[ʒ]	[zh]	ambrosia	['æm'brouʒə]	[ae m b r ow zh ax]
[tʃ]	[ch]	chicory	['tʃɪkɔɪ]	[ch ih k axr iy]
[dʒ]	[jh]	sage	[seɪdʒ]	[s ey jh]
[l]	[l]	licorice	['likɔɪʃ]	[l ih k axr ix sh]
[w]	[w]	kiwi	['kiwi]	[k iy w iy]
[r]	[r]	parsley	['parsli]	[p a a r s l iy]
[j]	[y]	yew	[yu]	[y uw]
[h]	[h]	horseradish	['hɔrsrædɪʃ]	[h ao r s r ae d ih sh]
[ʔ]	[q]	uh-oh	[?ʌ?ou]	[q ah q ow]
[r̥]	[dx]	butter	['bʌrə]	[b ah dx axr]
[ɹ̥]	[nx]	wintergreen	['wɪɾəgrɪn]	[w ih nx axr g r i n ]
[l̥]	[el]	thistle	['θɪsl]	[th ih s el]

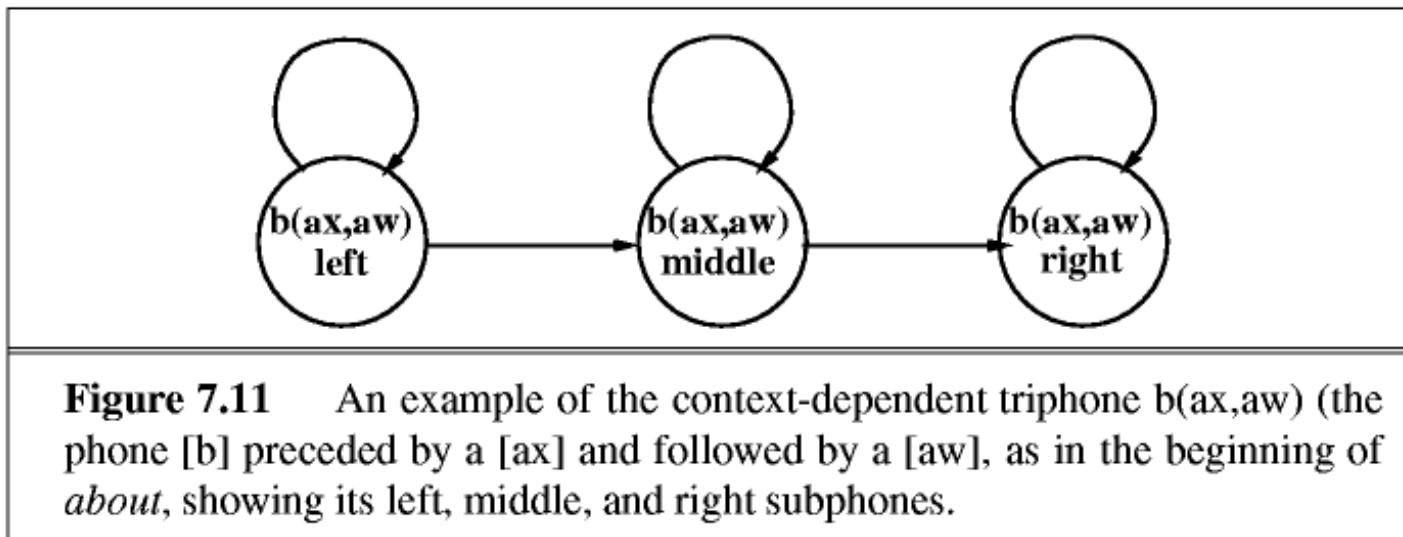
Figure 4.1 IPA and ARPAbet symbols for transcription of English consonants.

IPA	ARPAbet		IPA	ARPAbet
Symbol	Symbol	Word	Transcription	Transcription
[i]	[iy]	lily	['lili]	[l ih l iy]
[ɪ]	[ih]	lily	['lili]	[l ih l iy]
[eɪ]	[ey]	daisy	['deɪzi]	[d ey z i]
[ɛ]	[eh]	poinsettia	[pom'seriə]	[p oy n s eh dx iy ax]
[æ]	[ae]	aster	['æstə]	[ae s t axr]
[ɑ]	[aa]	poppy	['papi]	[p aa p i]
[ɔ]	[ao]	orchid	['ɔrkid]	[ao r k ix d]
[ʊ]	[uh]	woodruff	['wudrəf]	[w uh d r ah f]
[oʊ]	[ow]	lotus	['lourəs]	[l ow dx ax s]
[ʊ]	[uw]	tulip	['tulip]	[t uw l ix p]
[ʌ]	[uh]	buttercup	['bʌtər,kap]	[b uh dx axr k uh p]
[ɹ̥]	[er]	bird	['bəd]	[b er d]
[aɪ]	[ay]	iris	['aɪris]	[ay r ix s]
[aʊ]	[aw]	sunflower	['sʌnflauə]	[s ah n fl aw axr]
[ɔɪ]	[oy]	poinsettia	[pom'seriə]	[p oy n s eh dx iy ax]
[ju]	[y uw]	feverfew	['fivəfju]	[f iy v axr f y u]
[ə]	[ax]	woodruff	['wudrəf]	[w uh d r ax f]
[ə̥]	[axr]	heather	['heðə]	[h eh dh axr]
[i̥]	[ix]	tulip	['tulip]	[t uw l ix p]
[u̥]	[ux]		[]	[]

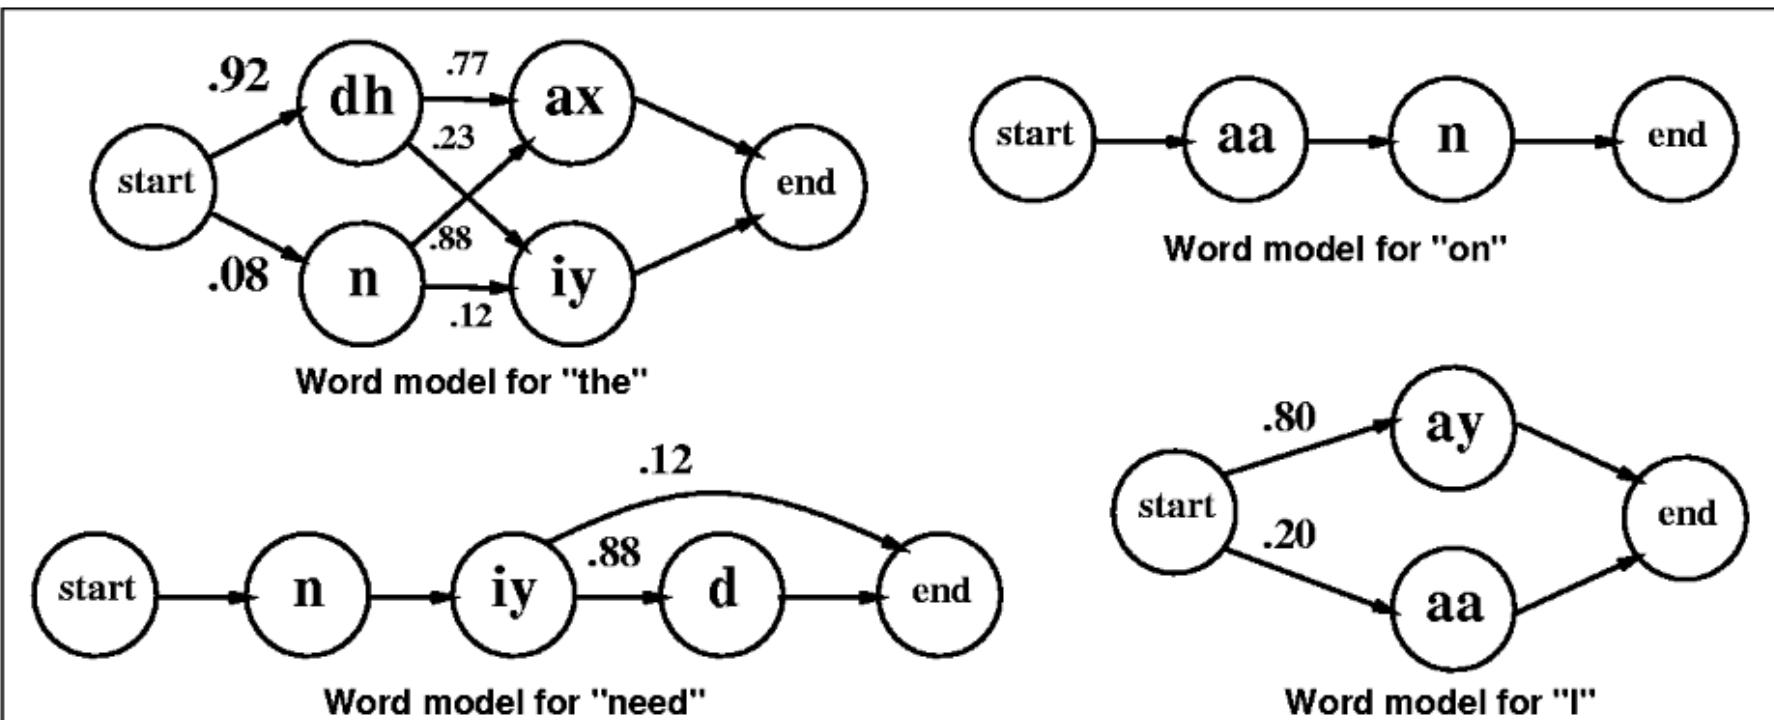
Figure 4.2 IPA and ARPAbet symbols for transcription of English vowels

# HMM models for phones

- HMM states in most speech recognition systems correspond to *subphones*
  - There are around 60 phones and as many as  $60^3$  context-dependent *tripphones*



# HMM models for words



**Figure 7.5** Pronunciation networks for the words *I*, *on*, *need*, and *the*. All networks (especially *the*) are significantly simplified.